


# The Effect of Compiler Optimizations on High-Level Synthesis for FPGAs

Qijing Huang, Ruolong Lian, Andrew Canis, Jongsok Choi,  
Ryan Xi, Stephen Brown, Jason Anderson

IEEE FCCM 2013

Seattle, Washington, April 29, 2013



Dept. of Electrical and Computer Engineering  
University of Toronto

# Motivation

- Compilers execute optimization passes:
  - dead code elim, constant propagation, loop rotation, inlining, exlining, etc.
- LLVM incorporates >50 different passes
- The familiar optimization levels -O2, -O3 correspond to recipes of such passes
- The levels are designed to optimize SW not HW
- **Question:** Which passes are good for HLS?


# Analysis of Compiler Passes

- We applied each pass (of 56) in isolation
- Compare with -O0 (no optimization)
- LegUp: Open-source high-level synthesis
  - <http://legup.eecg.toronto.edu>
- Target Cyclone II FPGA
- 11 CHStone benchmarks
  - Standard high-level synthesis suite
  - <http://www.ertl.jp/chstone>

# Analysis of Compiler Passes (2)

- Consider 4 hardware metrics:
  - Cycle latency (# of clock cycles)
  - FMax
  - Wall-clock time ( $= \text{Cycle latency} / \text{FMax}$ )
  - # of LEs (area)
- Large # of ModelSim + Quartus II runs
  - $56 \times 11 = 616$  synthesis, place, route runs
  - Used SciNet (Canadian cloud computing system)

# Representative Results: Cycle Latency



# Overall Results Summary

Across 56 passes		Clock cycles	FMax	Wall-clock time	LEs
	Min.	0.72	0.76	0.92	0.93
	Max.	1.83	1.05	2.24	1.34
	St. Dev.	0.12	0.04	0.17	0.05

- Widest swings observed in clock cycles, wall-clock time.

# Custom Recipes

- Create a recipe of passes to the “taste” of each CHStone benchmark
- For each benchmark (11), for each HW metric (4), create a recipe containing **just** the beneficial passes (in alphabetical order)



# Custom Recipes

- Sample results for `dfmul` benchmark:

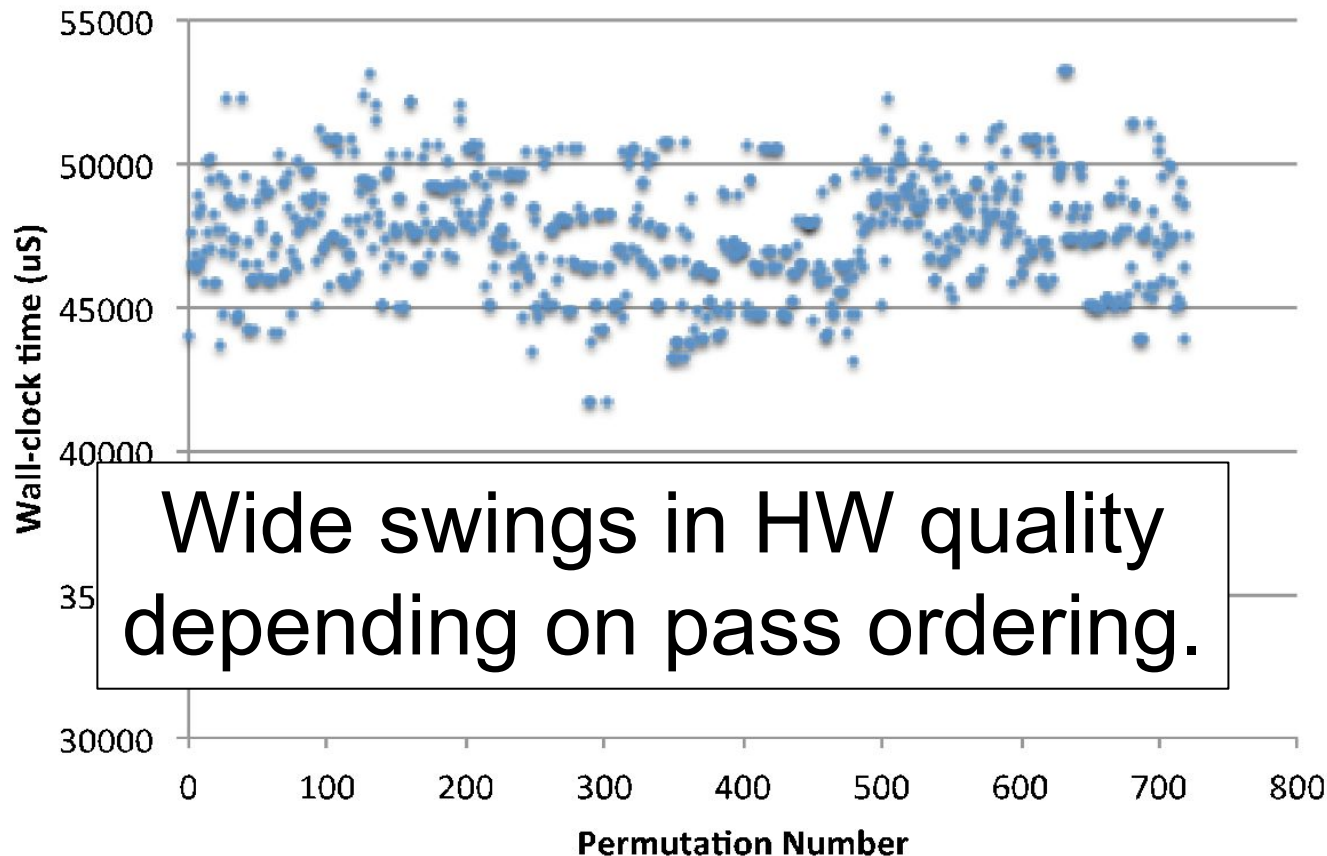
	Recipe	Clockcycles	FMax	Wall-clock time	LEs
clock					0.2
	FMax	1.42	<b>1.02</b>	1.39	1.29
	Wall-clock time	0.92	1.01	<b>0.91</b>	0.93
	LEs	1.02	0.99	1.02	<b>0.91</b>

Significant potential to improve upon `-O3` optimization!



# Order Dependence

- Consider **all** orderings of 6 passes shown beneficial to wall-clock time for jpeg benchmark ( $6! = 720$  runs)



# Dynamic Pass Recipes

- “Good” passes to apply:  
highly benchmark dependent
- Impact of passes highly order dependent
- **We need a dynamic  
feedback-driven approach**

# Early Prediction of Cycle Latency

- Basic block (BB): a section of code with a single entry/exit point.

We do not need to execute synthesis, place, route to determine cycle latency!


- Through partial HLS, we can determine the schedule length for each BB.

$$CycleLatency = \sum_{b \in BasicBlocks} \#Execs(b) \times SchedLength(b)$$


# Dynamic Optimization Flow

Unoptimized program

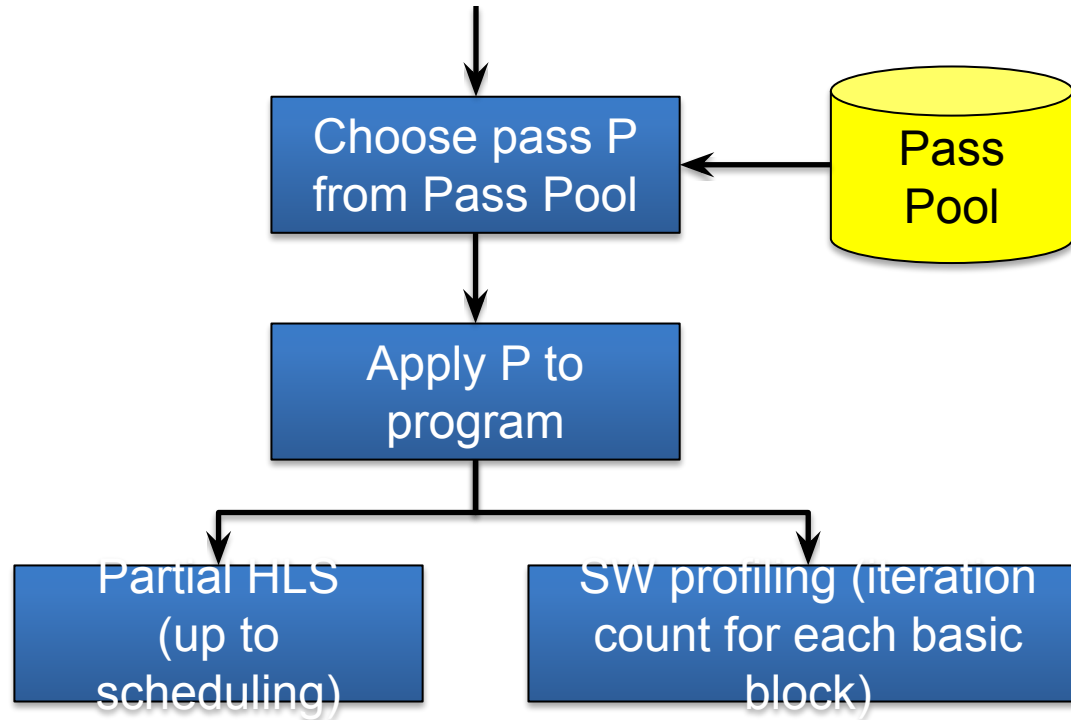
Unoptimized program



Unoptimized program

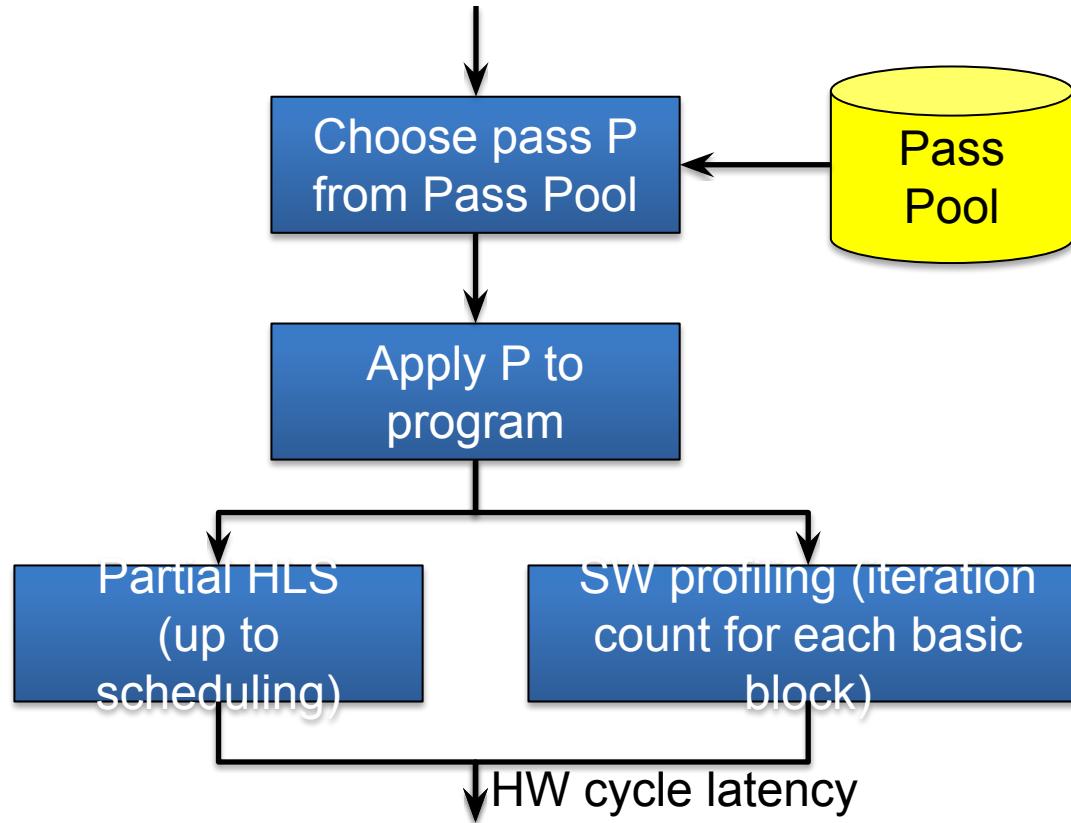


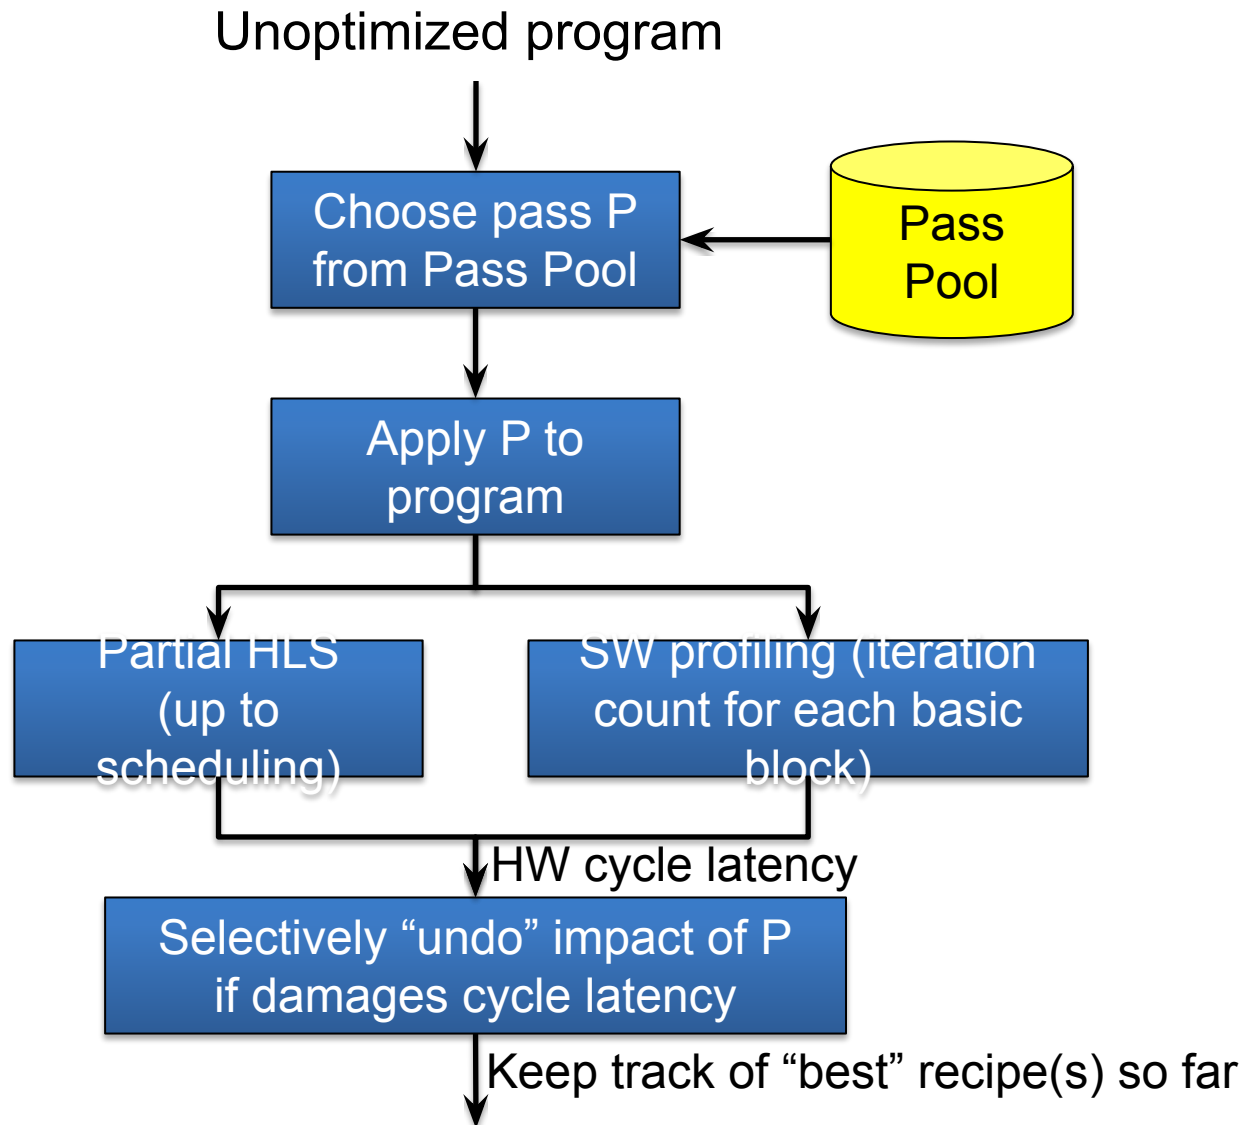
Unoptimized program



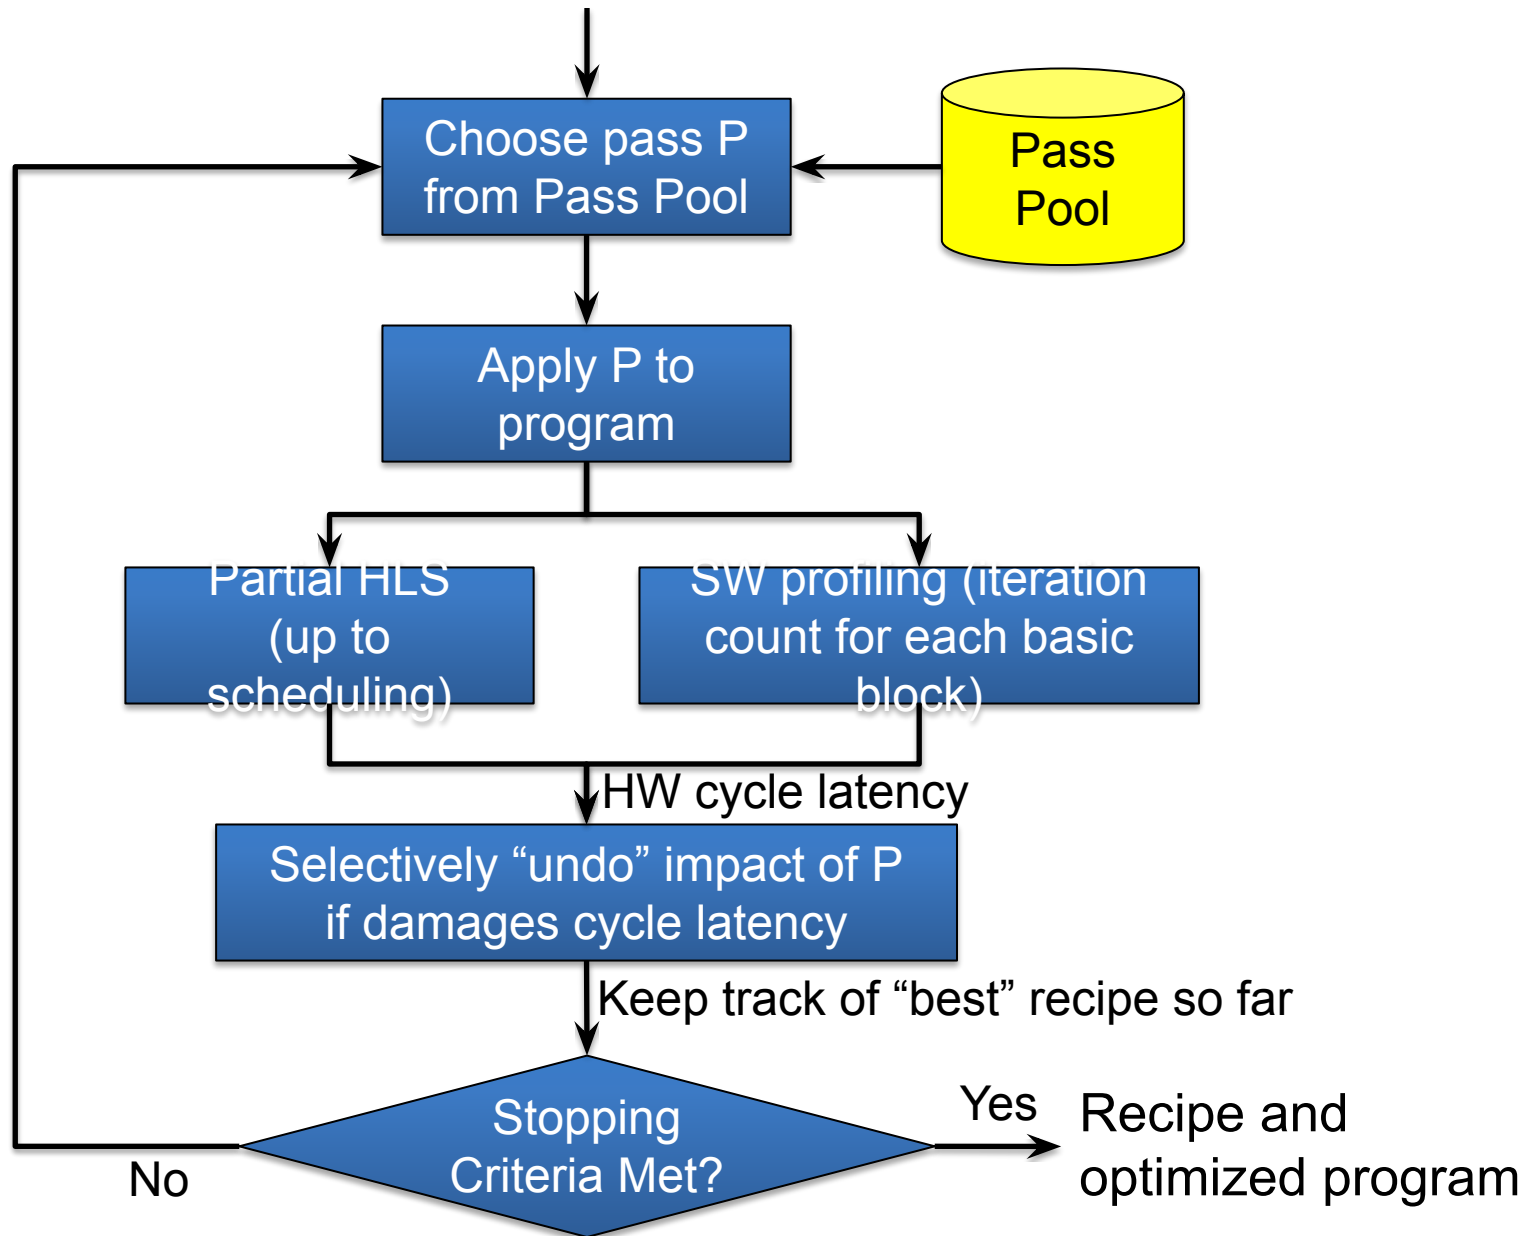


Unoptimized program





Unoptimized program



# Pass Pool

- Contains 41 of 56 passes as follows:
  - Any pass that showed impact on any benchmark vs. no optimization (-O0)
  - Any pass in -O3 that showed impact on any benchmark when removed from -O3
  - Passes not in -O3 that showed no impact in isolation (they may show impact when combined)
- Passes ordered based on “pairs” analysis
  - Analyzed all pairs of passes to determine “preferred” order

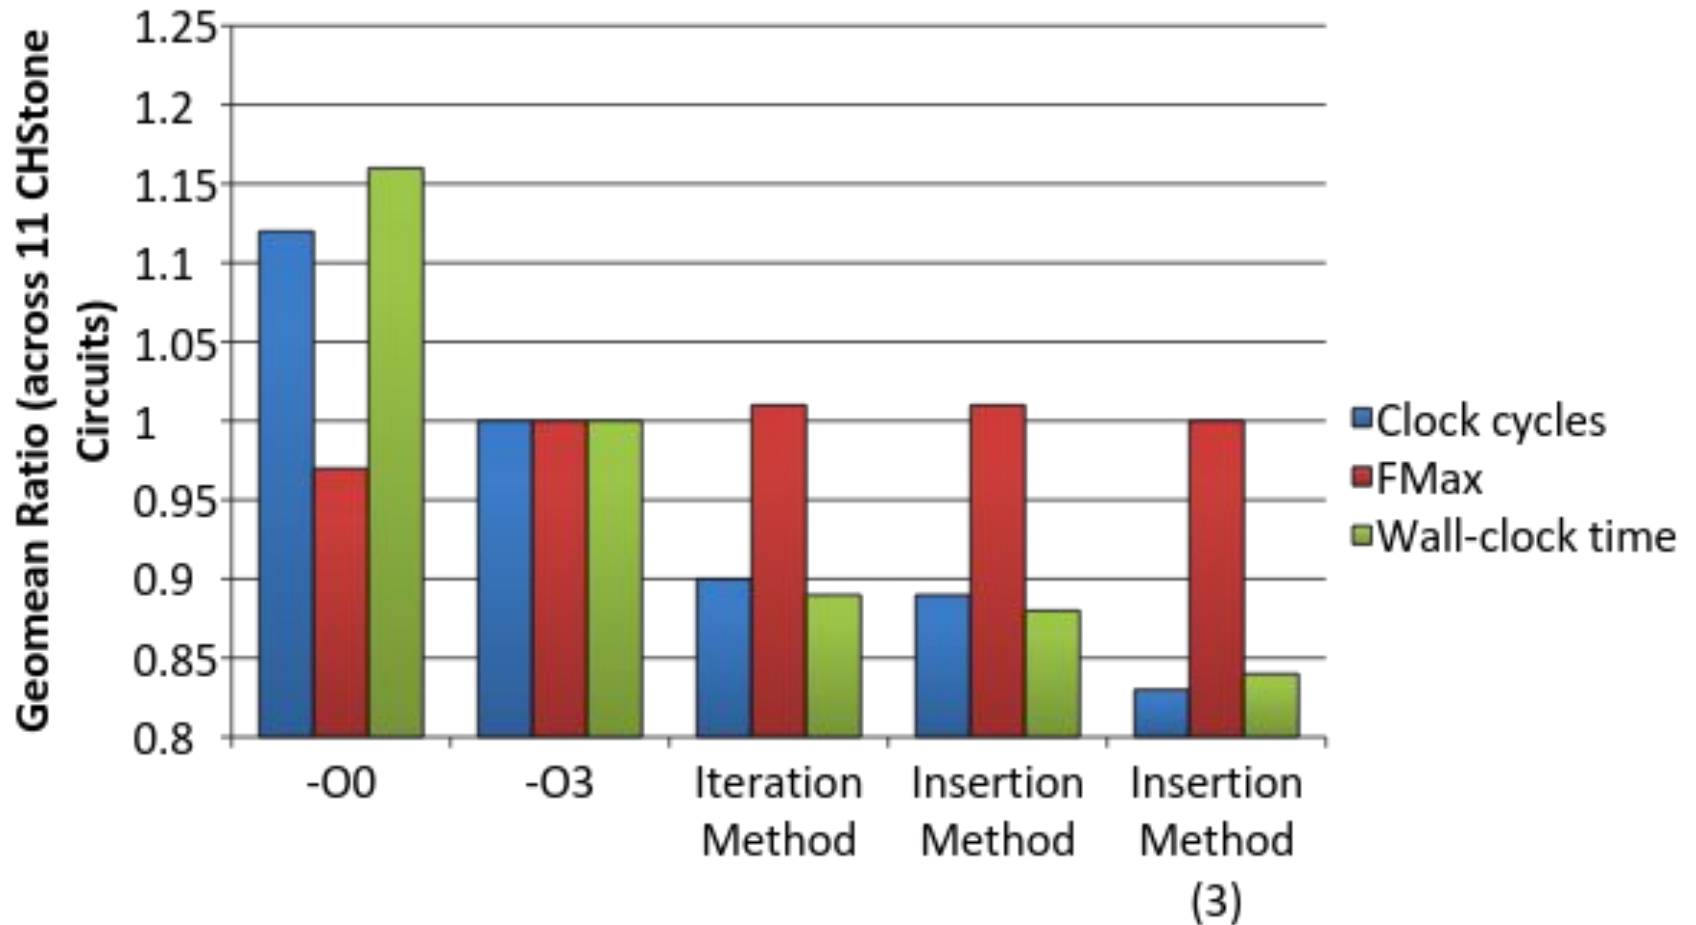
# Applying a Pass: 3 Ways

1. **Iteration method:**
  - i) Passes in pass pool are sorted (based on pairs analysis).
  - ii) Selected pass is added to end of “best” recipe
2. **Insertion method:** all possible insertion points in “best” recipe are considered for each selected pass
  - Order independent
3. **Insertion method (3):** same as #2, but store top 3 “best” recipes encountered


# Stopping Criteria

- Stop if no “walk” through pass pool improves results
- At most, 3 “walks” through the pass pool

# Results (Speed)

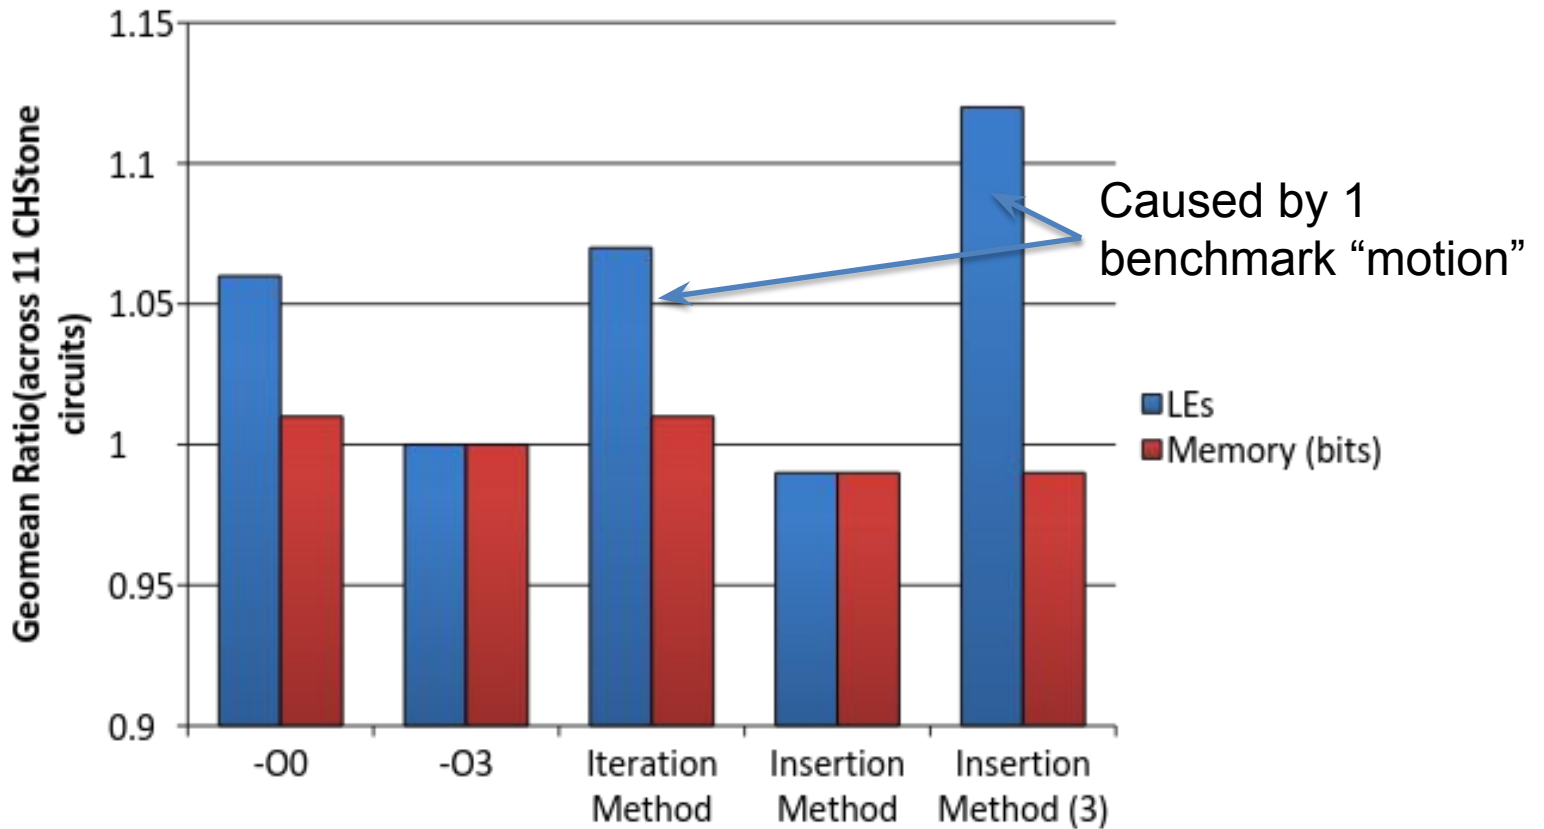


# Results (Speed)





# Results (Area)



# Summary

- Compiler passes have strong influence on quality of HLS-generated HW
- Impact of passes highly order dependent
- Significant potential to raise HW quality via dynamic feedback-driven approach:
  - Circuit wall-clock time can be reduced by 11-16%, on average, depending on run-time allowed (vs. -O3 recipe)

# Future Work

- Our approach is somewhat “brute force”
  - Heuristics to “predict” which passes benefit certain code styles
- Apply different pass recipes to different parts of code
- New passes specifically targeted to HLS hardware generation

Questions?

[legup.eecg.toronto.edu](http://legup.eecg.toronto.edu)