



# Centrifuge: Evaluating full-system HLS-generated heterogeneous-accelerator SoCs using FPGA-Acceleration

**Qijing Huang**, Christopher Yarp, Sagar Karandikar, Nathan Pemberton,  
Benjamin Brock, Liang Ma<sup>†</sup>, Guohao Dai<sup>‡</sup>, Robert Quitt,  
Krsti Asanović, John Wawrzynek

University of California, Berkeley

<sup>†</sup>Politecnico di Torino

<sup>‡</sup>Tsinghua University

# Outline

**1. Motivation**

**2. Centrifuge  
Design Flow**

**3. Centrifuge  
Design Space**

**4. Case Studies**

# Motivation

What is a good methodology for designing  
SoCs with many accelerators?

High-Performance System  
Low NRE costs  
Short Time-to-market

# Accelerator Design Flow

1. Workload Characterization
2. Accelerator Modeling
  - Analytical
  - Transaction-based
3. RTL Development
4. Emulation and Verification
5. Chip Tapeout
6. Software Development

# Accelerator Design Flow

## 1. Workload Characterization

### 2. Accelerator Modeling

- Analytical
- Transaction-based

### 3. RTL Development

### 4. Emulation and Verification

### 5. Chip Tapeout

### 6. Software Development

- Profiling on existing systems

Hotspot detection is only valid for the current system

- ISA-independent IR trace analysis

Lack of detailed information for hotspot detection

# Accelerator Design Flow

1. Workload Characterization
2. Accelerator Modeling
  - Analytical
  - Transaction-based
3. RTL Development
4. Emulation and Verification
5. Chip Tapeout
6. Software Development

- Roofline Models
- Dynamic Data Dependence Graphs (DDDG) Analysis

Inaccurate

Lack of insights for architectural variations in real implementation

# Accelerator Design Flow

## 1. Workload Characterization

## 2. Accelerator Modeling

- Analytical
- Transaction-based

## 3. RTL Development

## 4. Emulation and Verification

## 5. Chip Tapeout

## 6. Software Development

- A. Build Software Models for Hardware Components
- B. Gather Costs from RTL Synthesis
- C. Verify Functionality

Engineering Intensive

# Accelerator Design Flow

1. Workload Characterization
2. Accelerator Modeling
  - Analytical
  - Transaction-based
3. RTL Development
4. Emulation and Verification
5. Chip Tapeout
6. Software Development

Missing co-optimization opportunities

# Current Design Flow Complexities

1. Workload Characterization
2. Hardware Modeling and Verification
3. Software Integration

# Our Objectives

1. Tradeoffs informed with full-system evaluation
2. Fast development and verification cycle for both HW/SW
3. Large design space for rapid algorithm-hardware exploration

# Accelerator Design Flow

1. Workload Characterization
2. Accelerator Modeling
  - Analytical
  - Transaction-based
3. RTL Development
4. Emulation and Verification
5. Chip Tapeout
6. Software Development



## Proposal 1: Native Simulation

Run full-stack software with target SoC simulation

# Accelerator Design Flow

1. Workload Characterization
2. Accelerator Modeling
  - Analytical
  - Transaction-based
3. RTL Development
4. Emulation and Verification
5. Chip Tapeout
6. Software Development



## Proposal 1: Native Simulation

Run full-stack software with target SoC simulation

## Proposal 2: Rapid Prototyping

Combine 2 and 3 with one high-level abstraction

# Accelerator Design Flow

1. Workload Characterization
2. Accelerator Modeling
  - Analytical
  - Transaction-based
3. RTL Development
4. Emulation and Verification
5. Chip Tapeout
6. Software Development

Proposal 1: Native Simulation

Run full-stack software with target SoC simulation

Proposal 2: Rapid Prototyping

Combine 2 and 3 with one high-level abstraction

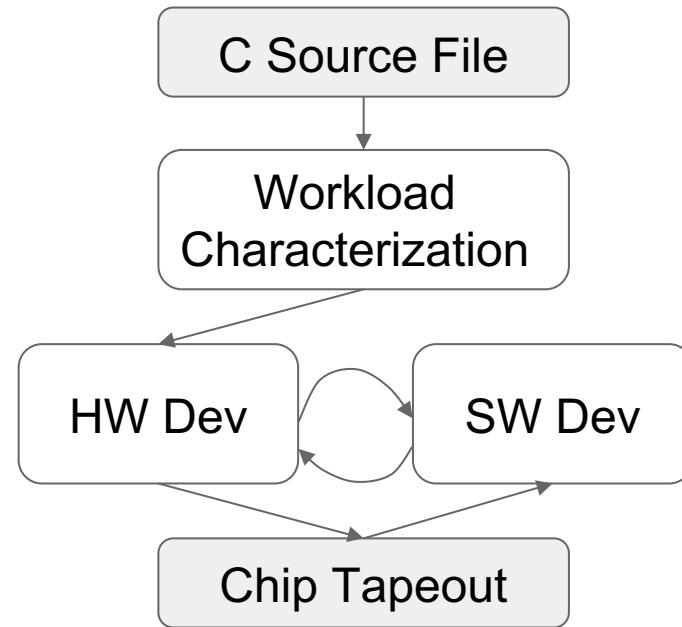
Proposal 3: Agile Development

Improve software stack concurrently with hardware development



# Centrifuge Accelerator Design Flow

1. Hotspot Detection on Native System Simulation
2. Fast Accelerator SoC Generation with HLS
3. Auto Generation of SW Communication Primitives
4. Continue HW and SW Development
5. Chip Tapeout



# Background: FireSim



**FireSim** is an open-source, FPGA-accelerated, cycle-exact, scalable hardware simulator:

- Ingests:
  - RTL Design (e.g. Rocket Chip, BOOM, NVDLA)
  - HW and/or SW IO Models (e.g. UART, Ethernet, DRAM)
  - Workload Descriptions
- Produces:
  - FPGA-accelerated Simulation (Not FPGA prototype)
  - Deployment on Cloud FPGAs

# Background: High-level Synthesis

**High-level Synthesis (HLS)** raises the level abstraction for hardware design and verification:

- Ingests:
  - High-level Algorithmic Description in C/C++/SystemC
- Produces:
  - Hardware RTL, C and RTL Verification Models
- Commercial Tools:



# Centrifuge Design Flow

**Centrifuge** is a unified accelerator design flow that generates the interfaces and design automation scripts that leverages existing tools:

- Injects:
  - High-level Algorithmic Description in C/C++
  - Accelerator Configuration
- Produces:
  - Full Functional Accelerator SoC
  - Corresponding Software Stack
  - Tooling Scripts

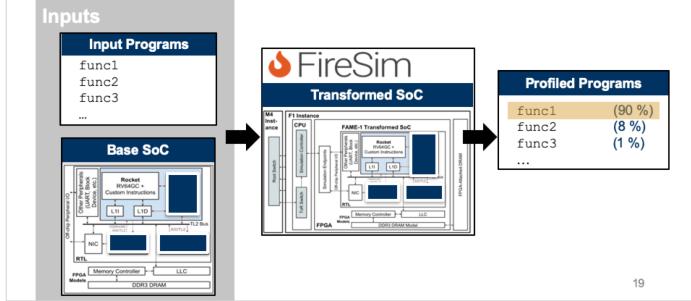
# Centrifuge Design Flow

**Centrifuge** provides the user with:

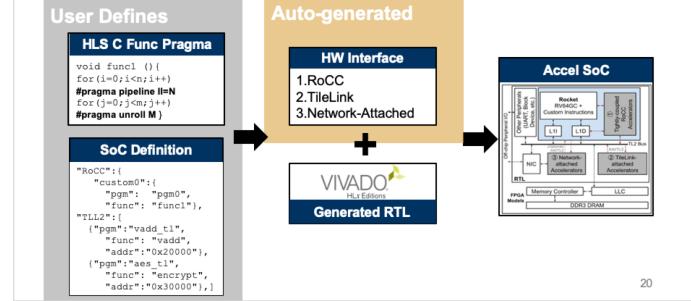
1. Full-system evaluation of the target workload
2. Fast development and verification cycle for both HW/SW
3. Large design space for rapid algorithm-hardware exploration
  - A. *Hardware Integration*
  - B. *Architectural Design Variation*
  - C. *Software Integration*

# Centrifuge Tool Flow

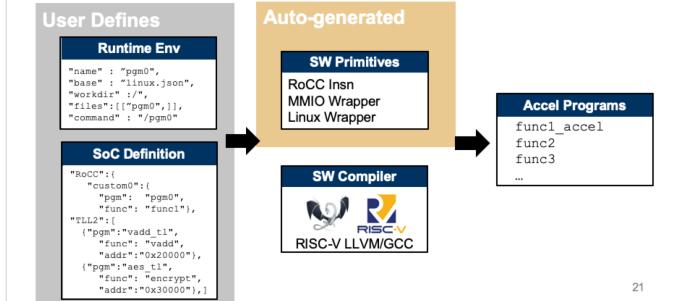
## Step 1: Hotspot Detection on FireSim



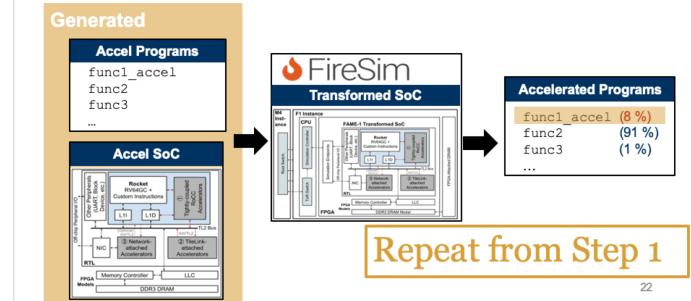
## Step 2: SoC Generation with HLS



## Step 3: SW Primitives Generation



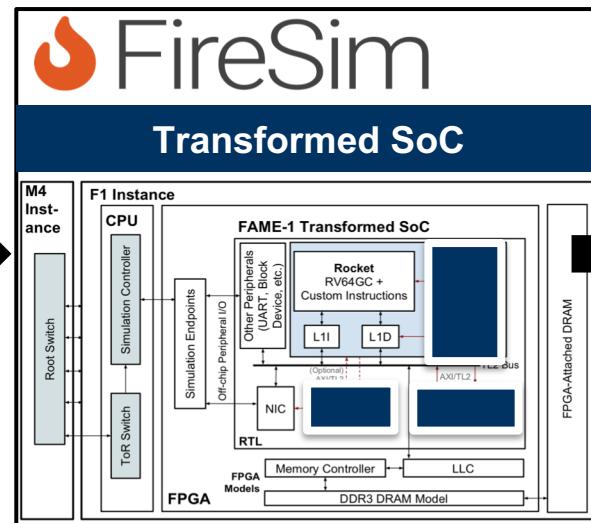
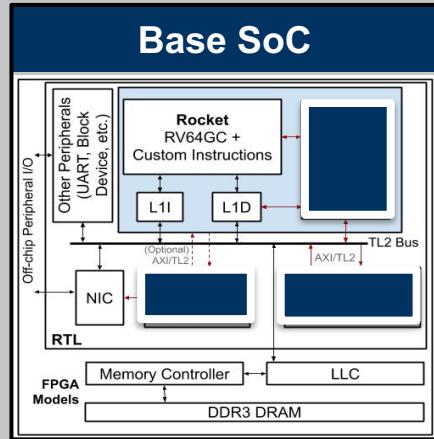
## Step 4: End-to-end Evaluation



# Step 1: Hotspot Detection on FireSim

## Inputs

Input Programs
func1
func2
func3
...



Profiled Programs	
func1	(90 %)
func2	(8 %)
func3	(1 %)
...	

# Step 2: SoC Generation with HLS

User Defines

## HLS C Func Pragma

```
void func1 (){  
    for(i=0;i<n;i++)  
#pragma pipeline II=N  
    for(j=0;j<m;j++)  
#pragma unroll M }
```

## SoC Definition

```
"RoCC": {  
    "custom0": {  
        "pgm": "pgm0",  
        "func": "func1"},  
    "TLL2": [  
        { "pgm": "vadd_t1",  
          "func": "vadd",  
          "addr": "0x20000"},  
        { "pgm": "aes_t1",  
          "func": "encrypt",  
          "addr": "0x30000"}]
```

Auto-generated

## HW Interface

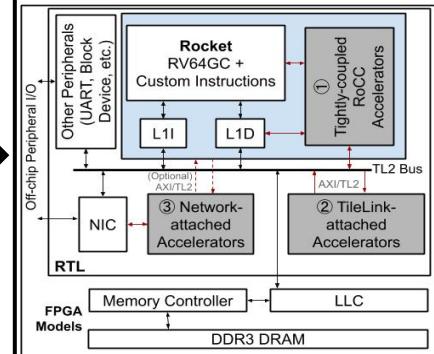
- 1.RoCC
- 2.TileLink
- 3.Network-Attached



VIVADO  
HLx Editions

Generated RTL

## Accel SoC



# Step 3: SW Primitives Generation

## User Defines

### Runtime Env

```
"name" : "pgm0",
"base" : "linux.json",
"workdir" : "/",
"files": [["pgm0", ]],
"command" : "/pgm0"
```

### SoC Definition

```
"RoCC": {
    "custom0": {
        "pgm": "pgm0",
        "func": "func1"
    },
    "TLL2": [
        {"pgm": "vadd_t1",
         "func": "vadd",
         "addr": "0x20000" },
        {"pgm": "aes_t1",
         "func": "encrypt",
         "addr": "0x30000" }
    ]
}
```

## Auto-generated

### SW Primitives

```
RoCC Insn
MMIO Wrapper
Linux Wrapper
```

### Accel Programs

```
func1_accel
func2
func3
...
```

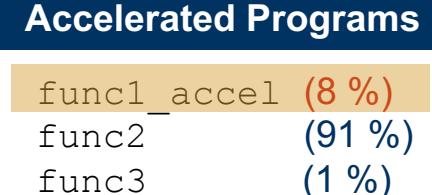
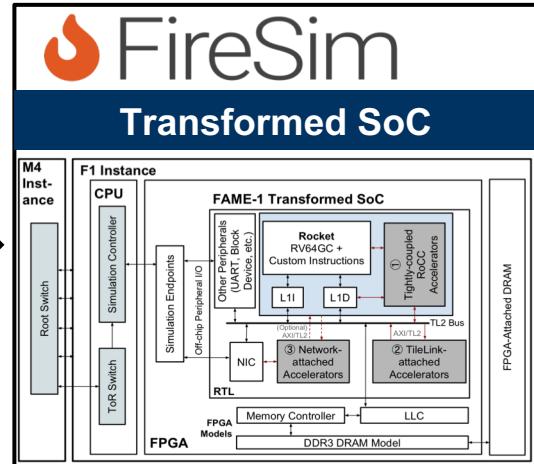
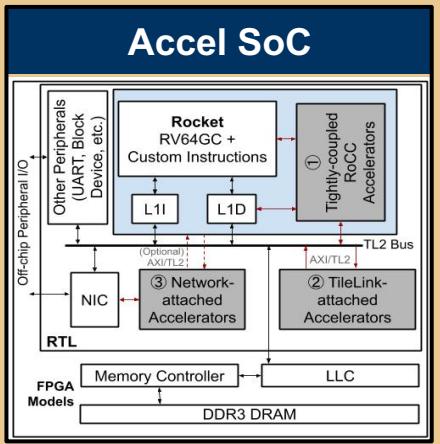
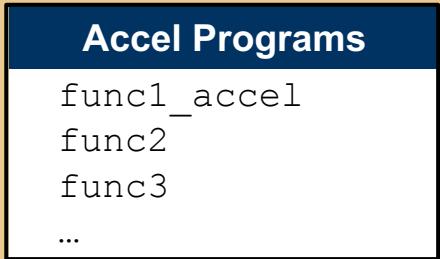
### SW Compiler



RISC-V LLVM/GCC

# Step 4: End-to-end Evaluation

Generated



Repeat from Step 1

# Exposed Design Space

1. Hardware Integration
2. Accelerator Design Variation
3. Software Integration

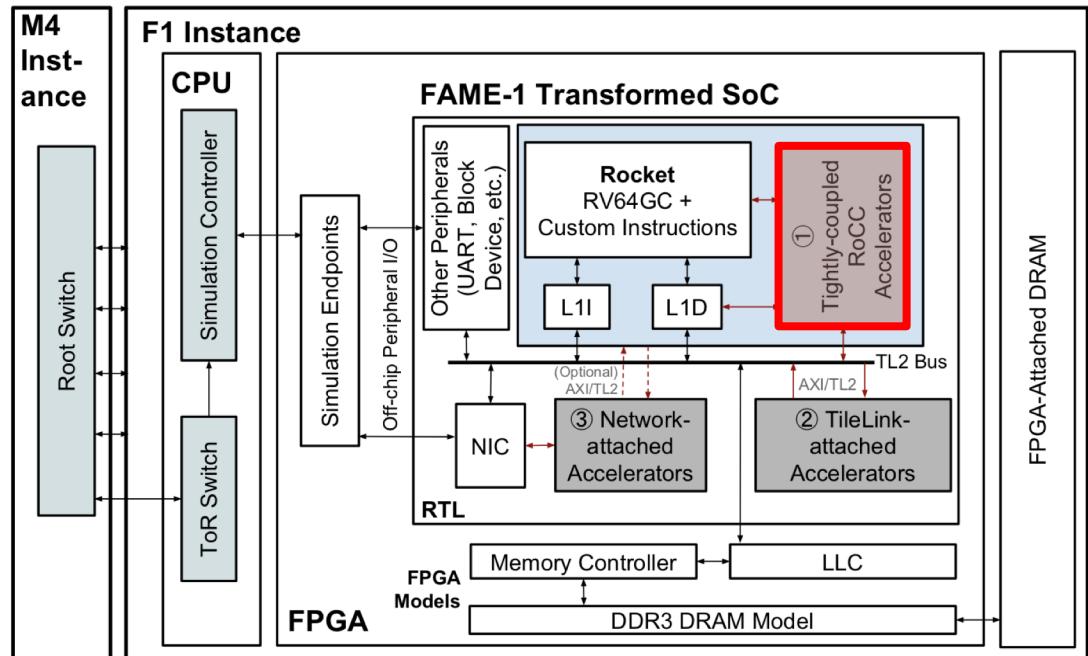
RoCC  
Coprocessors

TileLink  
Accelerators

Network-Attached  
Accelerators

# RoCC Coprocessors

- Invoked by RoCC instruction
- Sharing L1 and LLC with the CPU
- Sharing TLB with the CPU



# Exposed Design Space

1. Hardware Integration
2. Accelerator Design Variation
3. Software Integration

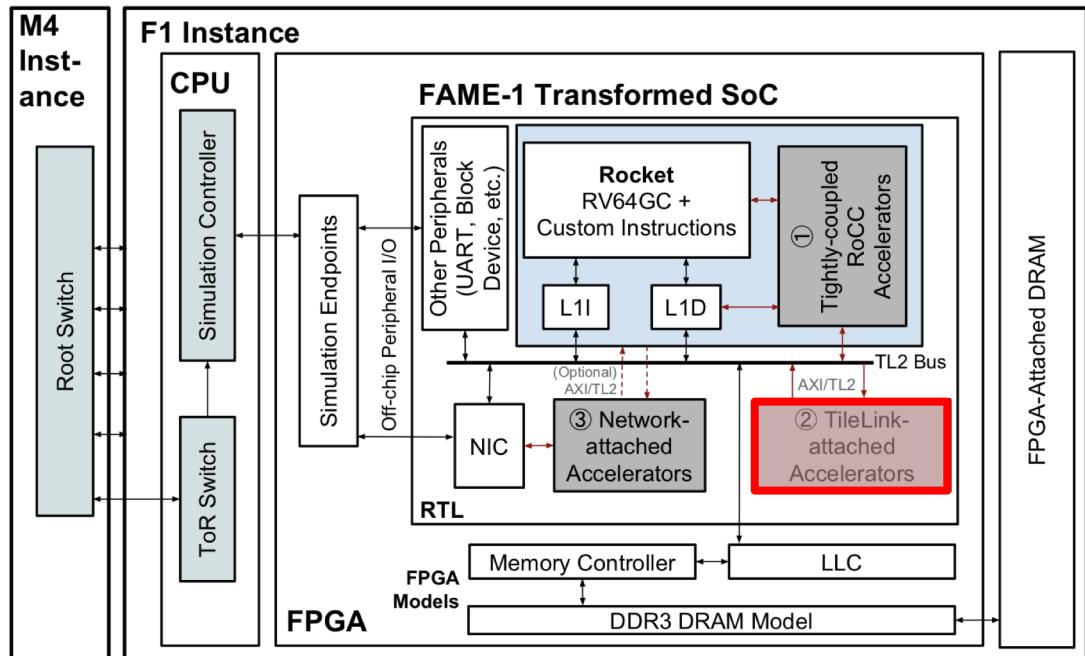
RoCC  
Coprocessors

TileLink  
Accelerators

Network-Attached  
Accelerators

# TileLink Accelerators

- Invoked by RoCC instruction or MMIO
- Sharing LLC with the CPU
- Physically-addressed



# Exposed Design Space

1. Hardware Integration
2. Accelerator Design Variation
3. Software Integration

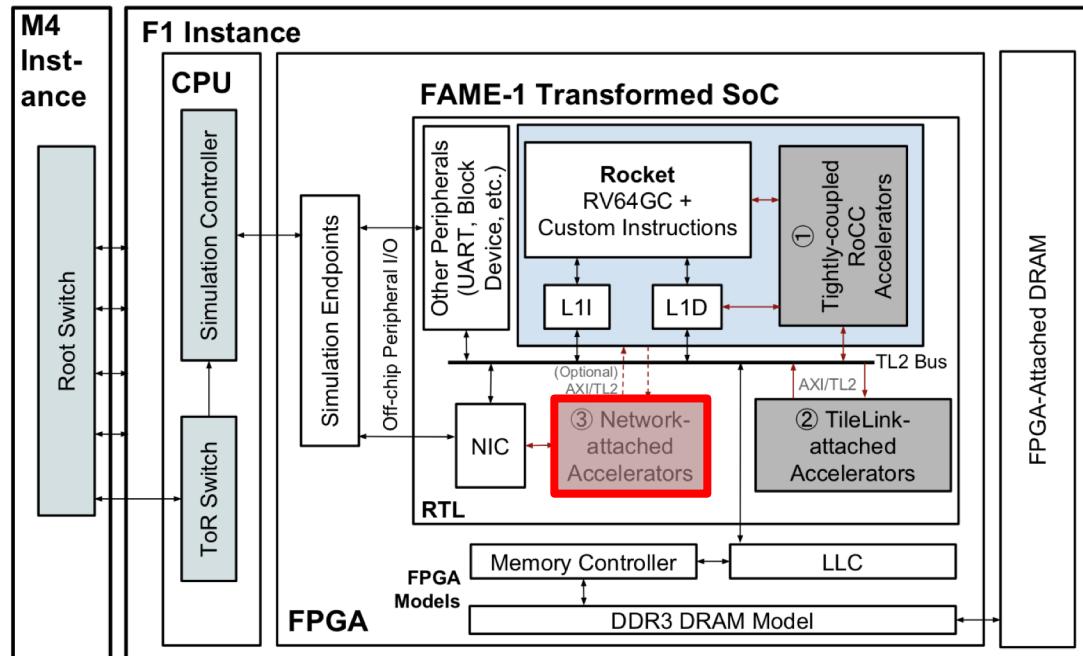
RoCC  
Coprocessors

TileLink  
Accelerators

Network-Attached  
Accelerators

# Network-attached Accelerators

- Same as TileLink Accelerators
- With Direct Access to the Ethernet



# Network-Attached Accelerator

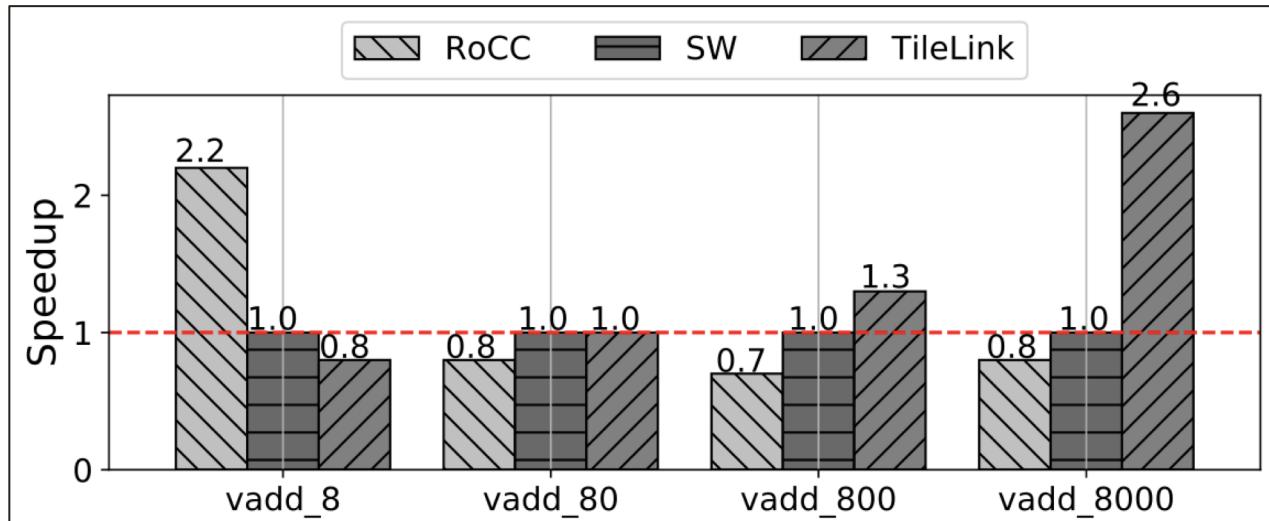
1. User includes the Centrifuge Ethernet streaming interfaces and packet parser functions in the HLS code

Centrifuge Ethernet  
Interfaces:

```
hls::stream<ap_uint<128> >& resp_head,  
hls::stream<ap_uint<65> >& resp_data,  
hls::stream<ap_uint<128> >& req_head,  
hls::stream<ap_uint<65> >& req_data,  
ap_uint<64>srcmac, ap_uint<64>dstmac
```

2. Centrifuge generates the interconnect between the accelerator and the Ethernet

# Hardware Integration DSE



Different Coupling for *vadd* Accelerator

- Three factors:
1. Cache Size
  2. Cache Latency
  3. Interface Bandwidth

Determine the best hardware integration strategy

# Exposed Design Space

1. Hardware Integration
2. Accelerator Design Variation
3. Software Integration

RoCC  
Coprocessors

TileLink  
Accelerators

Network-Attached  
Accelerators

# Exposed Design Space

- 1. Hardware Integration
- 2. Accelerator Design Variation
- 3. Software Integration
  - A. Loop Unrolling Factors
  - B. Loop Pipelining Factors
  - C. Memory Size and Parallel Ports
  - D. Resource Binding

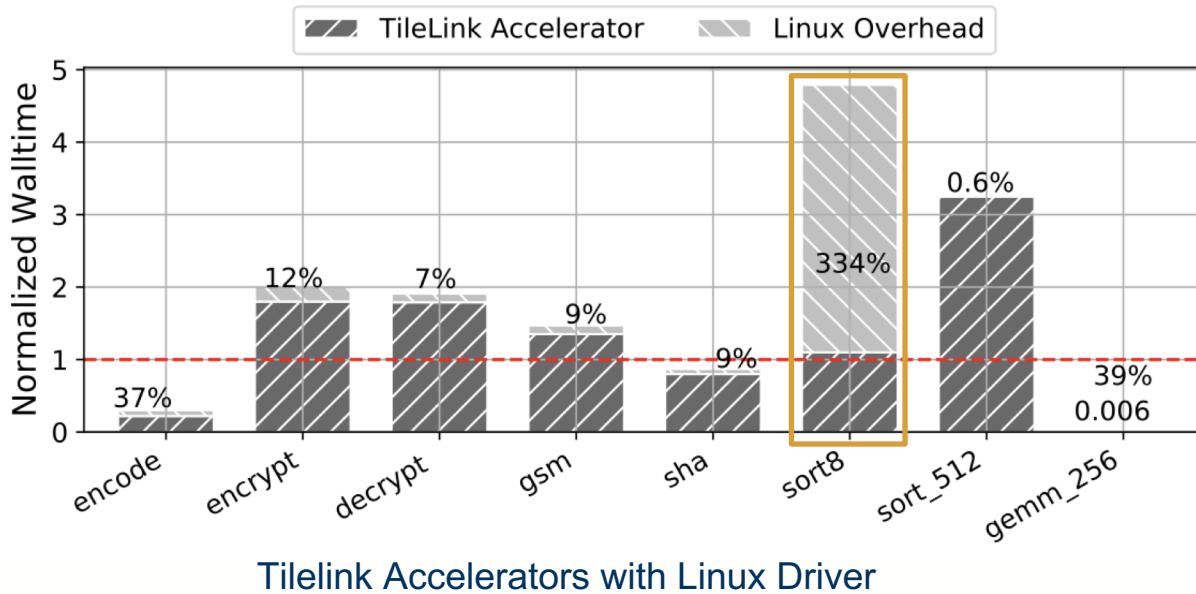
# Exposed Design Space

- 1. Hardware Integration
- 2. Accelerator Design Variation
- 3. Software Integration
  - A. Bare-metal
  - B. Linux

# Linux Wrapper Generation

- Allocates contiguous physical addresses
- Virtual to physical address lookup
- Map MMIO registers to the user space

# Software Integration DSE

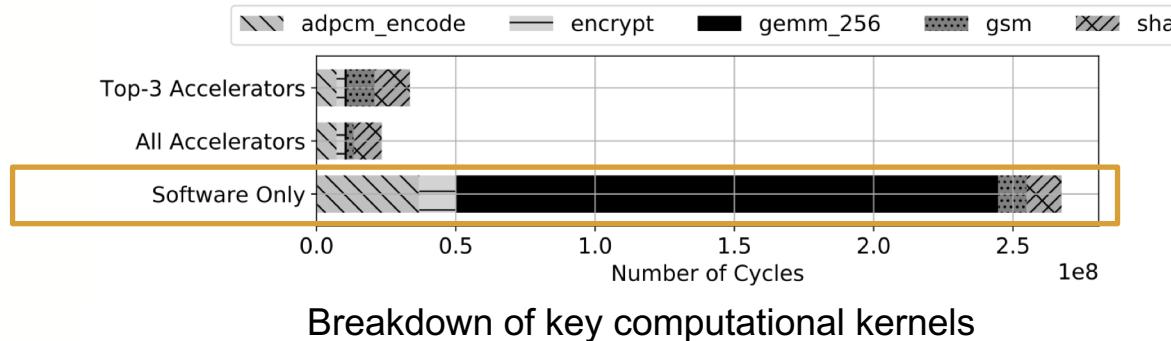


- In most cases, the Linux overhead is not significant
- RoCC-level integration might be better option for *sort8*

Determine whether the functions can be accelerated under Linux

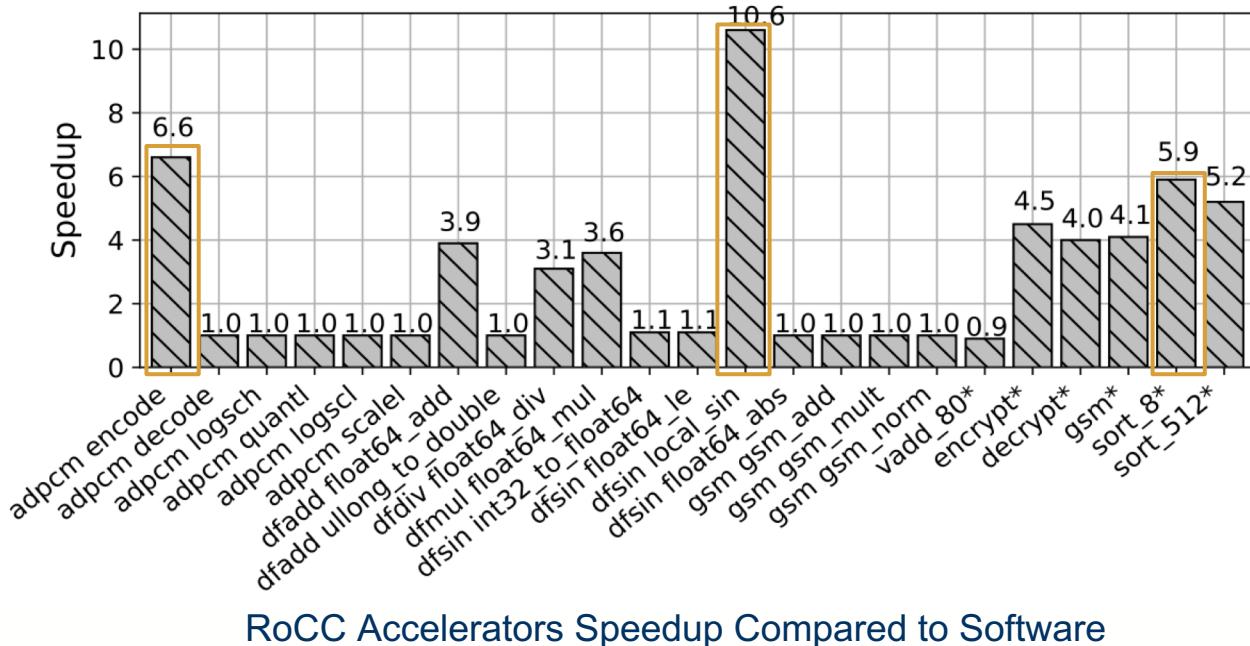
# Case Study 1: Smart-House Hub

- Application: A smart-house assistant



1. Apply Amdahl's law in hotspot detection

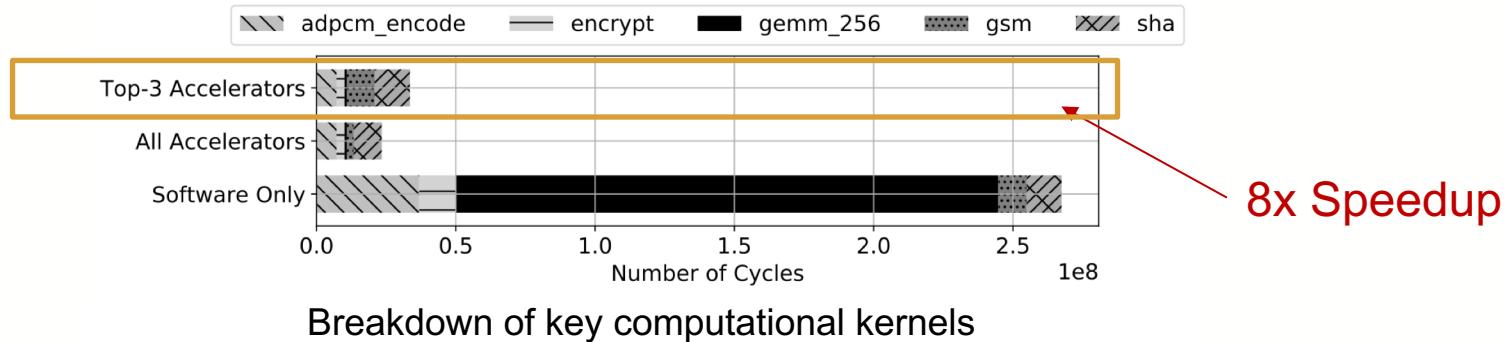
# Acceleration Region DSE



2. Determine whether the functions can be accelerated with HW

# Case Study 1: Smart-House Hub

- Application: A smart-house assistant



# Other Case Studies

## **Case Study 2: Distributed GEMM Accelerator**

1. Running Full-stack Linux with MPI
2. Distributed on 16 nodes

## **Case Study 3: Deep Learning Accelerators**

1. Different Hardware Utilization
2. Network-attached Accelerators

# Conclusion

We present a methodology and flow, *Centrifuge*, that can rapidly generate and evaluate heterogeneous SoCs by combining an HLS toolchain with the open-source FireSim FPGA-accelerated simulation platform

**Access to code:** <https://github.com/hqjenny/centrifuge>  
**Email:** [gijing.huang@berkeley.edu](mailto:gijing.huang@berkeley.edu)

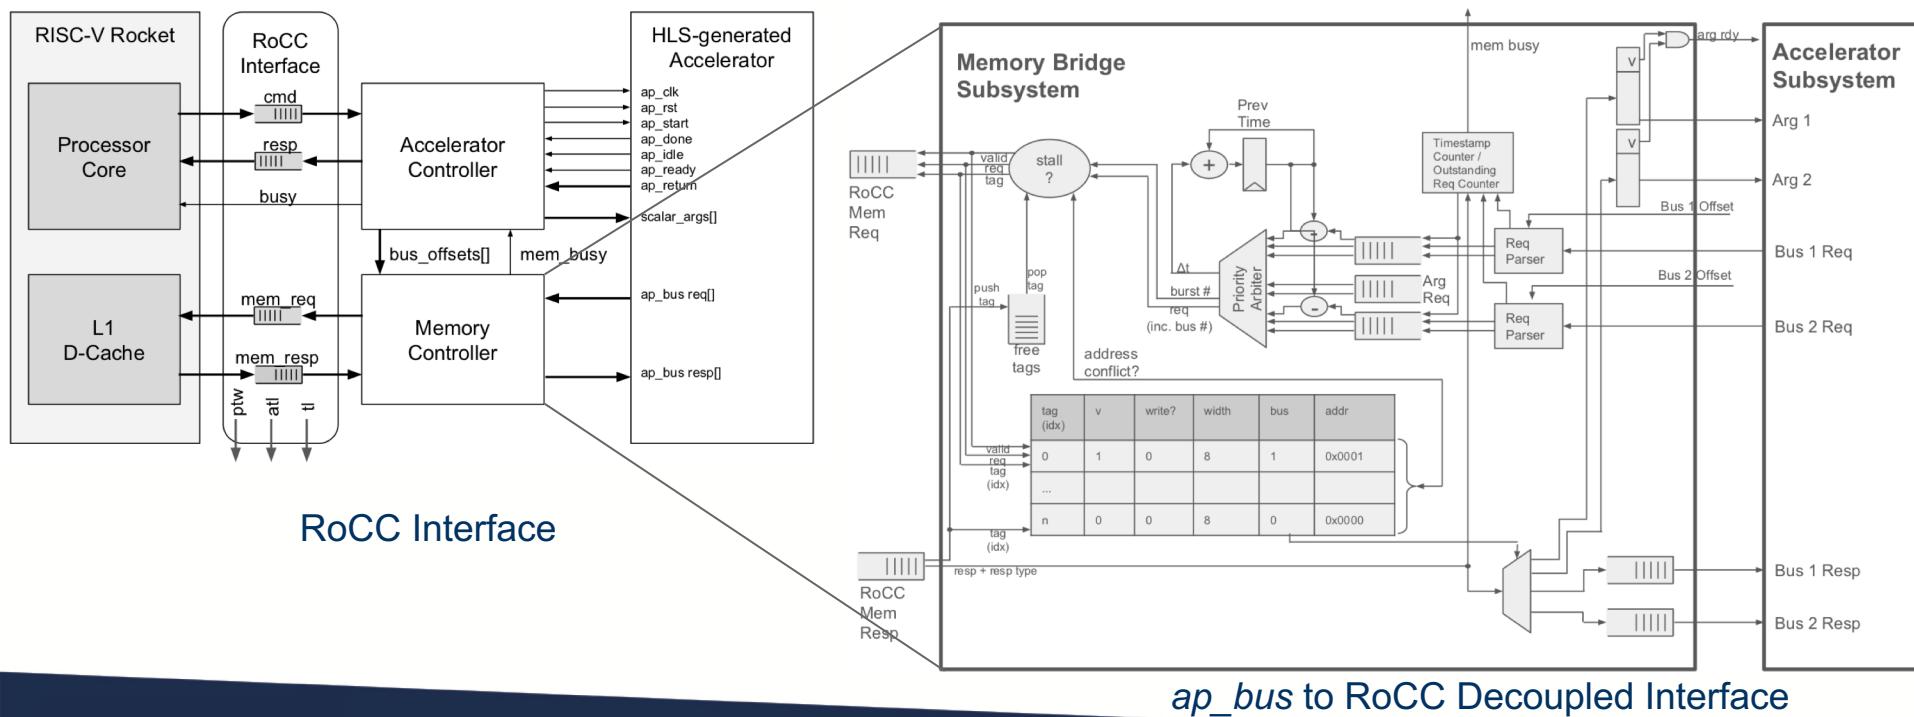
**FireSim:** <https://firesim.org>



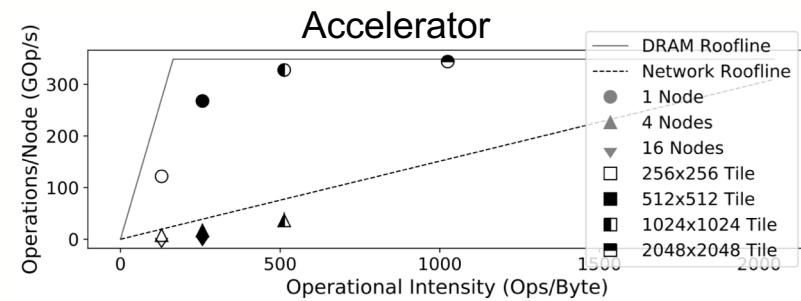
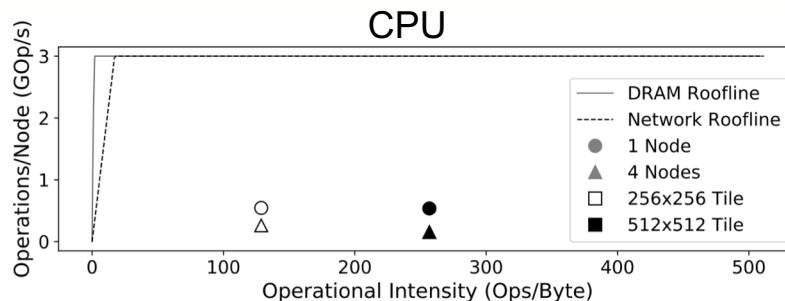


# Backup Slides

# RoCC Coprocessor Integration



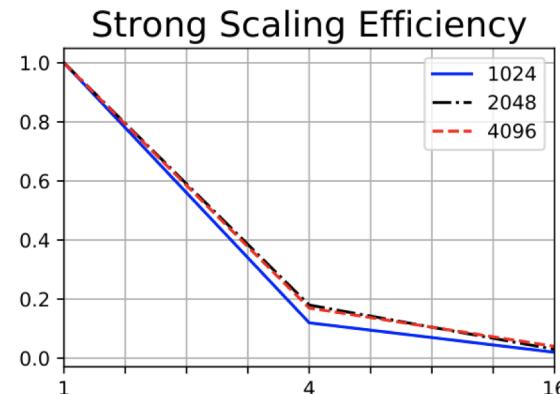
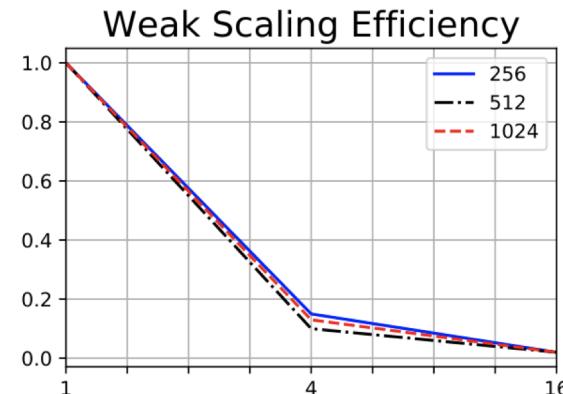
# Case Study 2: Distributed GEMM



Roofline Models

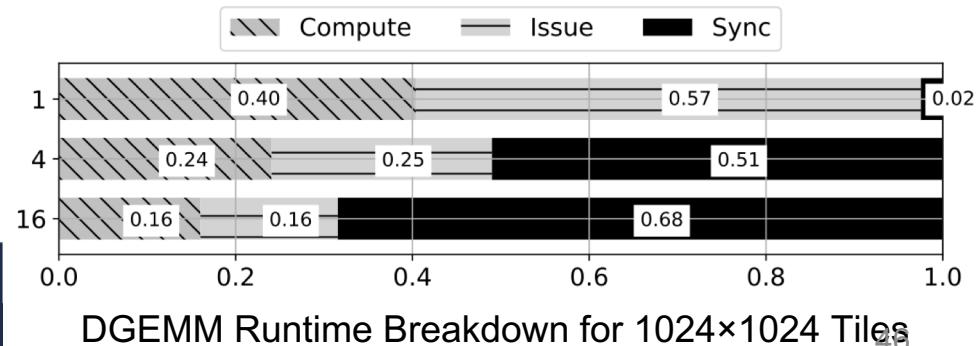
- Application: MPI-based Distributed Matrix Multiplication (DGEMM)
- Tools: STREAM and iperf are used to generate the rooflines
- Observation: With the GEMM accelerators, the distributed workload becomes more bandwidth-bound

# Case Study 2: Distributed GEMM



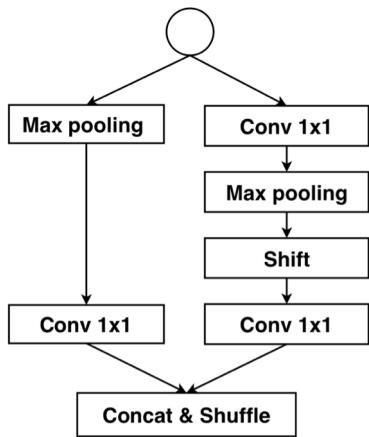
Scaling Efficiency for DGEMM with Accelerators

- Experiments on 1, 4, 16 nodes
- 2.0 GB/s measured DRAM bandwidth
- 1.2 Gbit/s measured network bandwidth

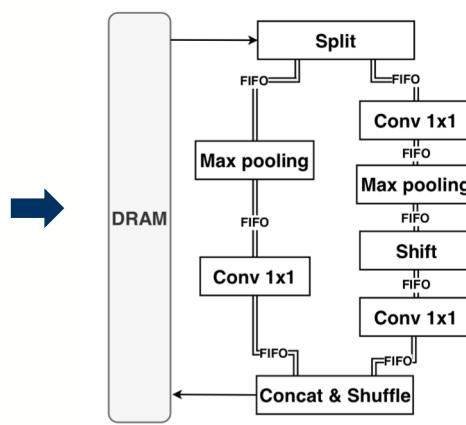


# Case Study 3: Deep Learning

## 1. HW Utilization for Different Layer Sizes



DNN Building Block



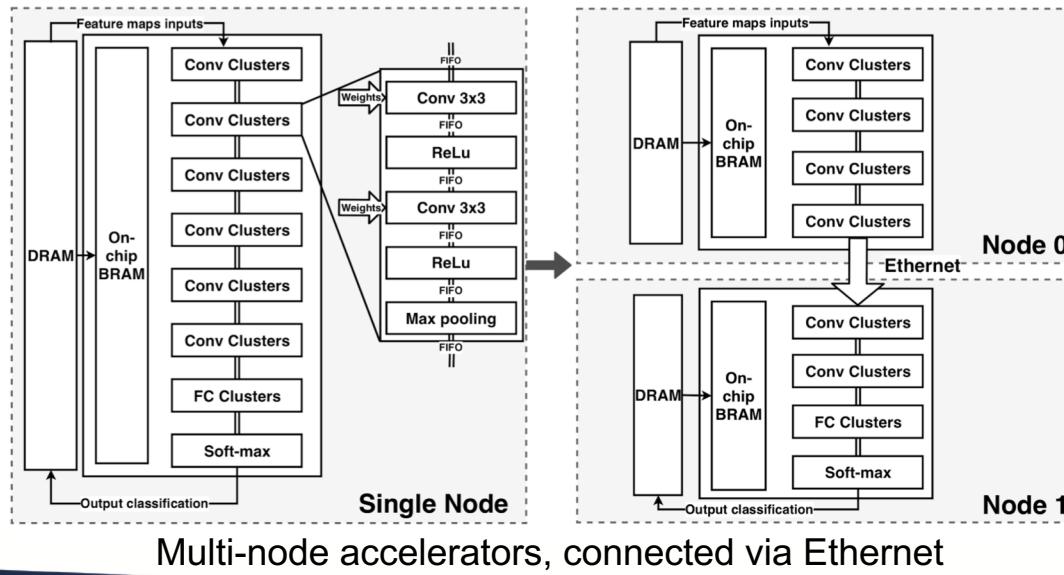
Spatial HW Design

Workload Size	Total Ops	Ops/cycle
$32 \times 16$	196608	4.55
$32 \times 32$	786432	15.12
$32 \times 64$	3145728	20.59
$16 \times 128$	3145728	21.35
$8 \times 64$	196608	17.09

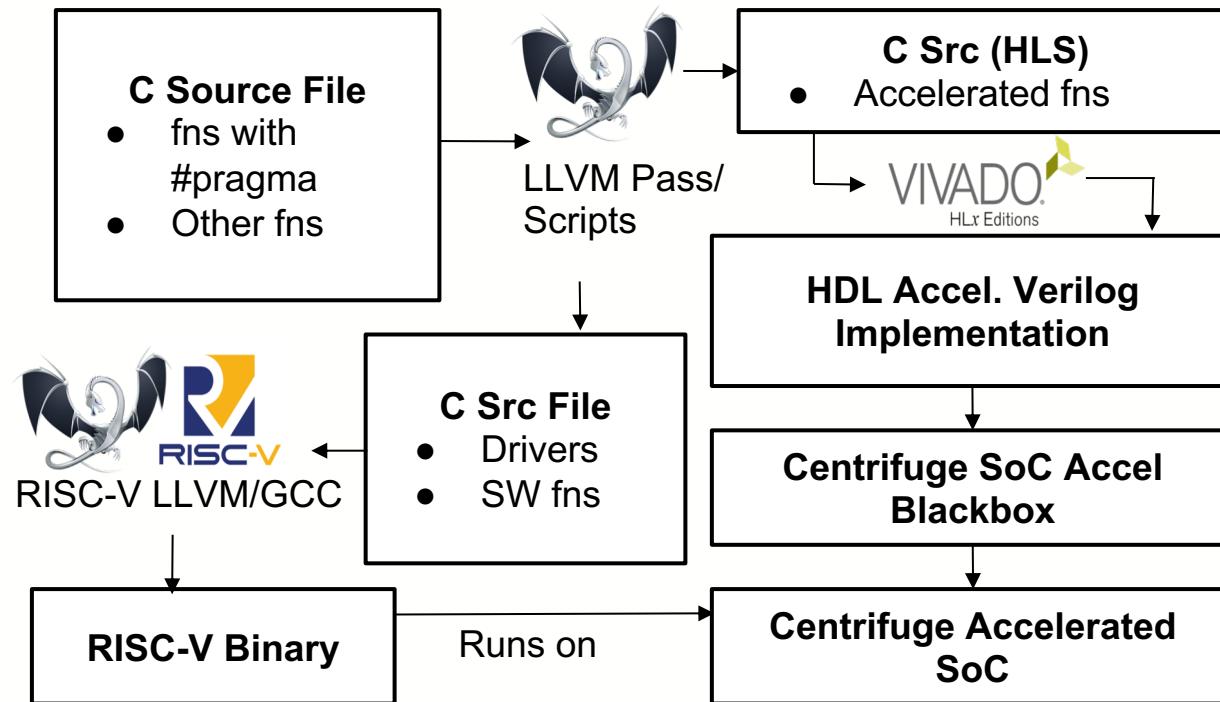
Ops/Cycle for Different Workload Sizes

# Case Study 3: Deep Learning

## 2. Distributed Accelerators



# Detailed Tool Flow



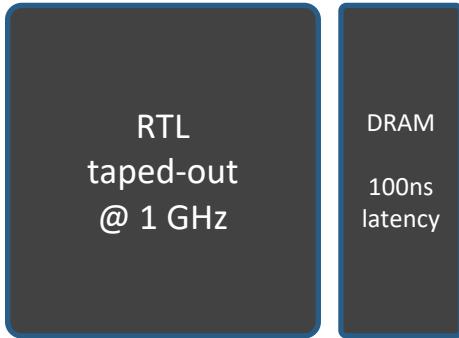
# FireSim



**FireSim** is an open-source, FPGA-accelerated, cycle-exact, scalable **hardware simulator**

- Produces:
  - FPGA-accelerated Simulation (Not FPGA prototype)
  - Deployment on Cloud FPGAs

Taped-out SoC Design



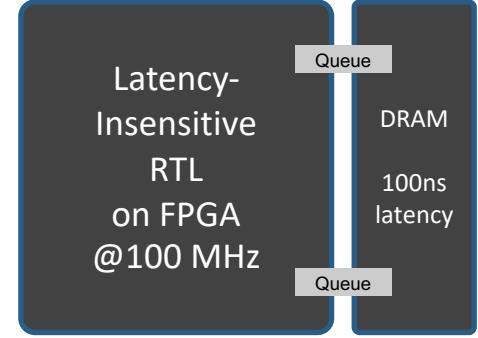
SoC sees **100** cycle DRAM latency

FPGA Prototyping



SoC sees **10** cycle DRAM latency

FPGA-accelerated Simulation

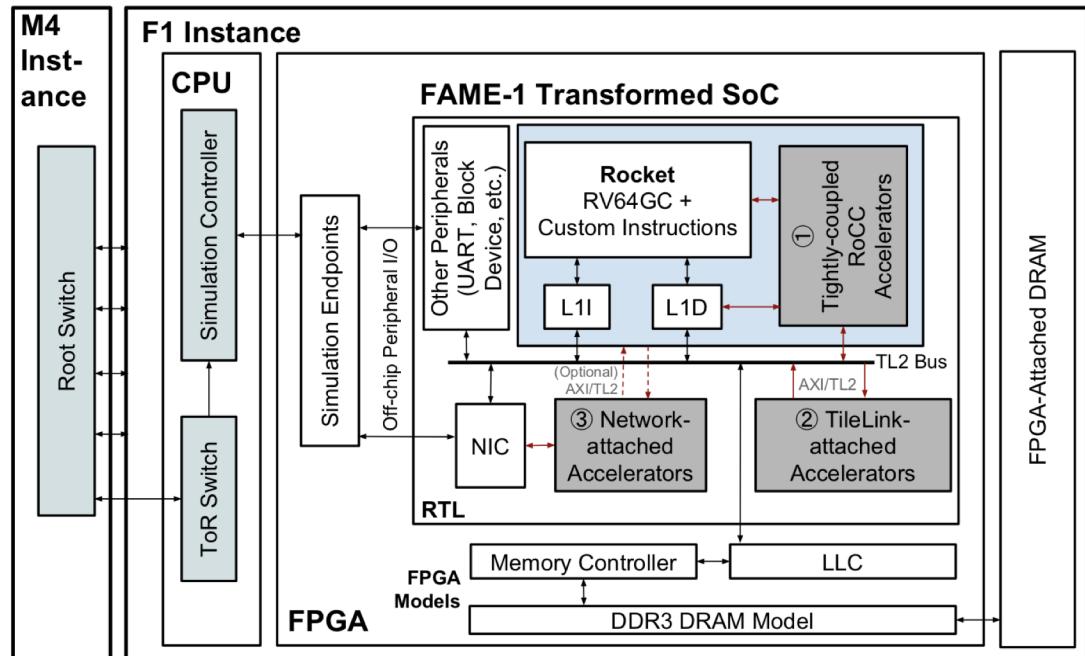


SoC sees **100** cycle DRAM latency

50

# Accelerator Coupling

1. RoCC  
Coprocessors
2. TileLink  
Accelerators
3. Network-attached  
Accelerators



# New Applications to Accelerate

- Packet Processing
- Video/Audio Compression/Decompression
- DNN Acceleration
- Graph Acceleration
- Database Systems

# Centrifuge Tool Flow

