

# AutoPhase: Juggling HLS Phase Orderings in Random Forests with Deep Reinforcement Learning

Ameer Haj-Ali\*, Qijing Huang\*,  
W. Moses, J. Xiang, K. Asanovic, J. Wawrzynek, I. Stoica

\*Equal Contribution

<https://github.com/ucb-bar/autophase>

*Toward automatic and intelligent optimization of off-the-shelf code*



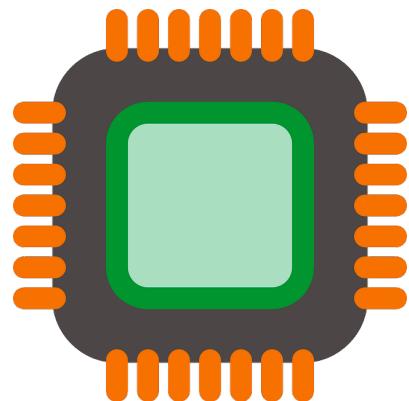
Berkeley  
Architecture  
Research



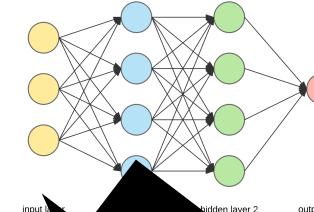
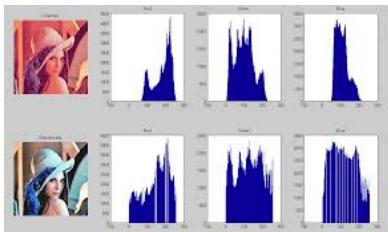
```
#include <stdio.h>
int main (void) {
    printf ("Hello, World!\n");
    return 0;
}
```



01001001  
11011101  
01001111  
01001001  
11011101



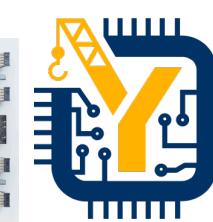
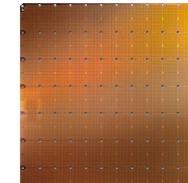
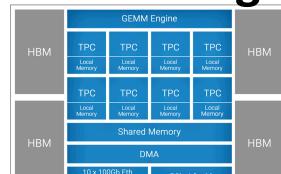
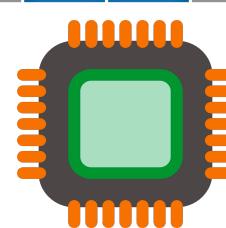
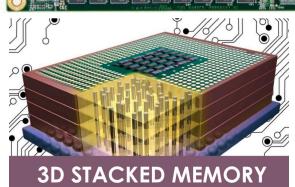
For years, software engineers used **heuristics** and **hand engineering** to optimize compilation!



This is  
confusing

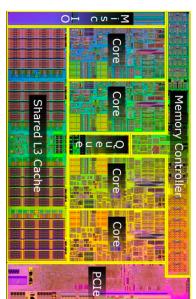
...

Application-Specific  
Integrated Circuits



Distributed Computing

Plethora of new applications  
Need a way to compile  
applications to rapidly  
changing hardware and  
software without  
compromising  
performance!



# Compiler Challenge Example: Phase Ordering (NP-Hard)

Which passes?  
In which order?  
More than  $2^{247}$   
possibilities ...

`clang program.c -flag1 -flag2 ...`

# Example – Normalizing a Vector

```
for (int i=0; i<n; i++) {  
    out[i] = in[i] / mag(in, n);  
}
```

Pass #1: -licm (loop-invariant code motion)

```
double precompute = mag(in, n);  
for (int i=0; i<n; i++) {  
    out[i] = in[i] / precompute;  
}
```

Pass #2: -inline

```
double precompute, sum = 0;  
for (int i=0; i<n; i++) {  
    sum += A[i] * A[i];  
}  
precompute = sqrt(sum);  
for (int i=0; i<n; i++) {  
    out[i] = in[i] / precompute;  
}
```

$\Theta(n)$

# Example – Normalizing a Vector

```
for (int i=0; i<n; i++) {  
    out[i] = in[i] / mag(in, n);  
}
```

Pass #1: -inline

```
for (int i=0; i<n; i++) {  
    double sum = 0;  
    for (int j=0; j<n; j++) {  
        sum += A[j] * A[j];  
    }  
    out[i] = in[i] / sqrt(sum);
```

~~$\Theta(n)$~~



$\Theta(n^2)$

10/12/20

```
double sum;  
for (int i=0; i<n; i++) {  
    sum = 0;  
    for (int j=0; j<n; j++) {  
        sum += A[j] * A[j];  
    }  
    out[i] = in[i] / sqrt(sum);
```



# Related Works to Remedy

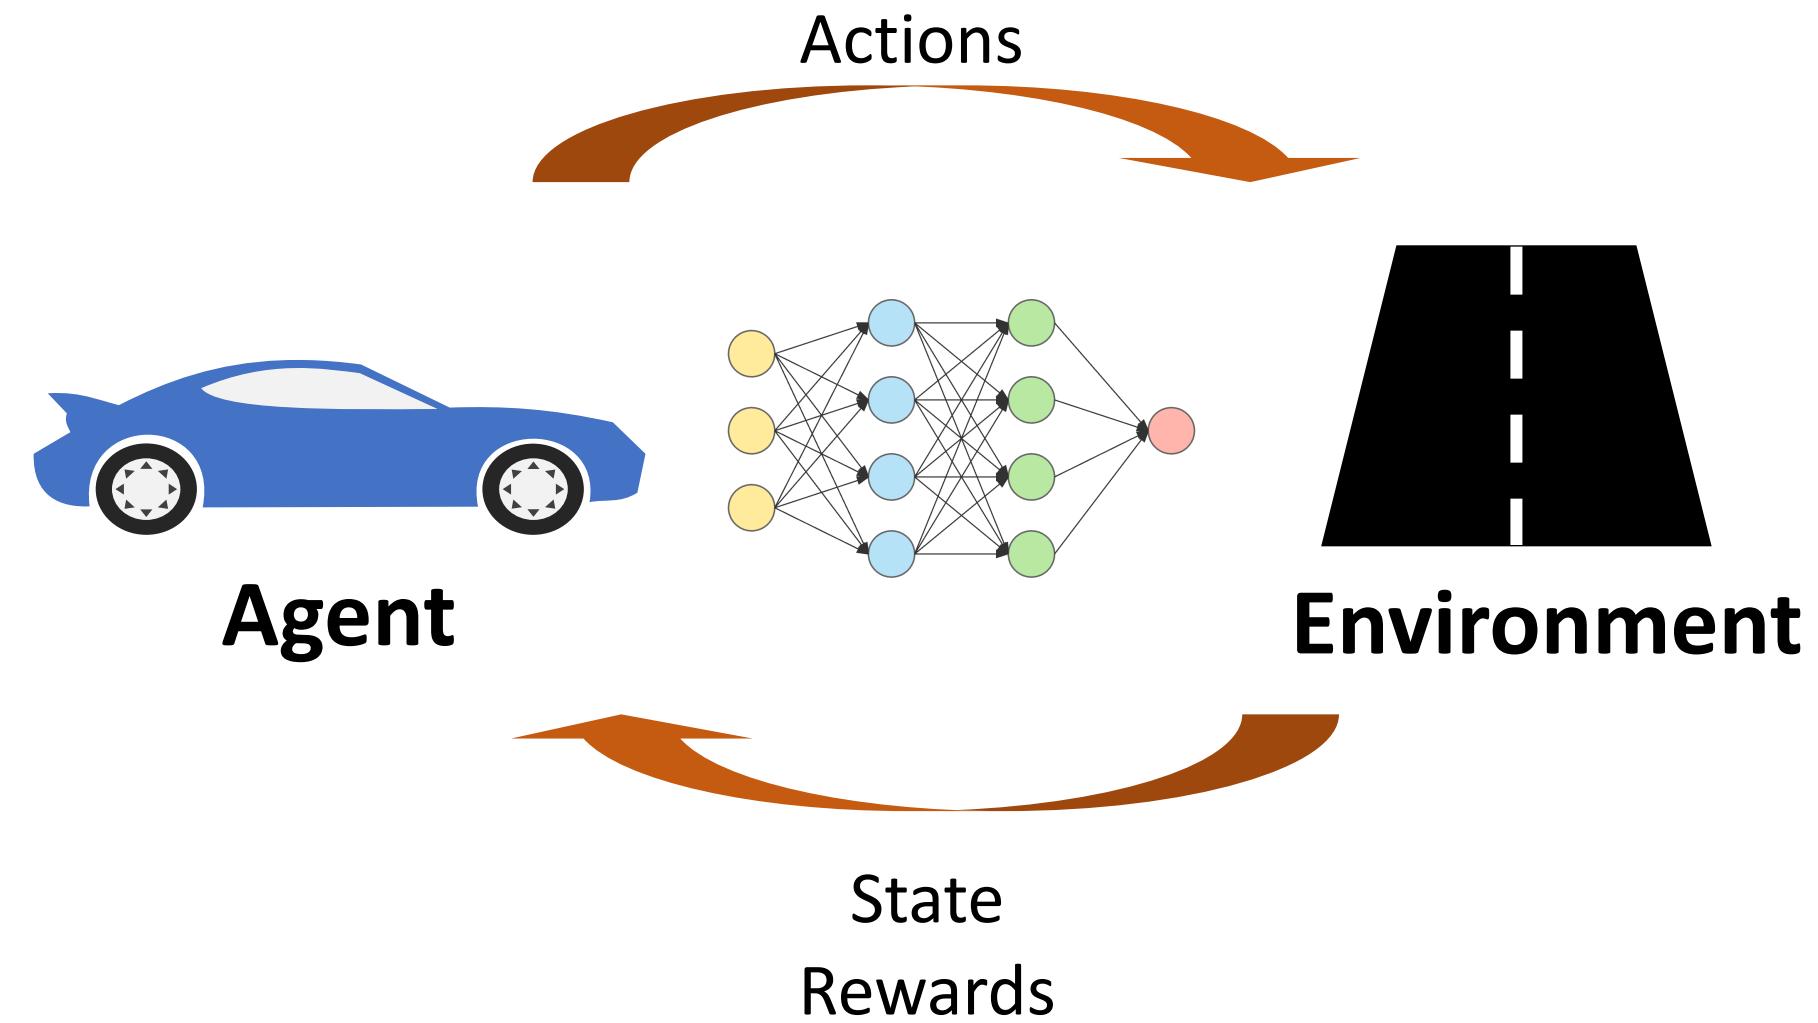
- Heuristics, Greedy algorithms, Genetic algorithms, Multi-armed Bandit
    - Hand engineered features
    - Error prone
    - Requires retraining for every program
    - Nothing is being “learned”
    - Cannot generalize
  - Supervised Learning
    - Requires hard/impossible-to-find labels
- 
- C. Cummins *et al.*, “End-to-End Deep Learning of Optimization Heuristics,” PACT 2017.
  - J. Ansel *et al.*, "OpenTuner: An Extensible Framework for Program Autotuning," PACT 2014.
  - Z. Wang *et al.*, "Machine Learning in Compiler Optimization." Proceedings of the IEEE 2018
  - **Q. Huang *et al.***, “The effect of compiler optimizations on high-level synthesis-generated hardware,” Trets 2015.
  - G. Fursin *et al.*, "MILEPOST GCC: machine learning based research compiler." 2008.
  - M. Stephenson et al. "Meta optimization: Improving compiler heuristics with machine learning," ACM sigplan notices 2013
  - F. Agakov et al. "Using machine learning to focus iterative optimization," CGO 2006.
  - S. Kulkarni, "Mitigating the compiler optimization phase-ordering problem using machine learning," OOPSLA 2012.

# The Ultimate Solution

- Handles complex search spaces and decisions
  - Automatic
  - Minimal supervision
  - Does not require retraining (generalizes)
  - Feedback/experience driven
  - Does not require labels
- Deep Reinforcement Learning



# Deep Reinforcement Learning



# Example

Decision(**Action**): Gas



Decision(**Action**): Brake

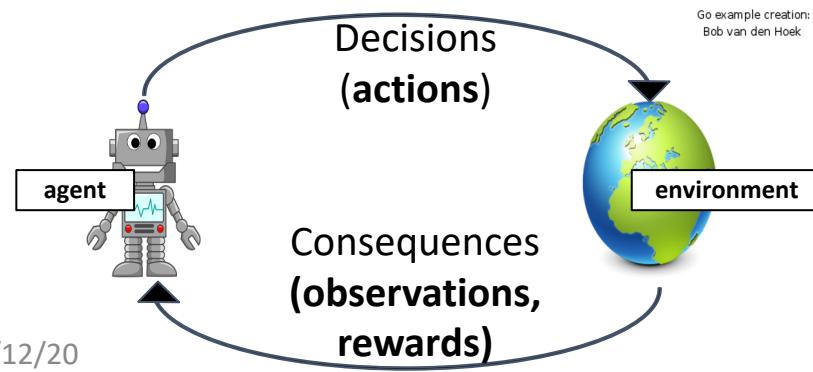
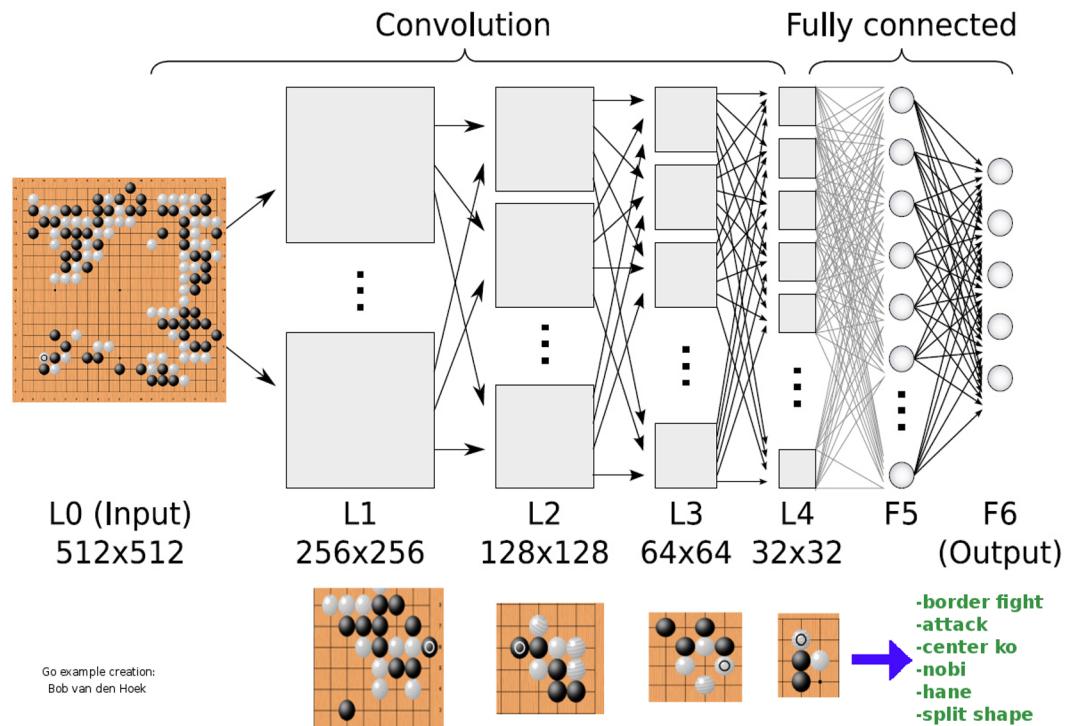


**Reward**

# Go as a Deep Reinforcement Learning Problem

AlphaGo (Silver et al. 2016)

- **Observations:**
  - board state
- **Actions:**
  - where to place the stones
- **Rewards:**
  - 1 if win
  - 0 otherwise



# Growing Number of RL Applications

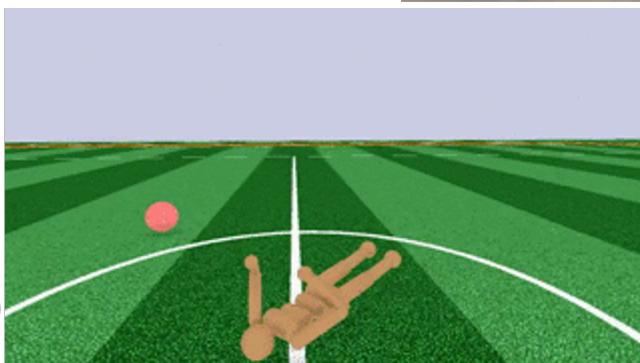
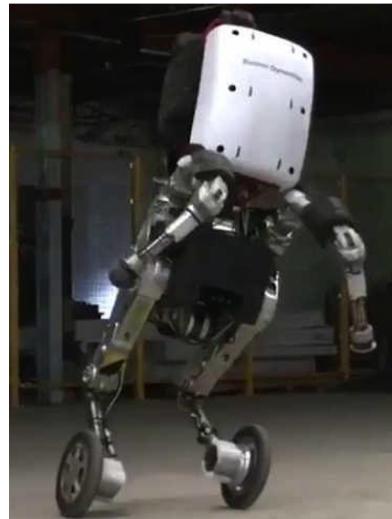
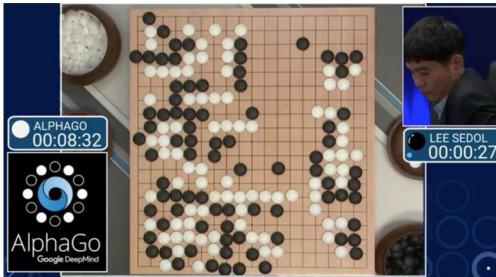
Robotics

Industrial  
Control

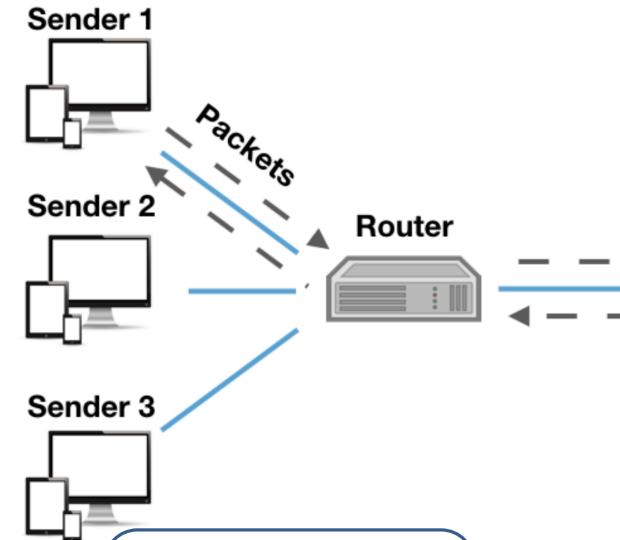
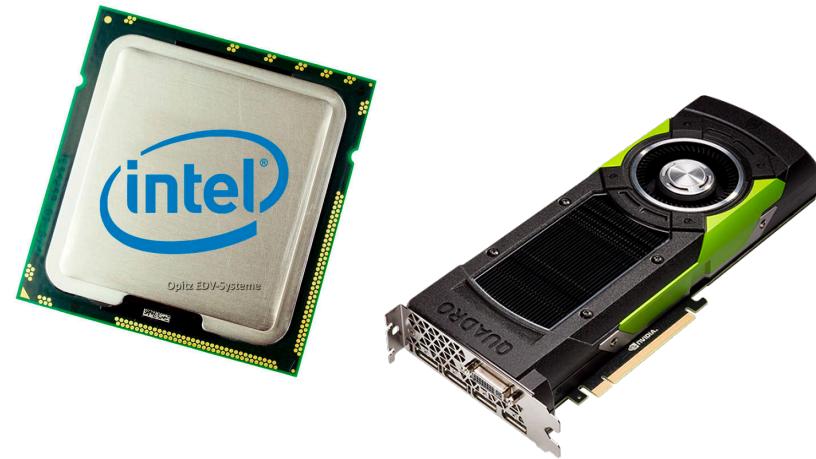
Advertising

System  
Optimization

Finance

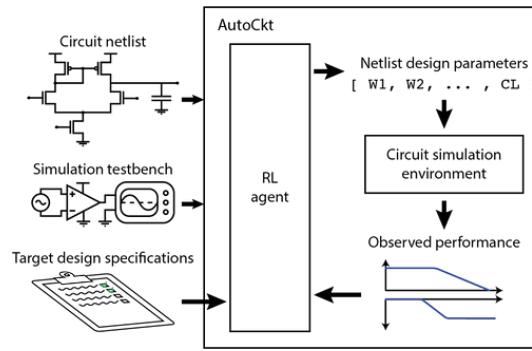


# Deep RL is Used Everywhere in Systems



Which passes?  
In which order?

gcc program.c -flag1 -flag2 ...



- Ameer Haj-Ali, et al., "A View on Deep Reinforcement Learning in System Optimization", (submitted). Available on arXiv.
- Qijing Huang\*, Ameer Haj-Ali\*, et al., "AutoPhase: Compiler Phase-Ordering for High-Level Synthesis with Deep RL," FCCM 2019.
- Keertana Settaluri, Ameer Haj-Ali, Qijing Huang, et al., "AutoCkt: Deep Reinforcement Learning of Analog Circuit Designs," DATE 2020.
- Bin Li, ..., Ameer Haj-Ali, et al., "RLDRM: Closed Loop Dynamic Cache Allocation with Deep RL for Network Function Virtualization", NetSoft 20'.
- Ameer Haj-Ali, et al., "NeuroVectorizer: End-to-End Vectorization with Deep Reinforcement Learning," CGO 2020, to appear.
- Jay, Nathan, et al. "A Deep Reinforcement Learning Perspective on Internet Congestion Control," ICML 2019.
- Mao, Hongzi, et al. "Resource management with deep reinforcement learning." HotNets 2016.
- Krishnan, Sanjay, et al. "Learning to optimize join queries with deep reinforcement learning." arXiv preprint 2018.
- Liang, Eric, et al., "Neural Packet Classification." arXiv preprint 2019.
- Xu, Cheng-Zhong, et al., "URL: A unified reinforcement learning approach for autonomic cloud management.", JPDC 2012.

# AutoPhase: Juggling HLS Phase Orderings in Random Forests with Deep Reinforcement Learning

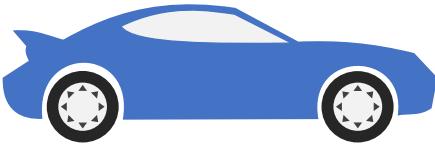
- A framework that uses Deep RL to address the phase-ordering challenge
- An importance analysis on the features
- Can generalize to different workloads
- Achieves 28% performance speedup over -O3



# AutoPhase: Phase Ordering with Deep RL

Decision (**Action**):

- Gas
- Brake
- Steering



Fee (**Reward**): \$100

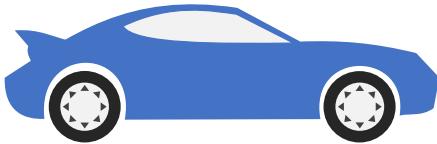
GPS Location (**State**):

37.87631055, -122.2388

# AutoPhase: Phase Ordering with Deep RL

Decision (**Action**):

- Gas
- Brake
- Steering



Fee (**Reward**): \$100

GPS Location (**State**):  
37.87631055, -122.2388



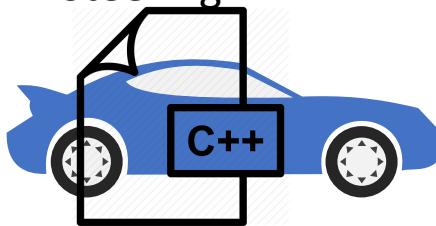
Fee (**Reward**): -\$500

GPS Location (**State**):  
-30.43131384, 73.1693

# AutoPhase: Phase Ordering with Deep RL

Decision (**Action**):

**Action:** Optimization Pass #  
- Gas (e.g. -licm, -inline, -sroa, mem2reg)  
- Brake  
- Steering



Source Code

Fee (**Reward**): \$100



GPS Location (**State**):  
37.87631055, -122.2388

**Reward:**

Cycle Count Improvement

**State:** a) Program Features  
b) Applied Passes

Fee (**Reward**): -\$500



GPS Location (**State**):  
-30.43131384, 73.1693

# Generating the States

## 1. Program Features:

- 56 static features (# of instructions, # of loops, ...)
- An LLVM analysis pass is built

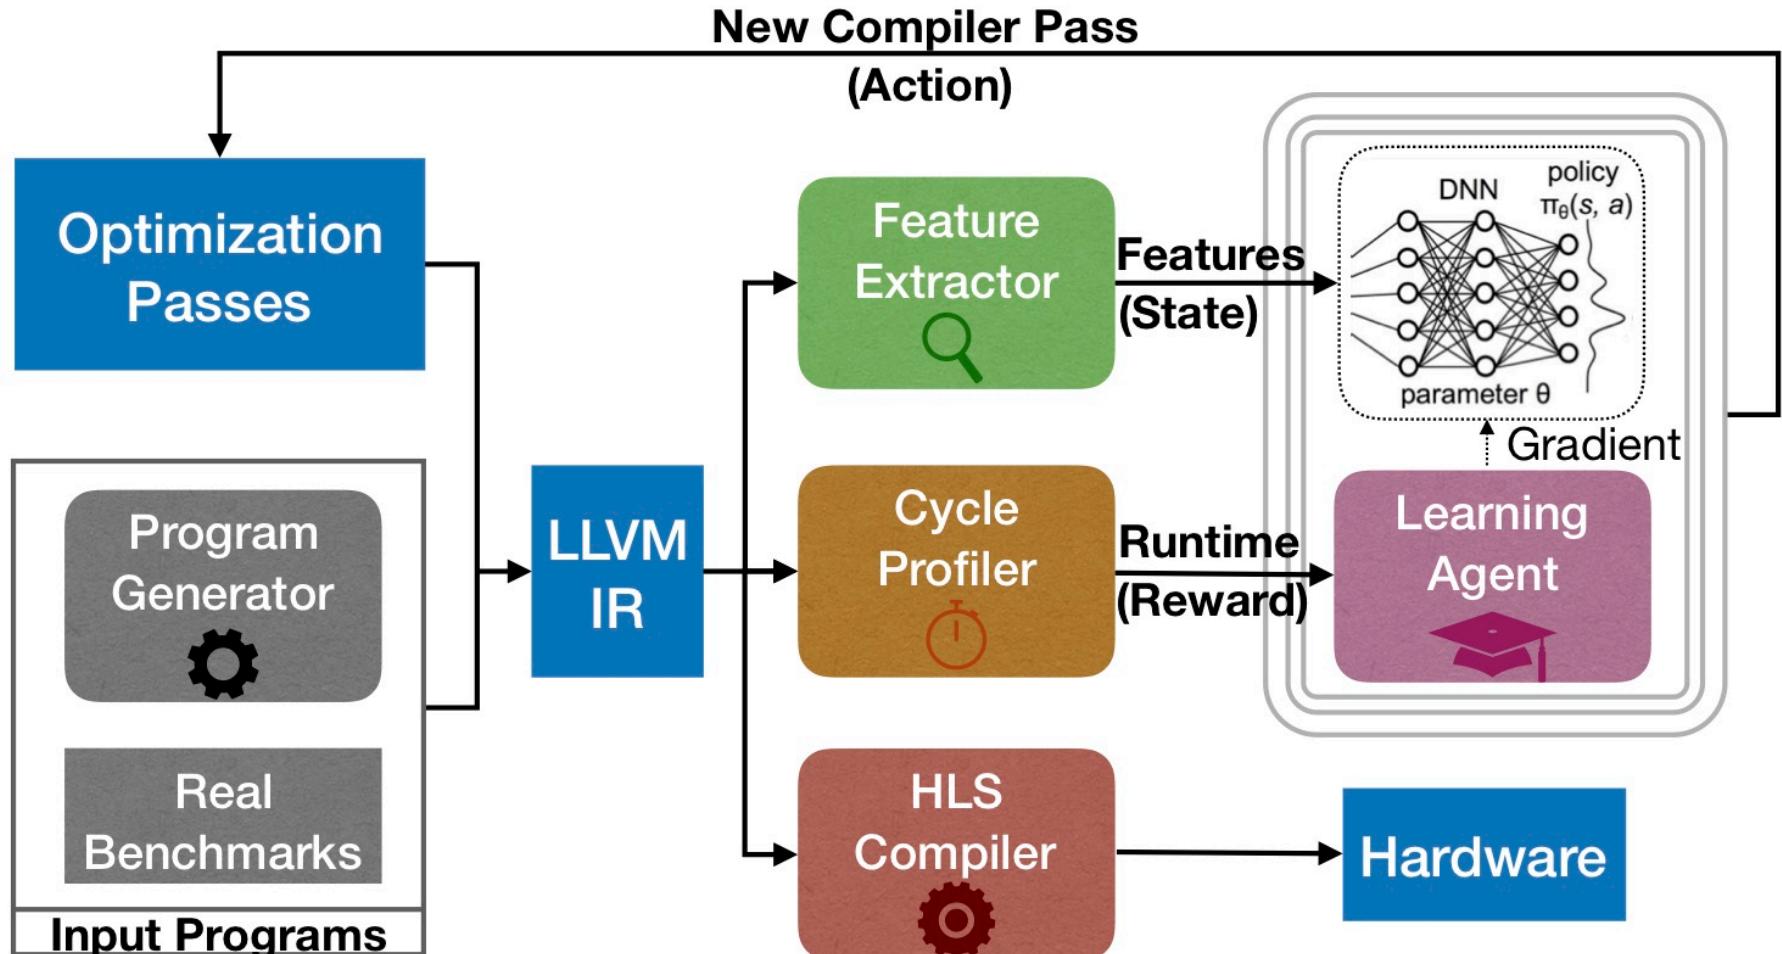
## 2. Histogram of Applied Passes:

- 45 optimization passes
- Once pass  $i$  is applied,  $Histogram[i]++$   
(**Histogram**: a vector of length 45,  $i$ : index of a pass)

# Generating the Rewards

- Reward definition:
  - cycles before a pass is applied**
    - **cycles after a pass is applied**
- LegUp estimates the circuit cycles from C
  - Leveraging SW run information
  - 20× faster than RTL simulation
  - 0.48% error rate on our benchmarks

# AutoPhase Framework



# AutoPhase Example

**Action:** -licm

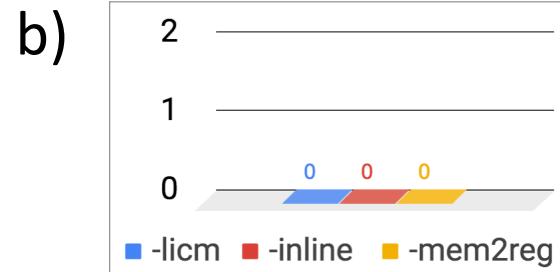
**Before:** 100 Cycles

**After:** 100 Cycles

**Rewards:**  $100 - 100 = 0$

**State:**

a) # loops = 3, # insn = 20



Trajectory Length = 3

Pass Sequence:

-licm



# AutoPhase Example

Action: -mem2reg

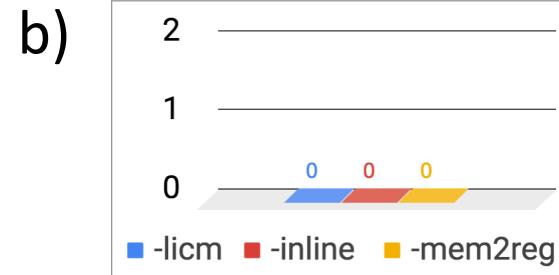
Before: 100 Cycles

After: 90 Cycles

Rewards:  $100 - 90 = 10$

State:

a) # loops = 3, # insn = 20



Trajectory Length = 3

Pass Sequence:

-licm



-mem2reg



# AutoPhase Example

Action: -inline

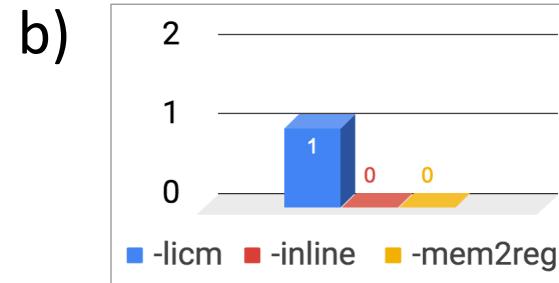
Before: 99 Cycles

After: 66 Cycles

Rewards:  $99 - 66 = 33$

State:

a) # loops = 3, # insn = 20



Trajectory Length = 3

Pass Sequence:

-licm



-mem2reg



-inline



# AutoPhase Example

Action:

Before: 88 Cycles

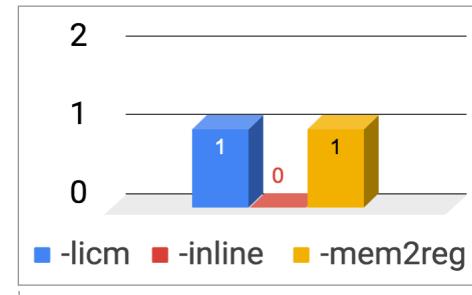
After: 56 Cycles

Rewards:  $88 - 56 = 32$

State:

a) # loops = 3, # insn = 30

b)



Trajectory Length = 3

Pass Sequence:

-licm



-mem2reg



-inline



# Evaluated Algorithms

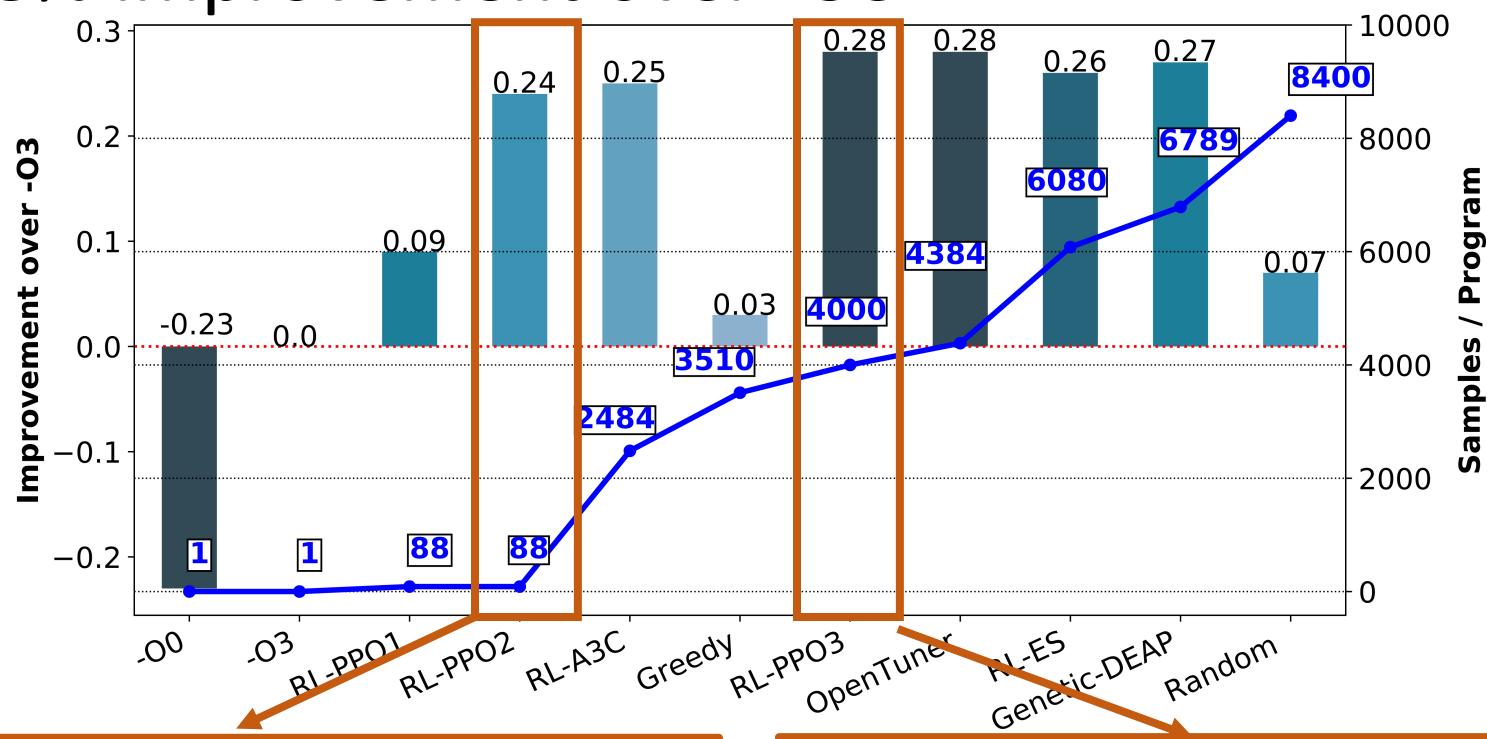
- Random Search
- Genetic Algorithms
- OpenTuner – Multiarmed Bandit/PSO
- Greedy Algorithms/Beam Search
- Deep Reinforcement Learning
  - Proximal Policy Optimization (PPO)
  - Asynchronous Actor-Critic Agents (A3C)
  - Evolutionary Strategies (ES)

# Methodology

- 9 benchmarks from CHStone
- A3C, PPO, ES vs. OpenTuner, Genetic, Greedy, Random, -O3
- The network:
  - $256 \times 256$  FC
  - Trajectory length: 45
- RLLib/Ray, 4-core Intel i7-4765T CPU with a Tesla K20c GPU for training and inference

# Performance (Execution Time)

- Fewer samples required
- 28% improvement over -O3



- Uses only **Histogram of Applied Passes**
- Compiles and runs **once** per trajectory

- Uses **Program Features** and **Histogram of Applied Passes**

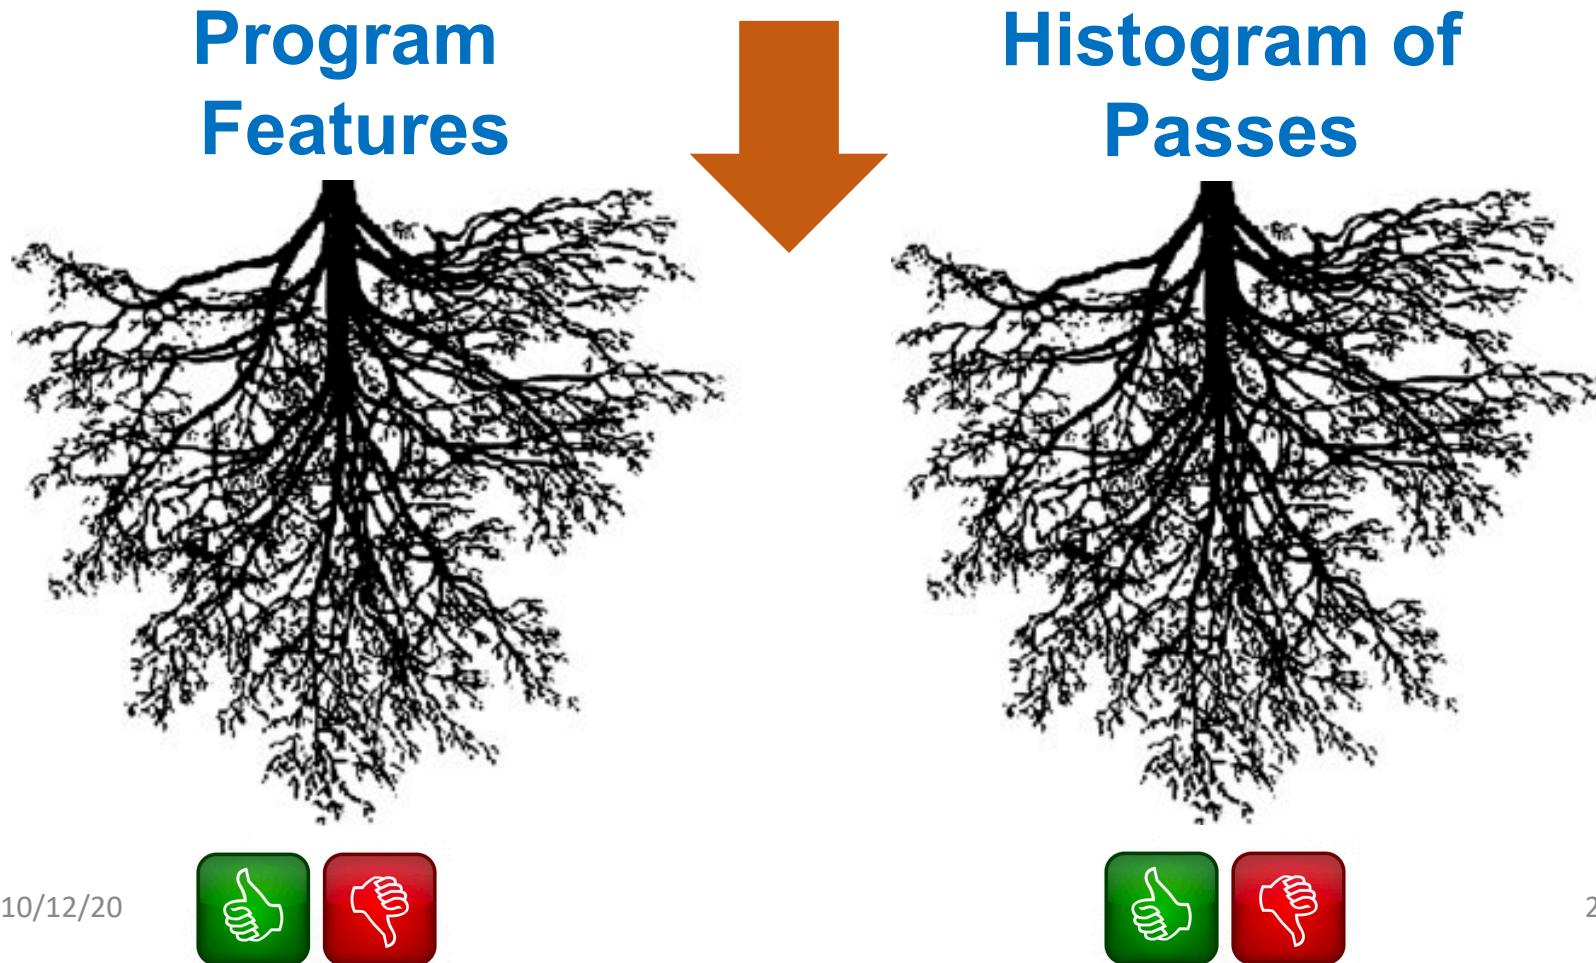
\* Results are normalized to **-O3**

# Can We Generalize?

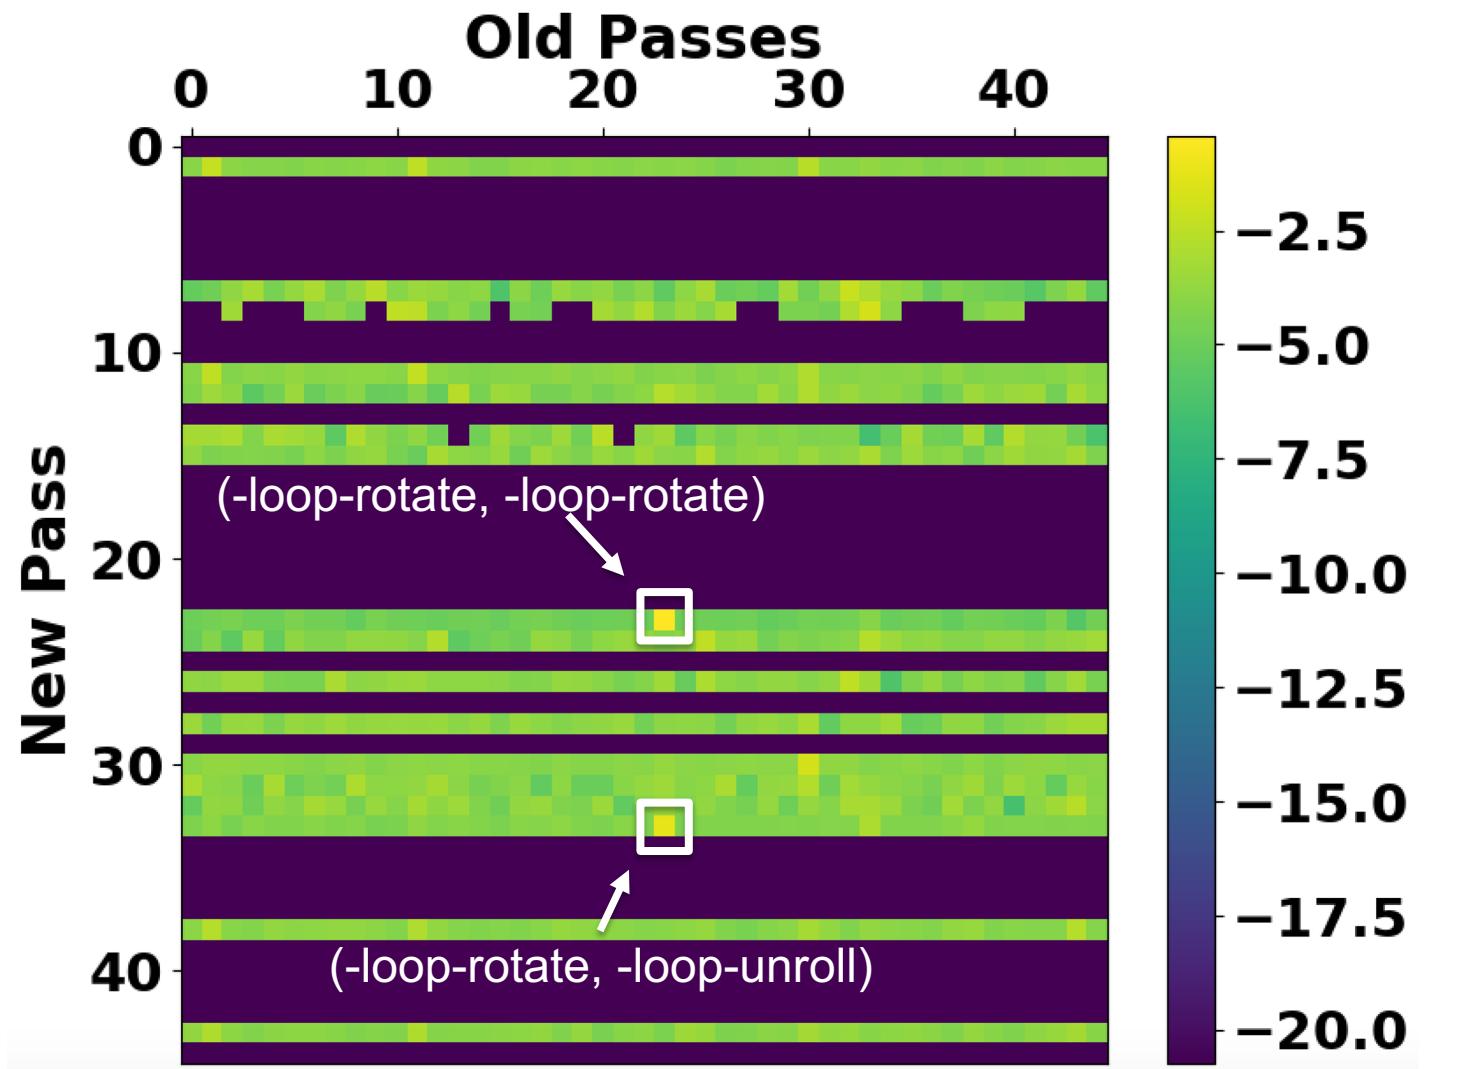
- Can we train a network and use it to predict the optimal sequence of any program in a single compilation? **Very Difficult!**
  - More than  $45^{45}$  possible phase orderings
  - Limited observations
  - Limited data
  - Decisions are not explainable

# Understanding the Correlations between Actions and States

- Two random forests for each optimization pass



# Importance Analysis

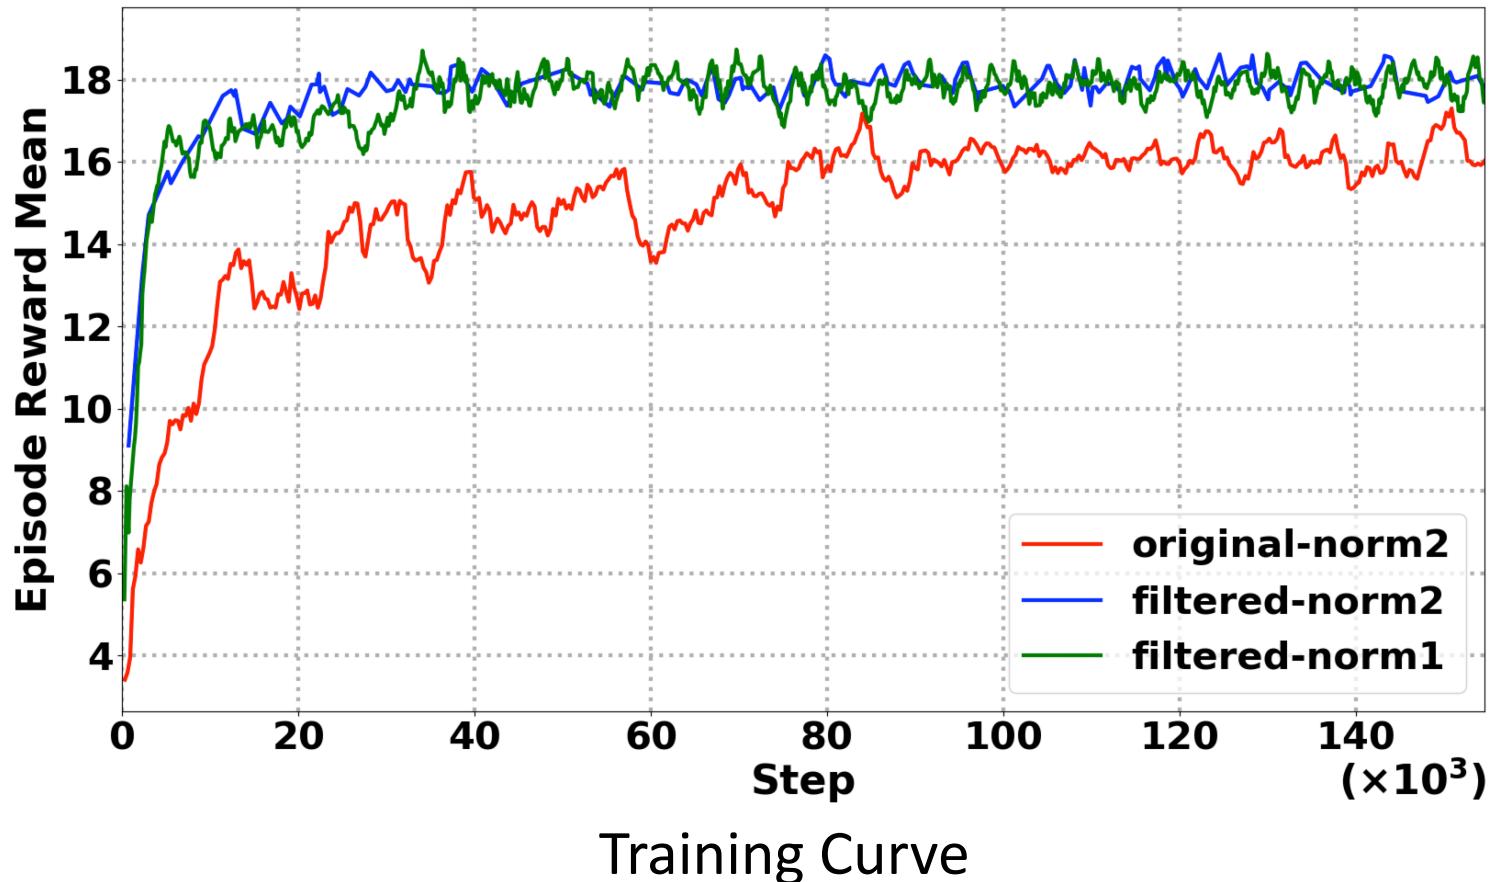


# Normalization Schemes

1. Logarithm of program features and rewards:
  - Reduces the magnitude of features/rewards
  - The neural network learns to correlate the products of features instead of a linear combination of them
    - $w_1 \log(o_{f_1}) + w_2 \log(o_{f_2}) = \log(o_{f_1}^{w_1} \cdot o_{f_2}^{w_2})$
2. Dividing the program feature values by the total number of instructions

# Training Results

- On 100 randomly-generated programs



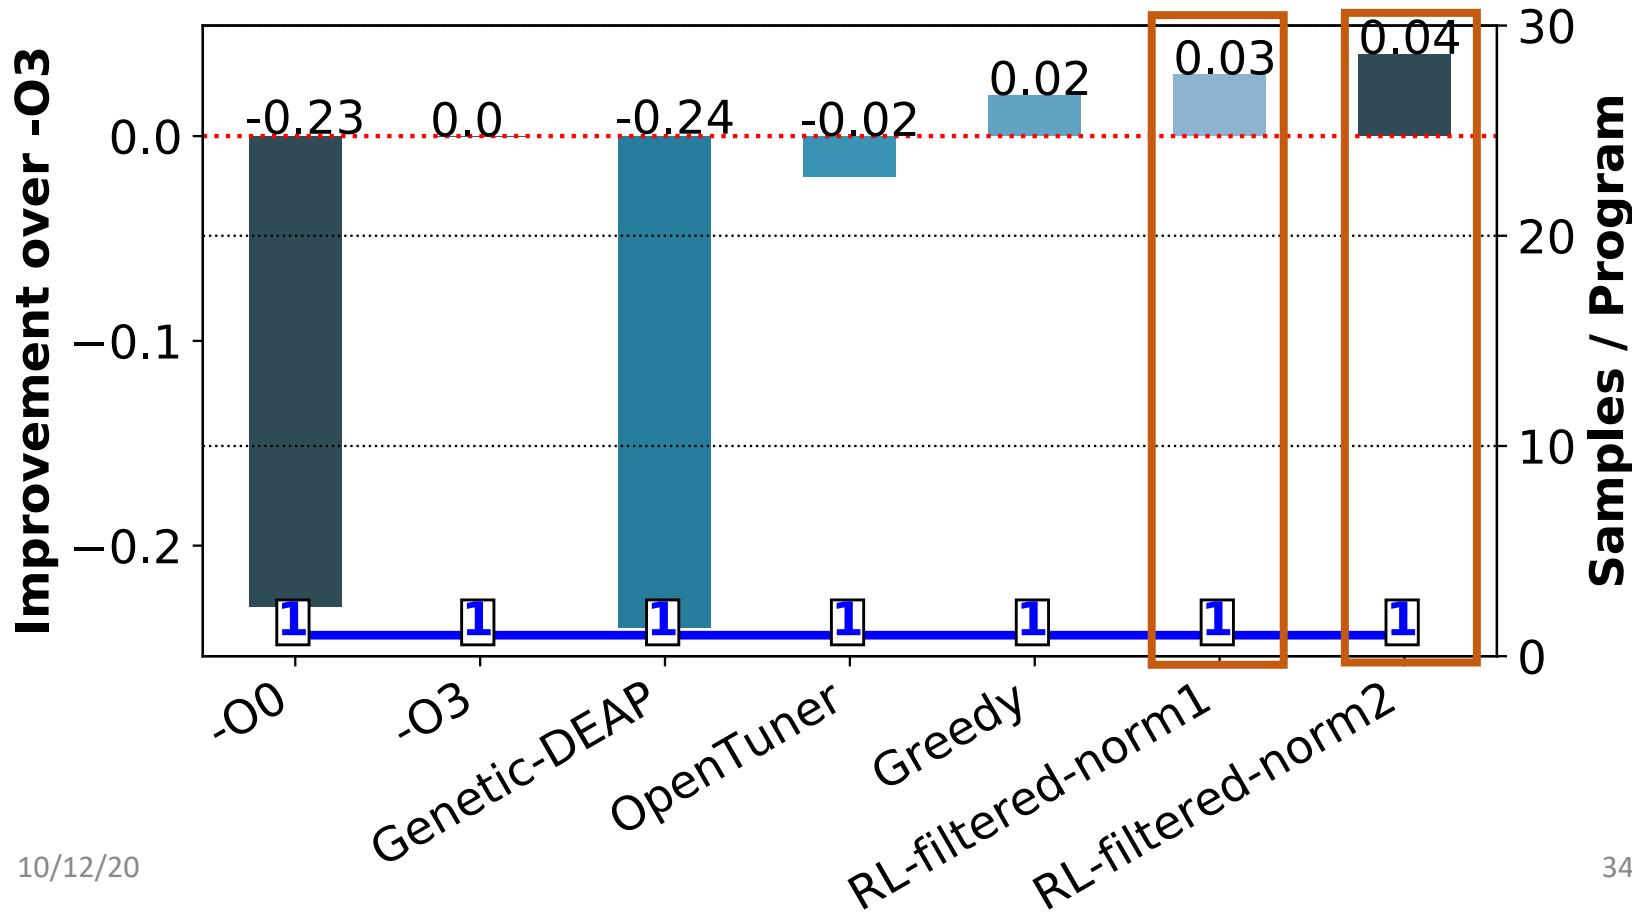
# Generalization Results

## Inference

- 13,000 randomly-generated programs
  - 6% better than –O3
- CHStone benchmarks
  - different from the randomly-generated programs

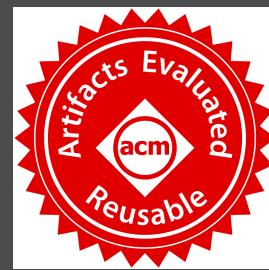
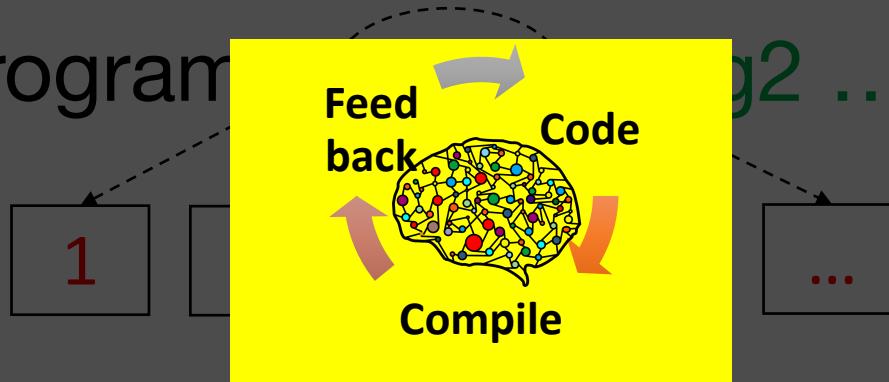
# Generalization Results

- ~4% improvement over -O3 on CHStone
- -O3 fails to compile the gsm benchmark



# It iNeuroVectorizer: End-to-End Vectorization with Deep Reinforcement Learning

clang program



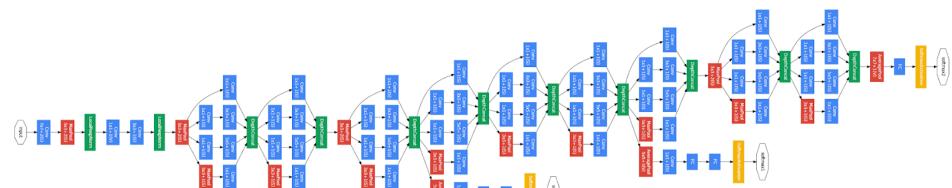
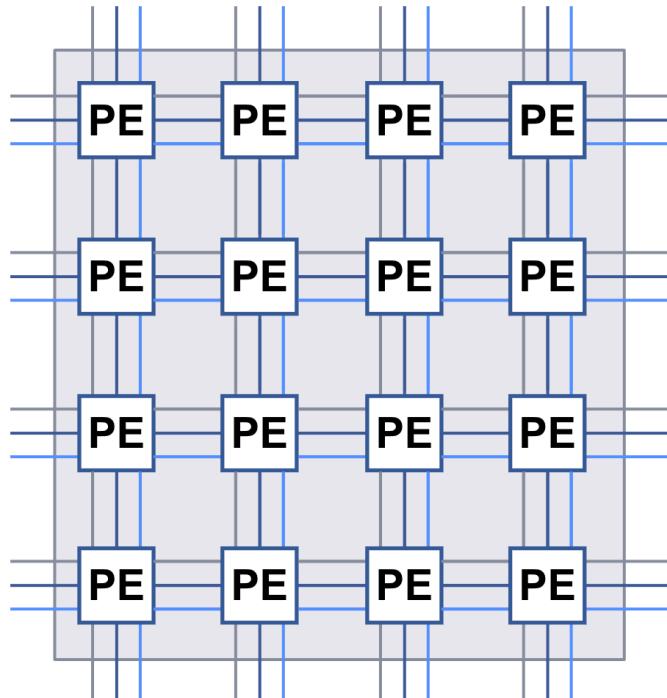
# Ongoing Work

More optimizations on CPUs/GPUs:

- Scheduling
- Multithreading
- Loop transformations
  - Vectorization
  - Parallelization
  - Tiling
  - Interchange
  - Fusion

# Ongoing Work

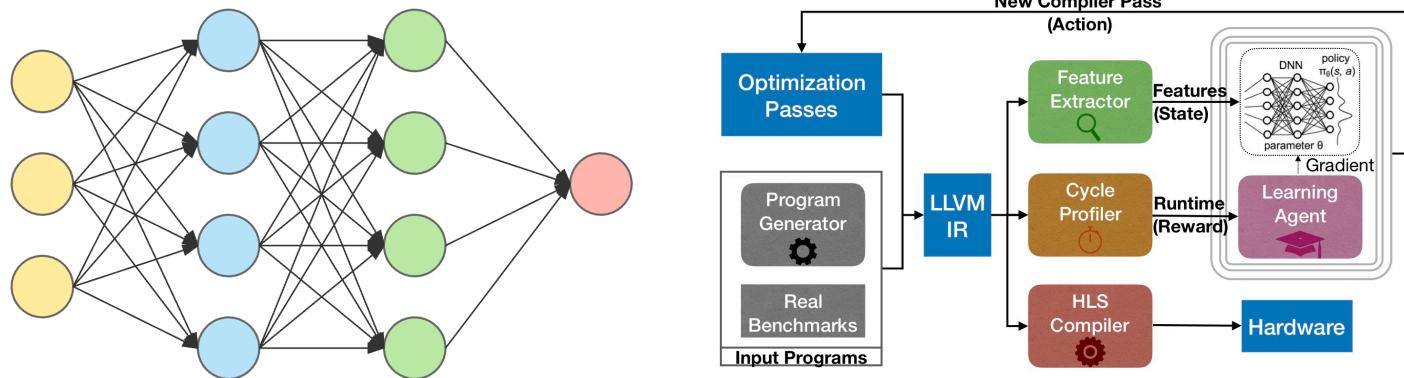
## NoC Scheduling



- 1. Temporal Organization**
  - Tasks to execute next
  
- 2. Spatial Organization**
  - Allocation & Placement

# Conclusions

- RL shows potential to tackle the phase ordering problem with fewer samples
- AutoPhase achieves 28% circuit speedup over -O3
- AutoPhase generalizes to real benchmarks from training on random programs
- AutoPhase is compatible with SW optimization



Access to code: <https://github.com/ucb-bar/autophase>

# Thanks!

ameerh@berkeley.edu

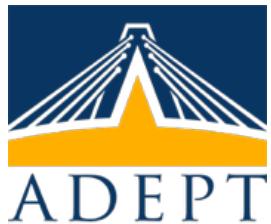


[ameerhajali.com](http://ameerhajali.com)

qijing.huang@berkeley.edu



Berkeley  
Architecture  
Research



 **rise lab**  
UC Berkeley