




Mind the Gap: Attainable Data Movement and Operational Intensity Bounds for Tensor Algorithms

Qijing Huang, Po-An Tsai, Joel Emer, Angshuman Parashar


NVIDIA, MIT CSAIL

Motivation: A design challenge

Algorithm




GPTx




Motivation: A design challenge

Algorithm




GPTx




Motivation: A design challenge

Algorithm

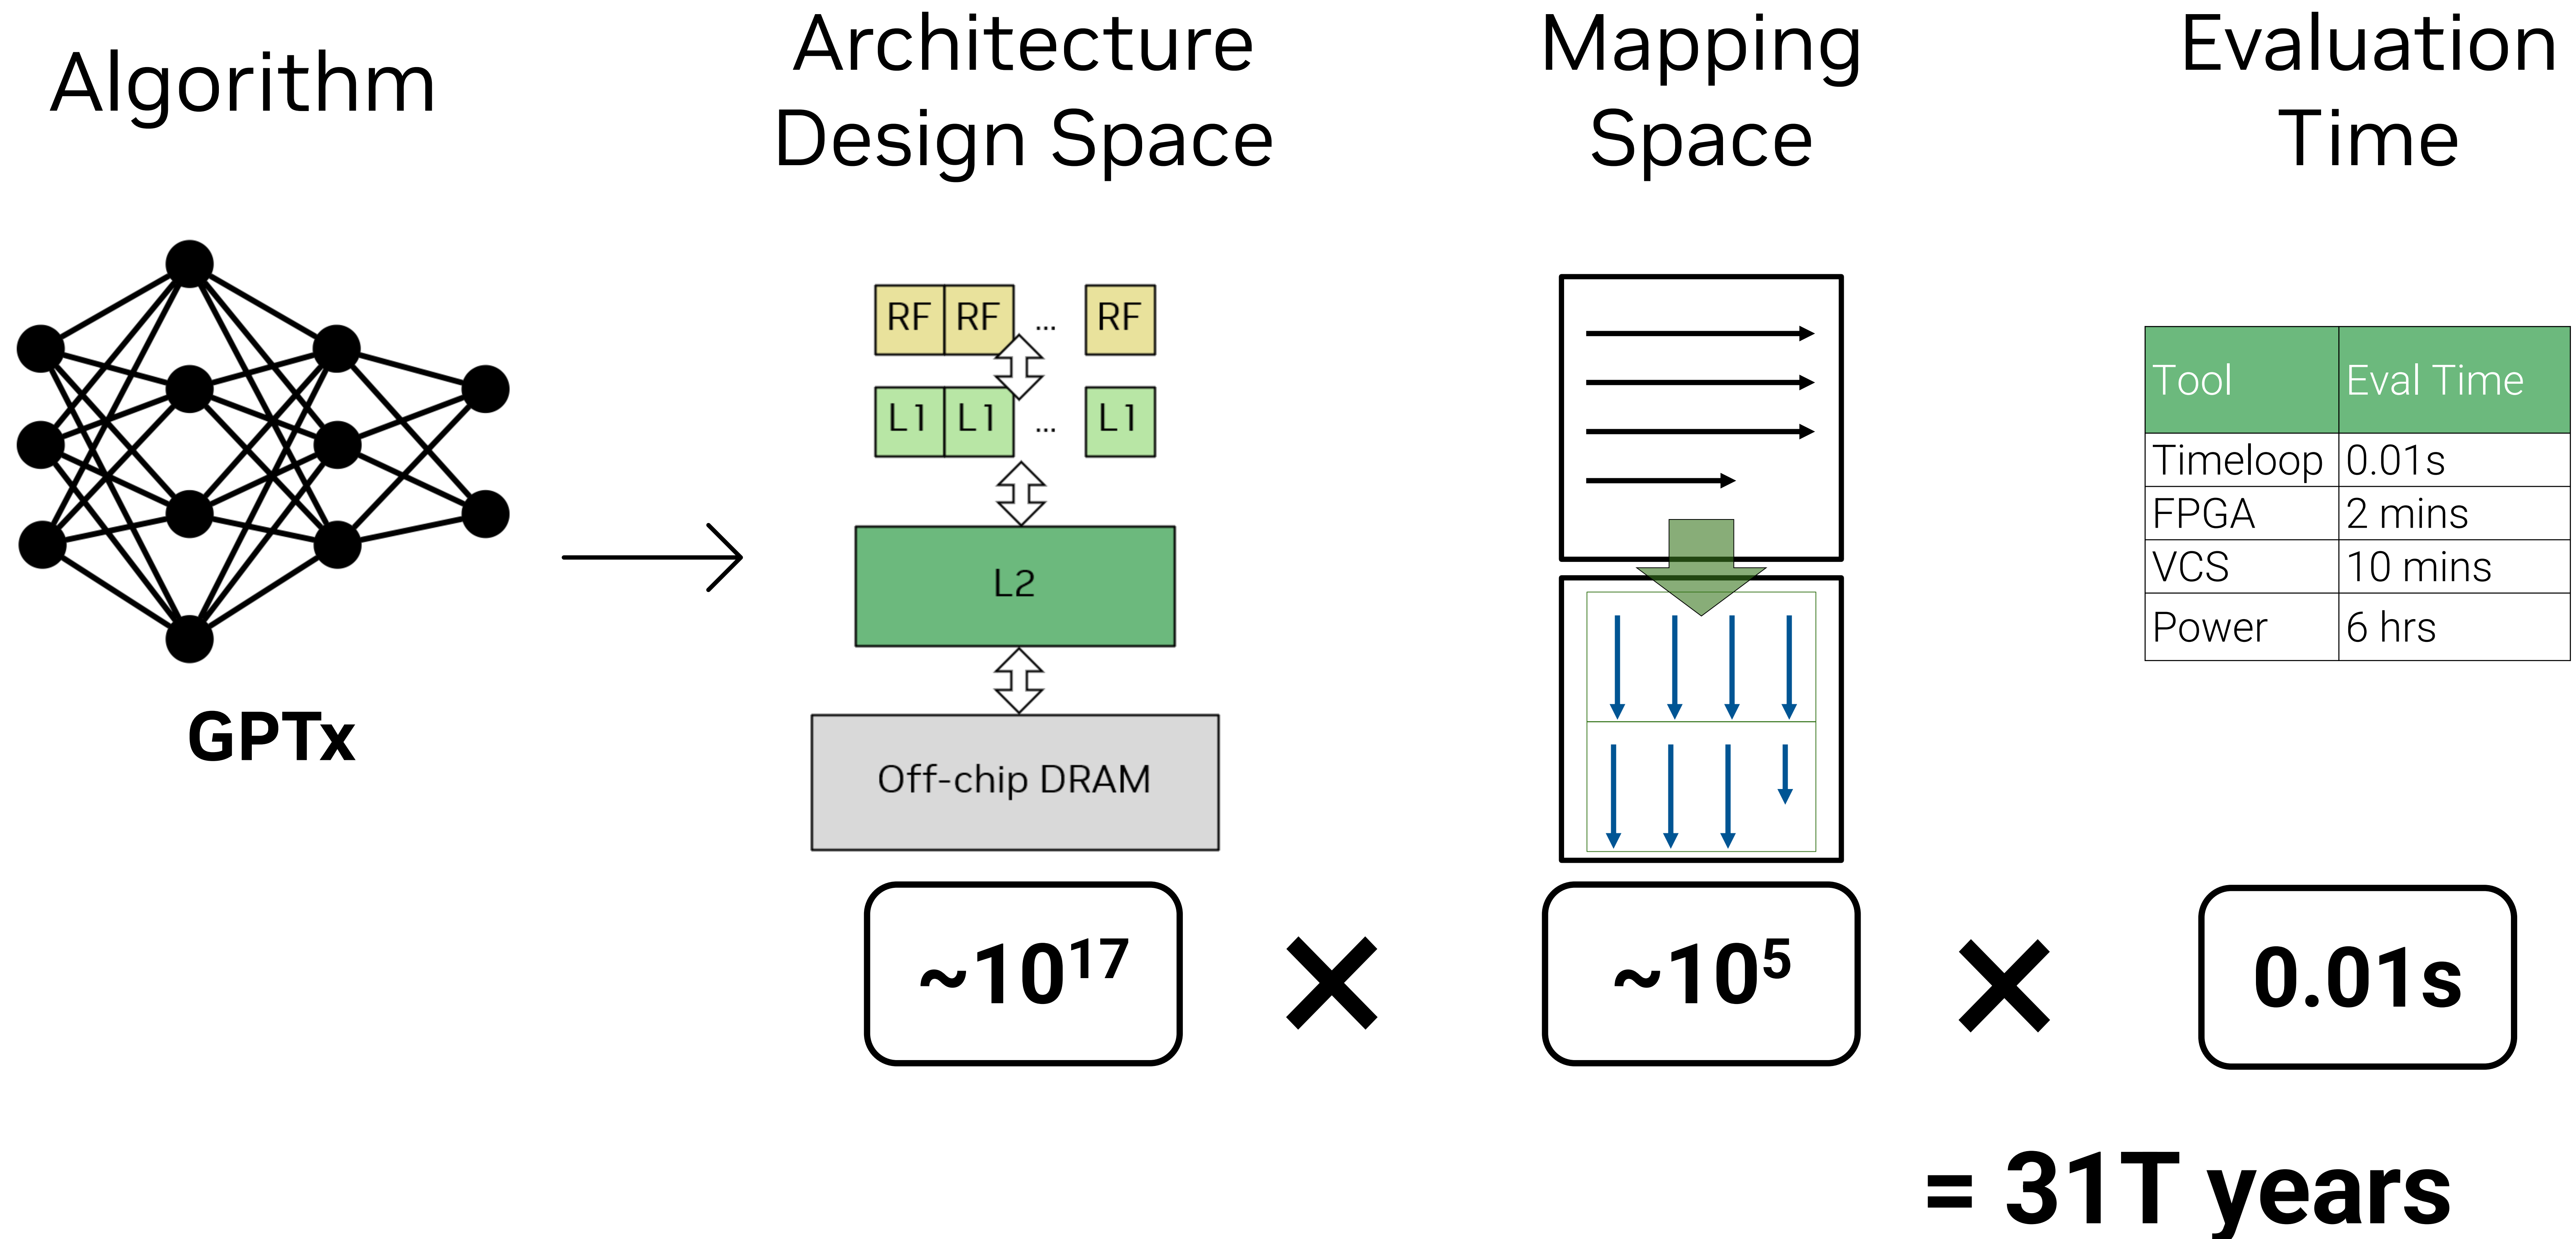


GPTx




How to provision chip area
between storage and compute?

Approach 1: Design space exploration




Approach 1: Design space exploration

Algorithm




GPTx

Architecture
Design Space



Mapping
Space



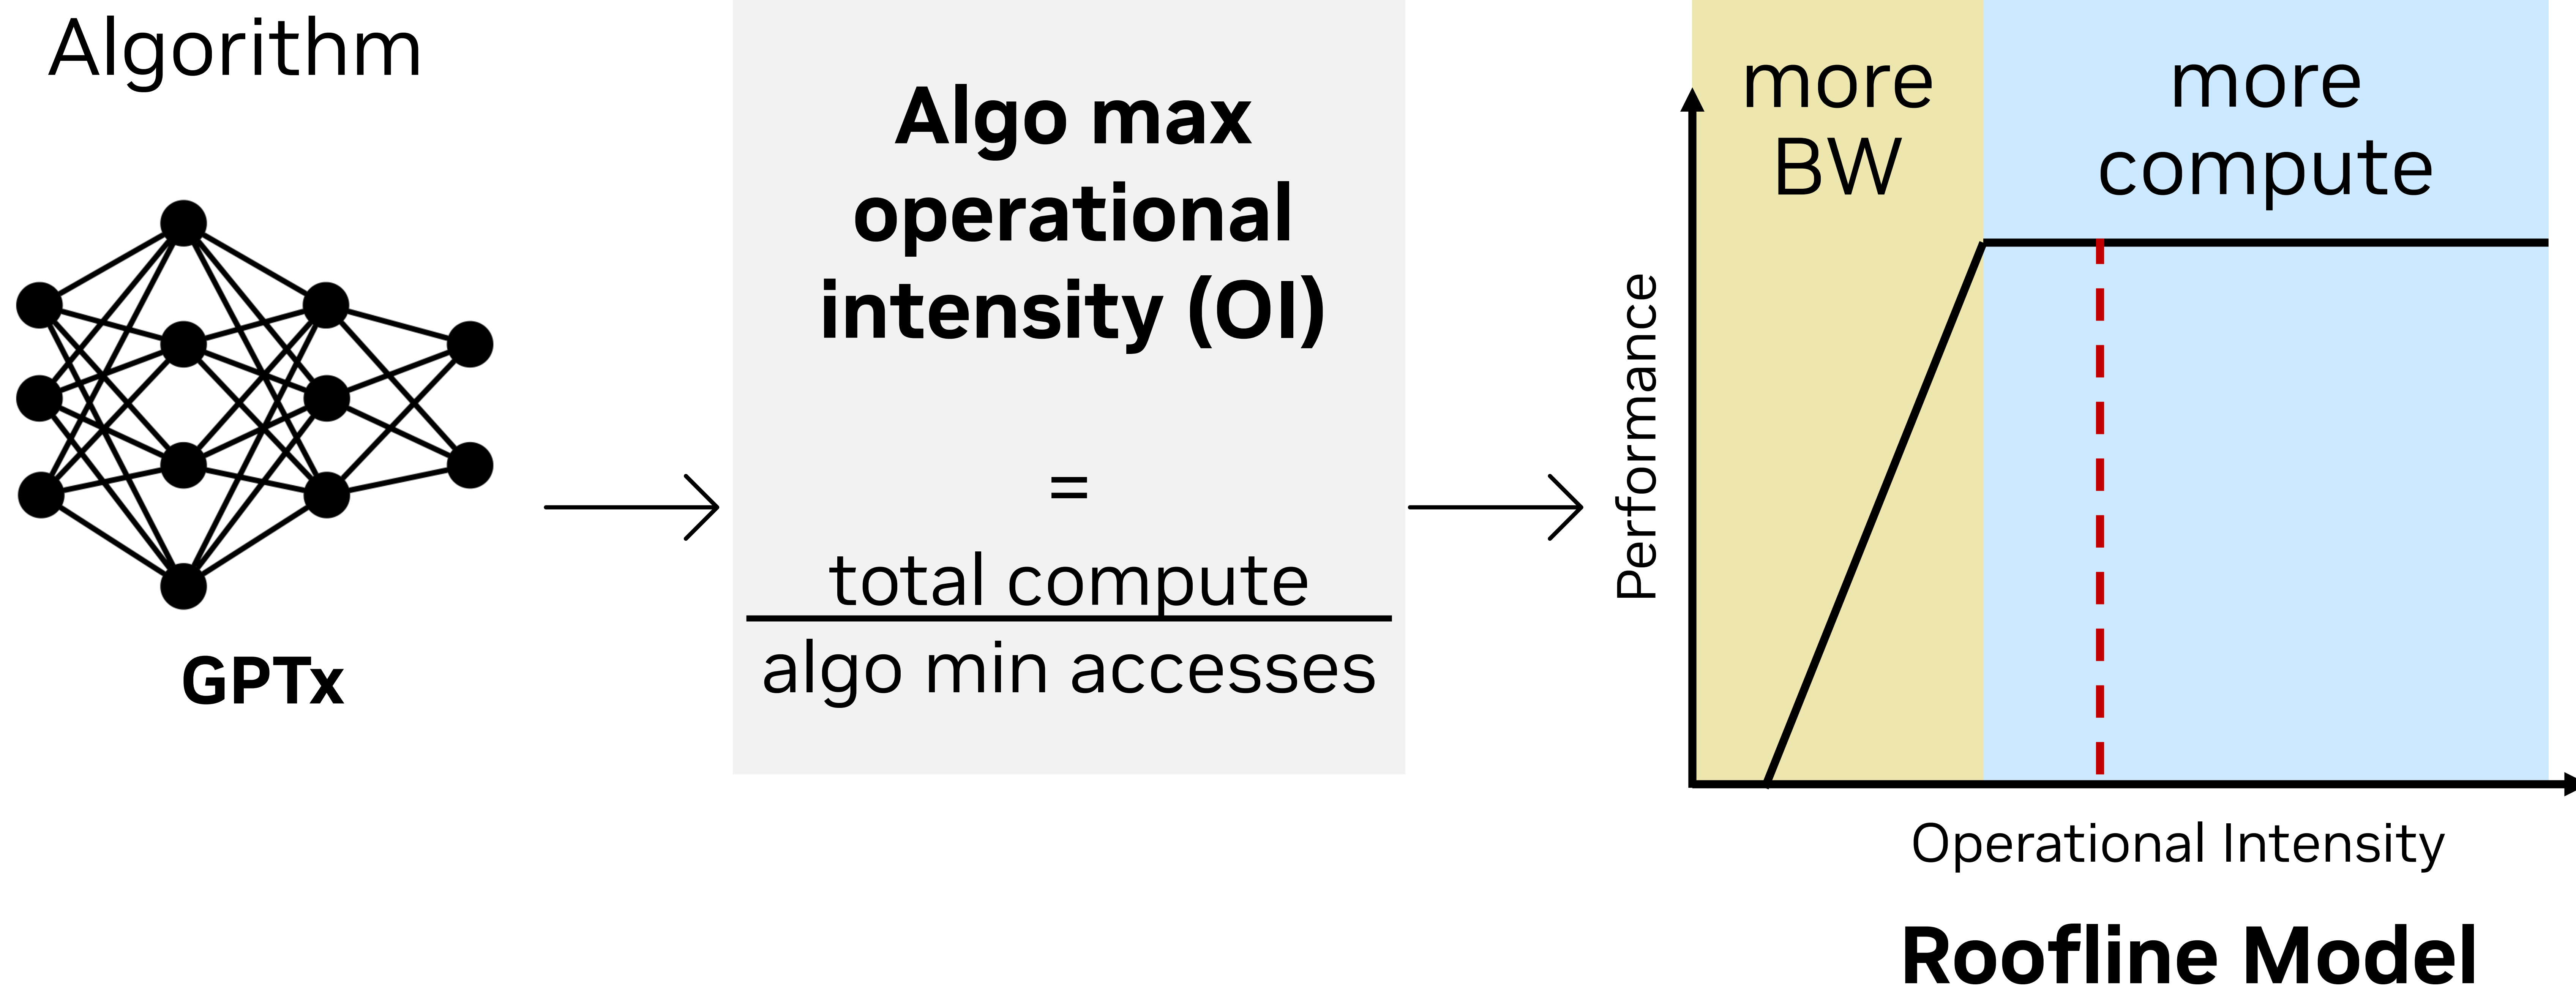
Evaluation
Time

Tool	Eval Time
Timeloop	0.01s
FPGA	2 mins
VCS	10 mins
Power	6 hrs

- Time-consuming and costly
- No optimality guarantee
- Lack of design insight

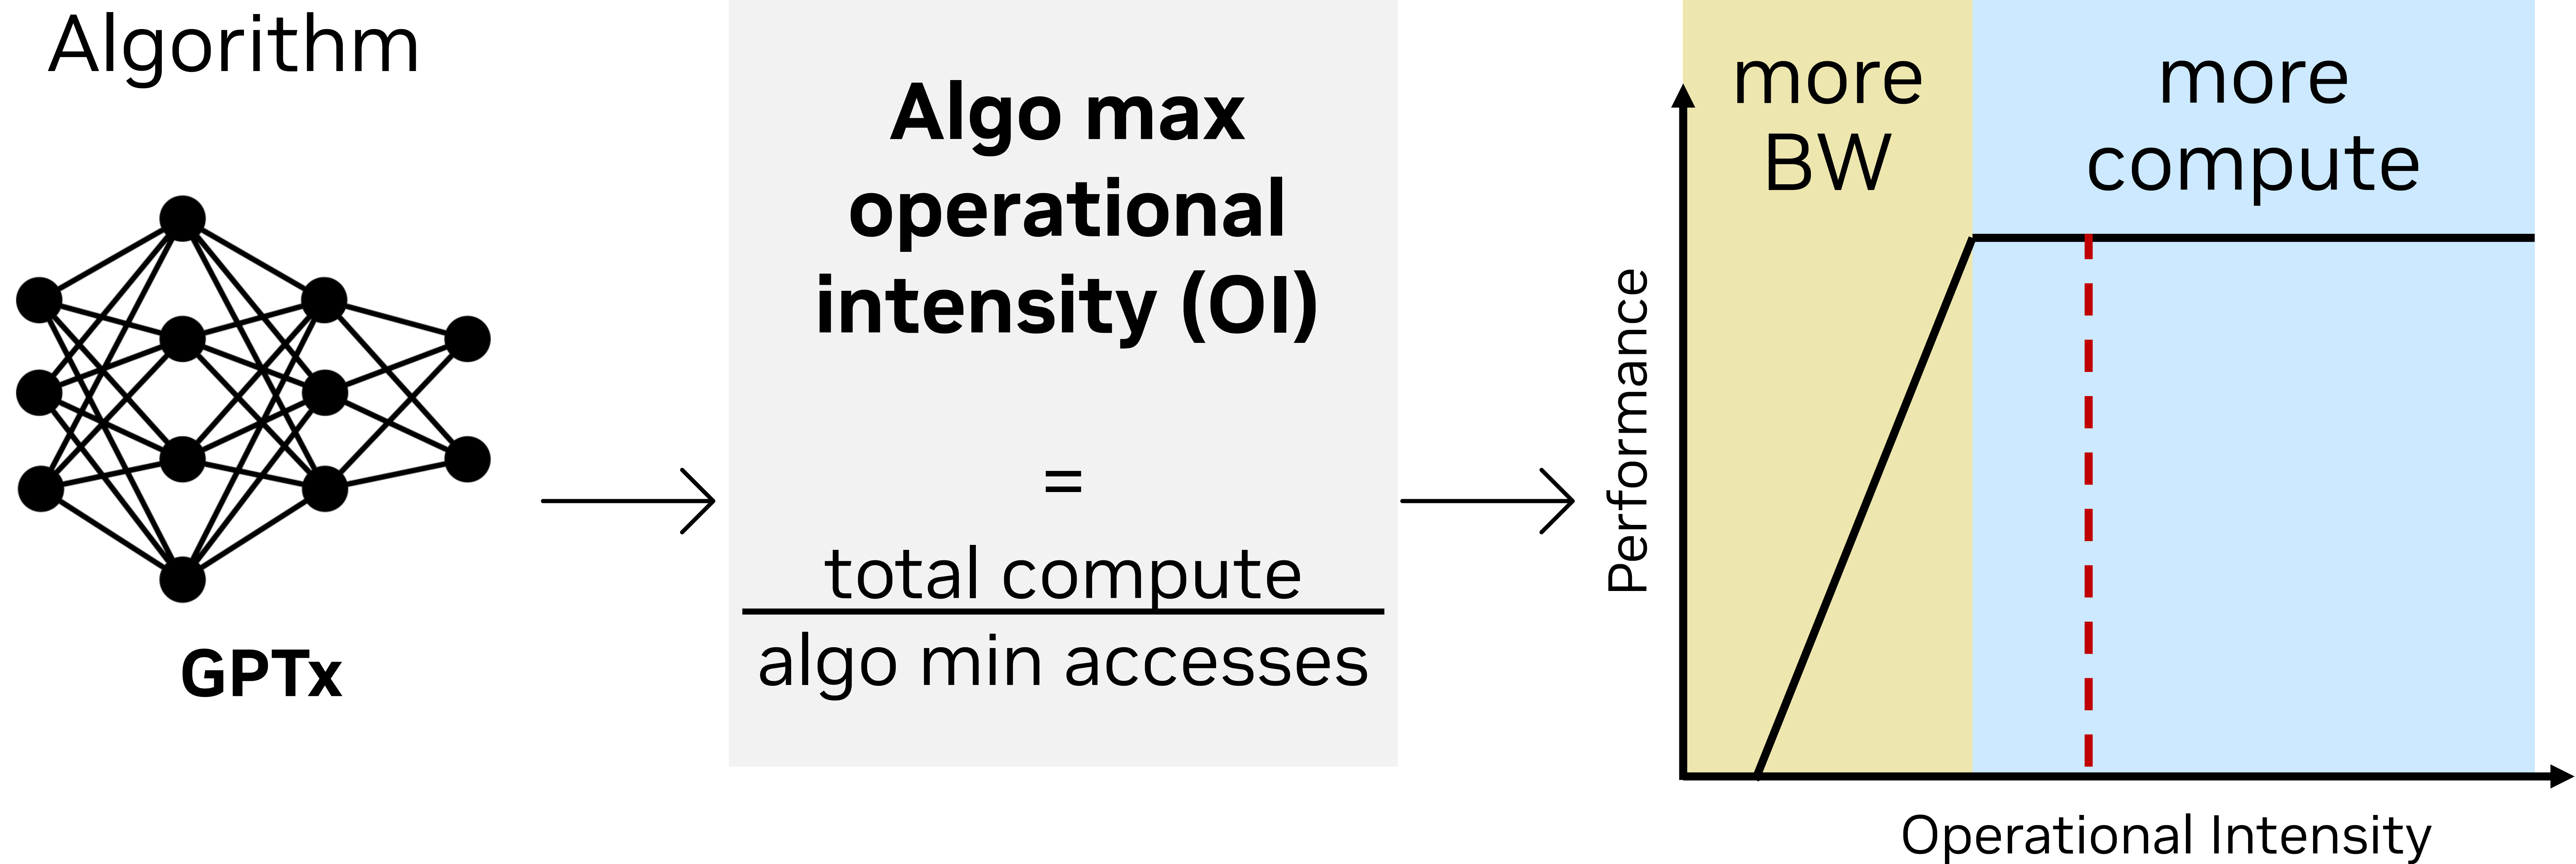
Approach 2: Roofline model analysis

“Speeds and feeds”



Approach 2: Roofline model analysis



“Speeds and feeds”



- No buffer storage tradeoffs are present in the analysis

Roofline Model

Motivation: Lack of design tools




How to provision chip area
between storage and compute?



What is missing?


- The workload **does not** always operate with *algorithmic minimal accesses*, or equivalently, *algorithmic maximal OI*.



- Actual *backing-store accesses* and *OI* depend on the **mapping** and **buffer sizes**.

A desirable data movement bound

Matrix Multiplication
(GEMM) Einsum:




$$O_{m,n} = I_{m,k} W_{k,n}$$

M – output row dim


K – reduction dim

N – output column dim




“Ski-slope Diagram”

Mind the gap: key design questions



[Gap 0] Given a buffer capacity, what is the minimal attainable *data access count*?


Mind the gap: key design questions



[Gap 0] Given a buffer capacity, what is the minimal attainable *data access count*?

[Gap 1] How much buffer capacity is required to achieve full data reuse?

Mind the gap: key design questions



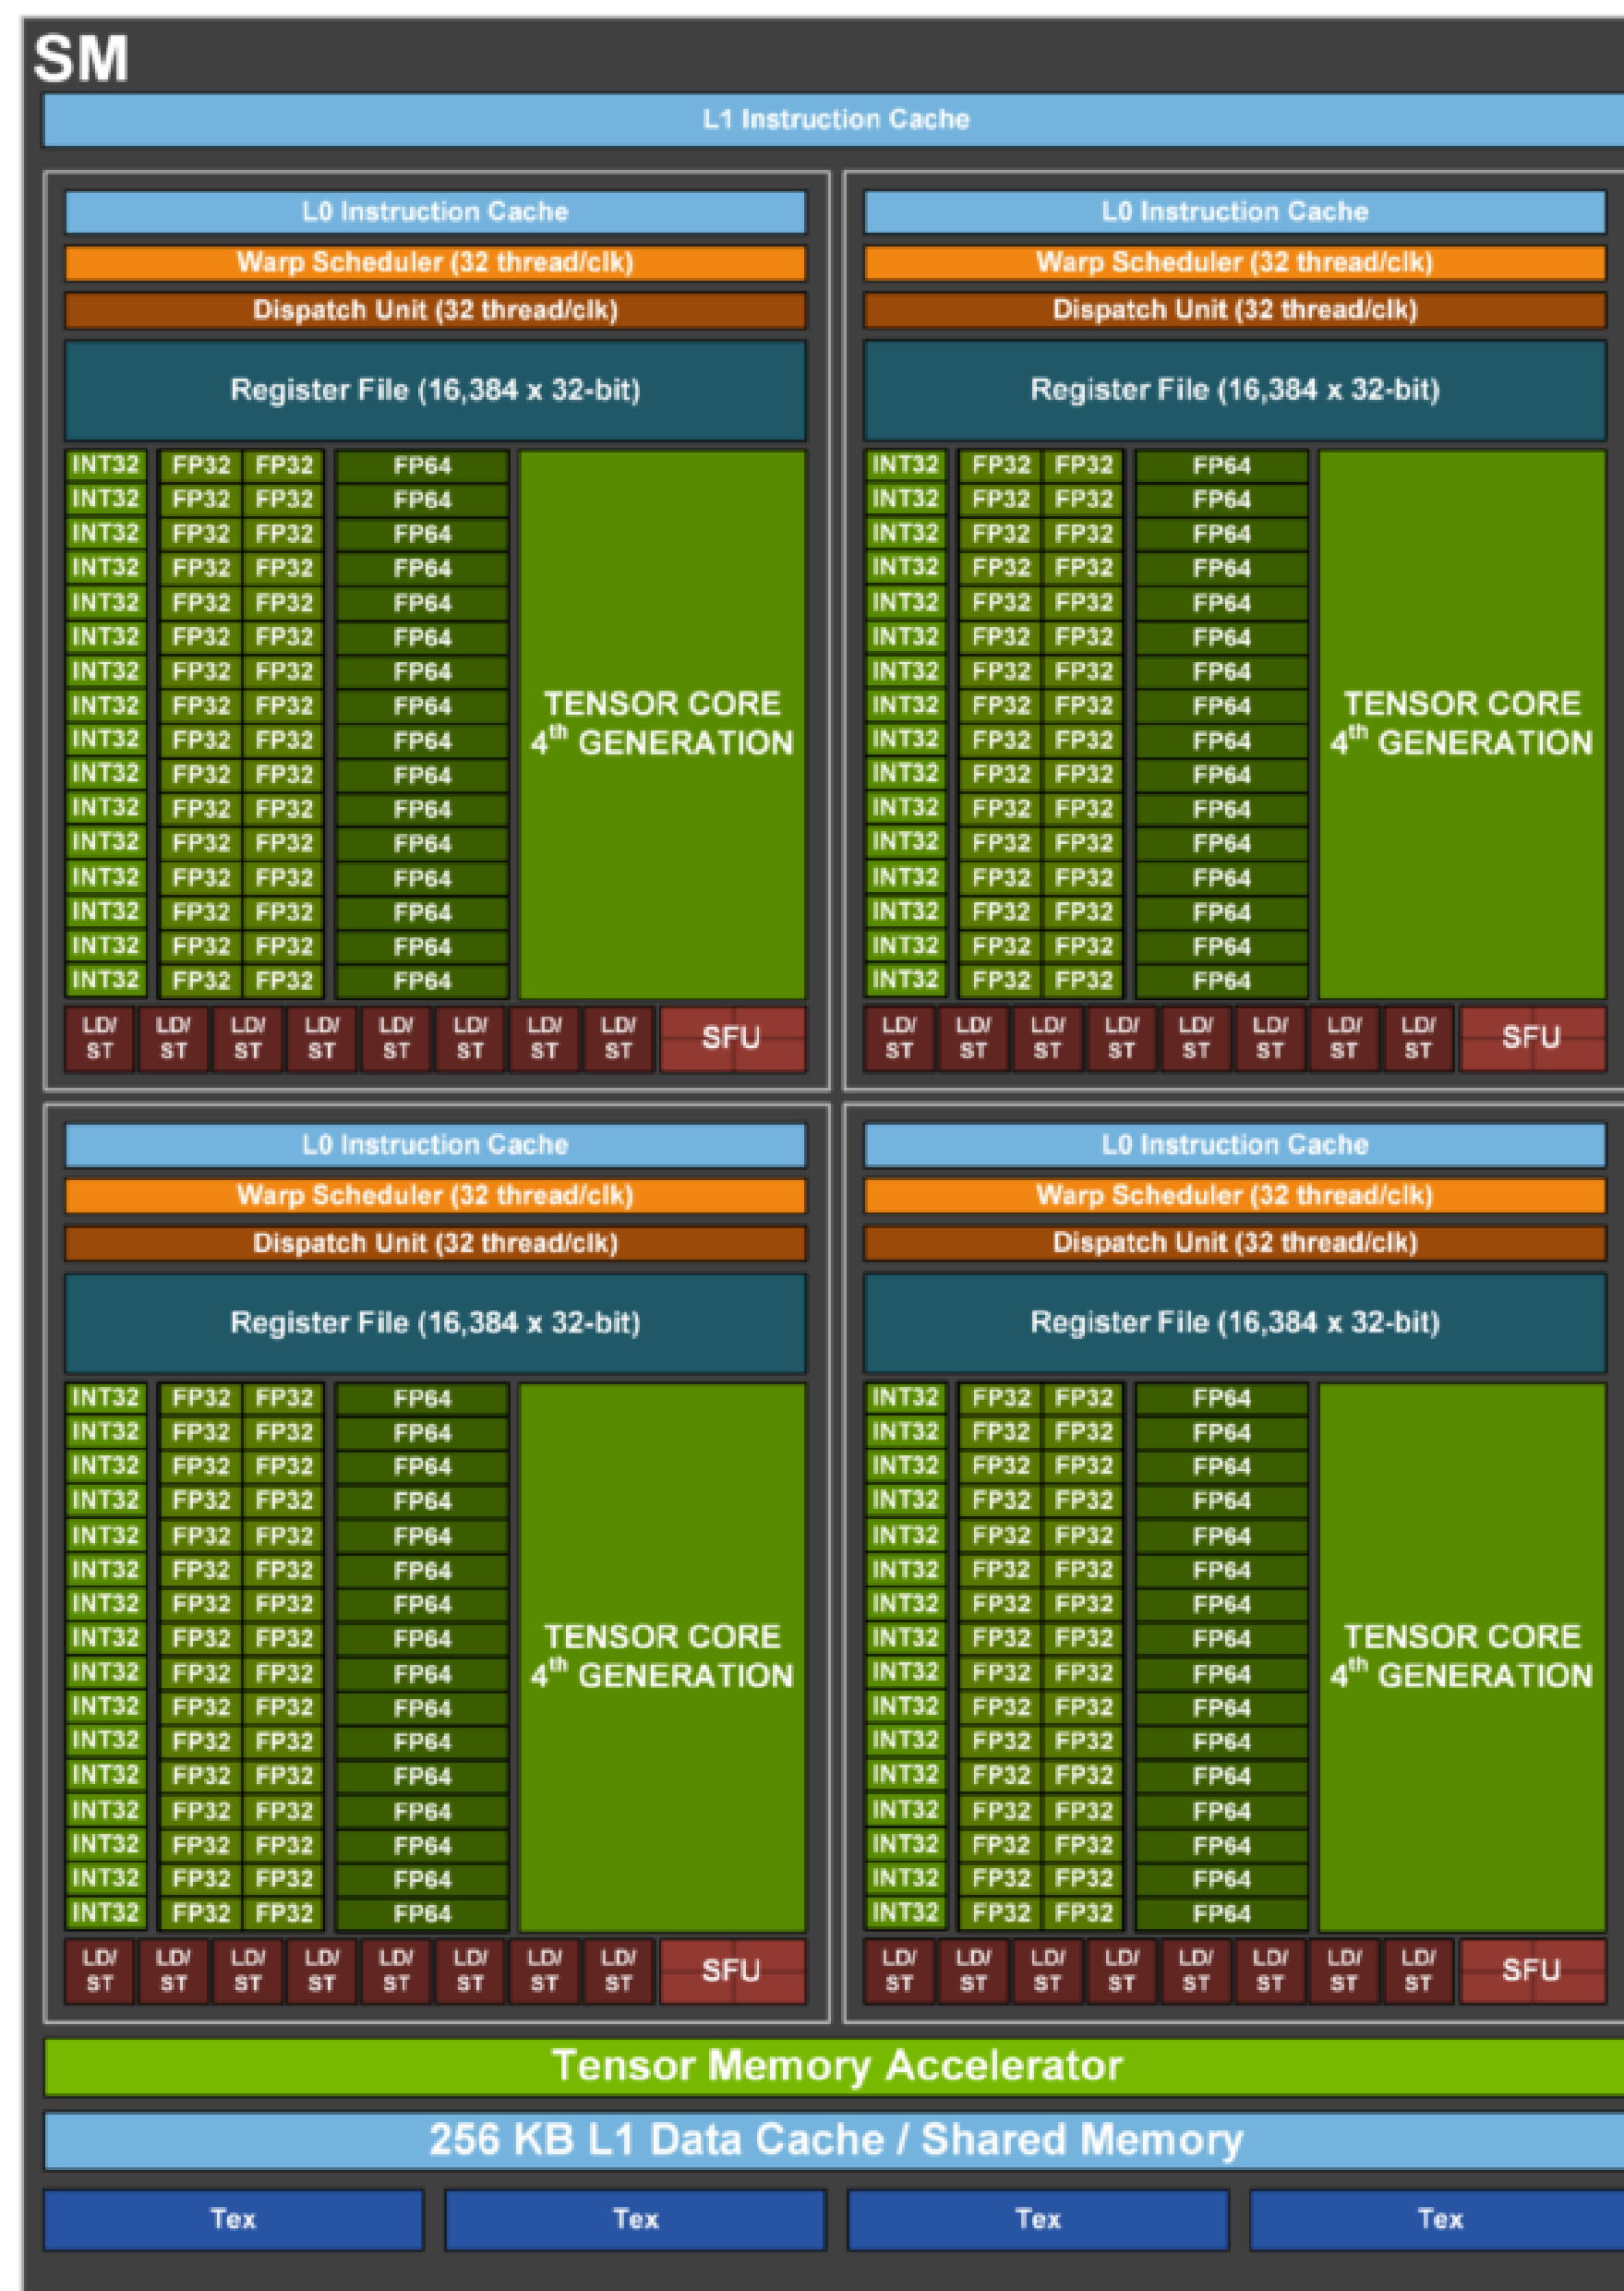
[Gap 0] Given a buffer capacity, what is the minimal attainable *data access count*?

[Gap 1] How much buffer capacity is required to achieve full data reuse?

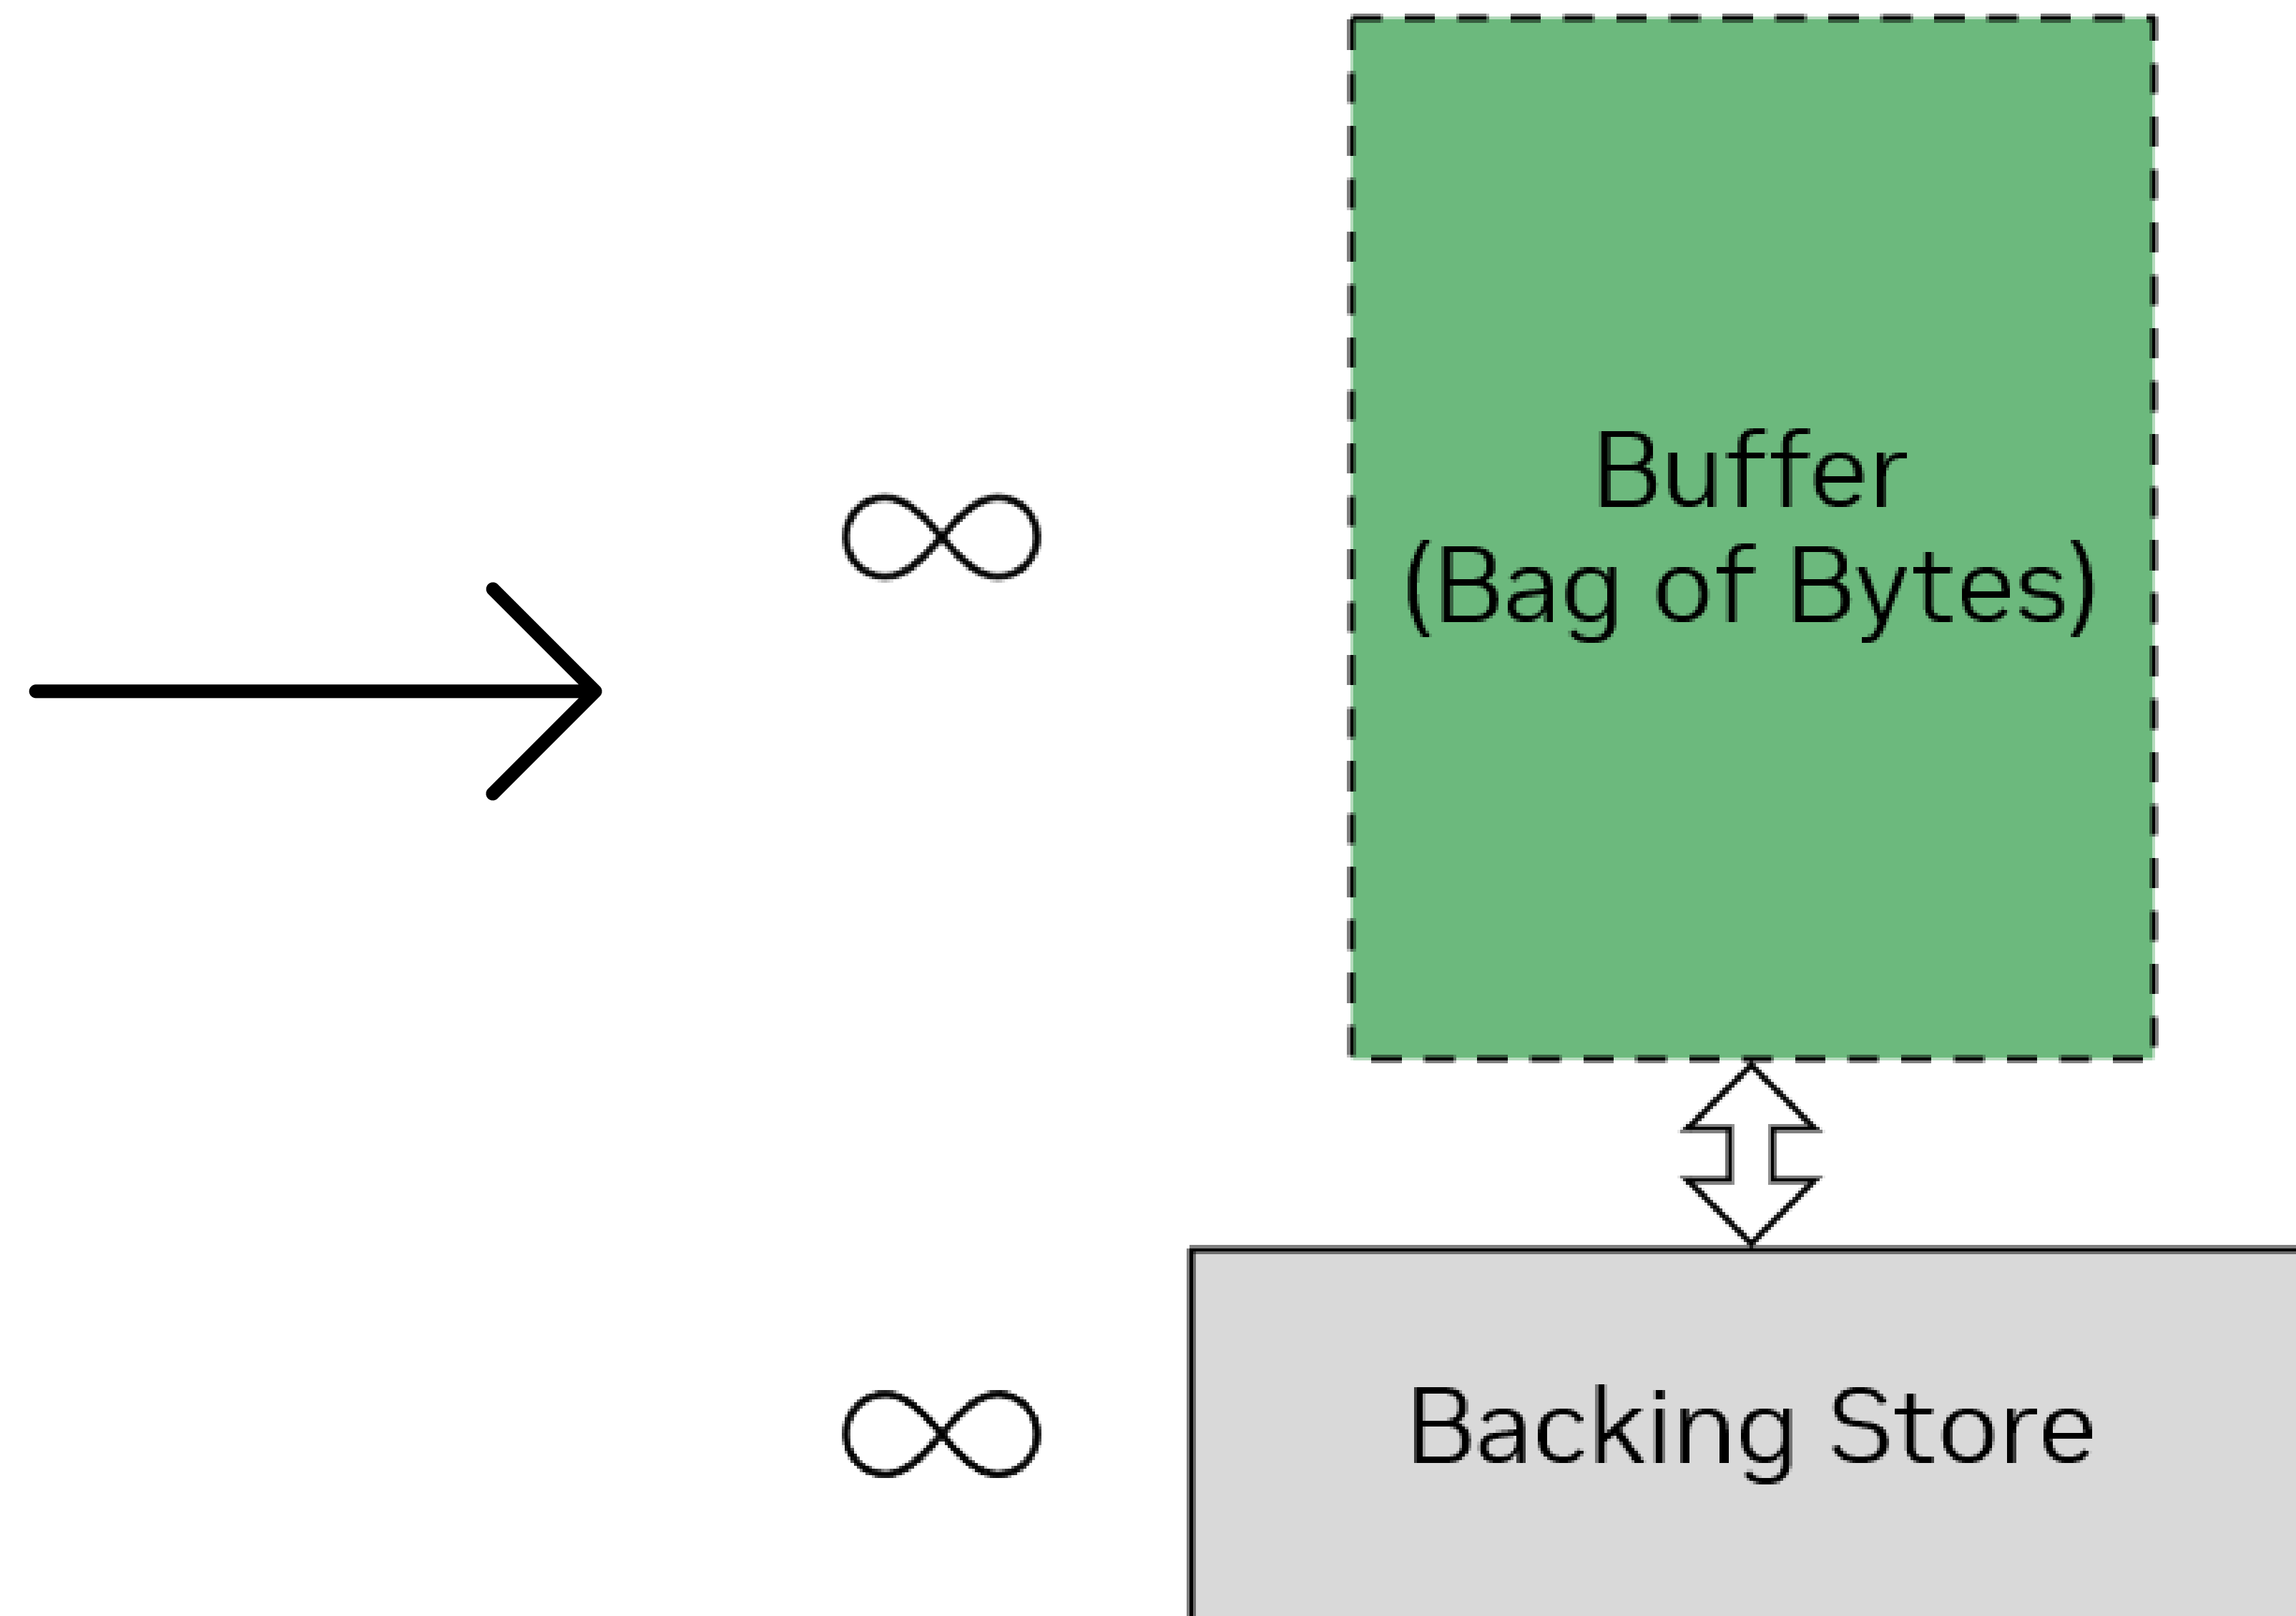
[rate of change of Gap 0] How does a workload benefit from incremental increase in buffer capacity?

The *Snowcat* Architecture

Enables exhaustive search



Real Design



Snowcat
Architecture

The *Orojenes* Methodology

A single exhaustive mapping search per workload


Inputs

Single Einsum ■

Chain of Einsums ■ ■ ■ ■

Exhaustive Search

Snowcat Arch



Mapping 0

Buffer Util

Accesses

Mapping 1

Buffer Util

Accesses


...

Mapping N

Buffer Util

Accesses


Ski-slope Diagram




$\min(\text{buffer util}, \text{accesses})$

OI Bound Derivation

Data Movement Bound




OI Bound



→
Inverse
&
Multiply by
total operations

Example *Orogenesis* Analysis

GEMM Einsum:



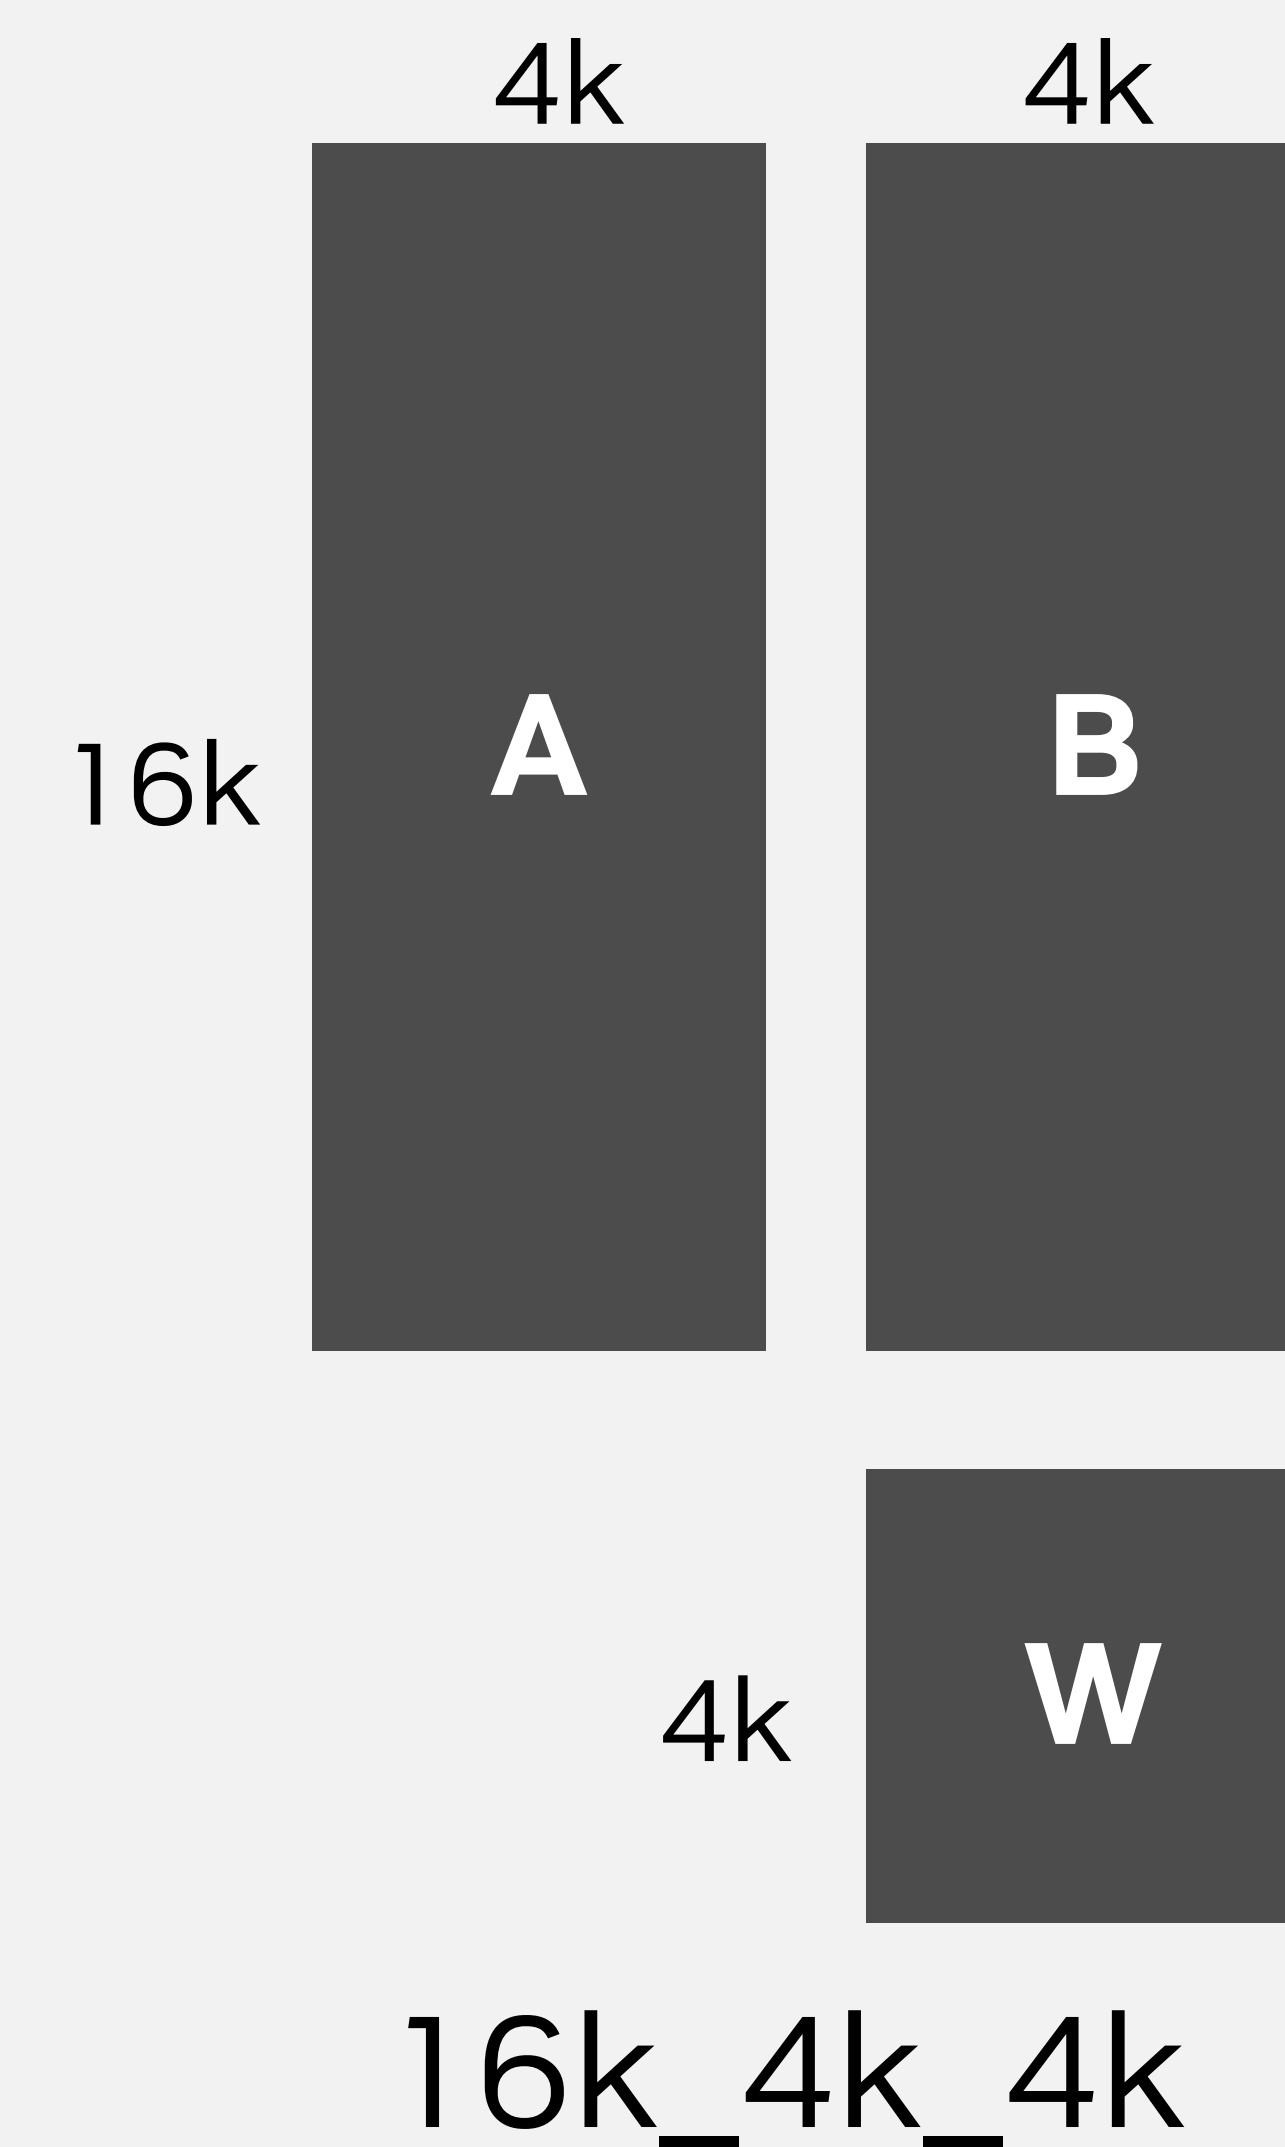


$$O_{m,n} = I_{m,k} W_{k,n}$$

M – output row dim

K - reduction dim

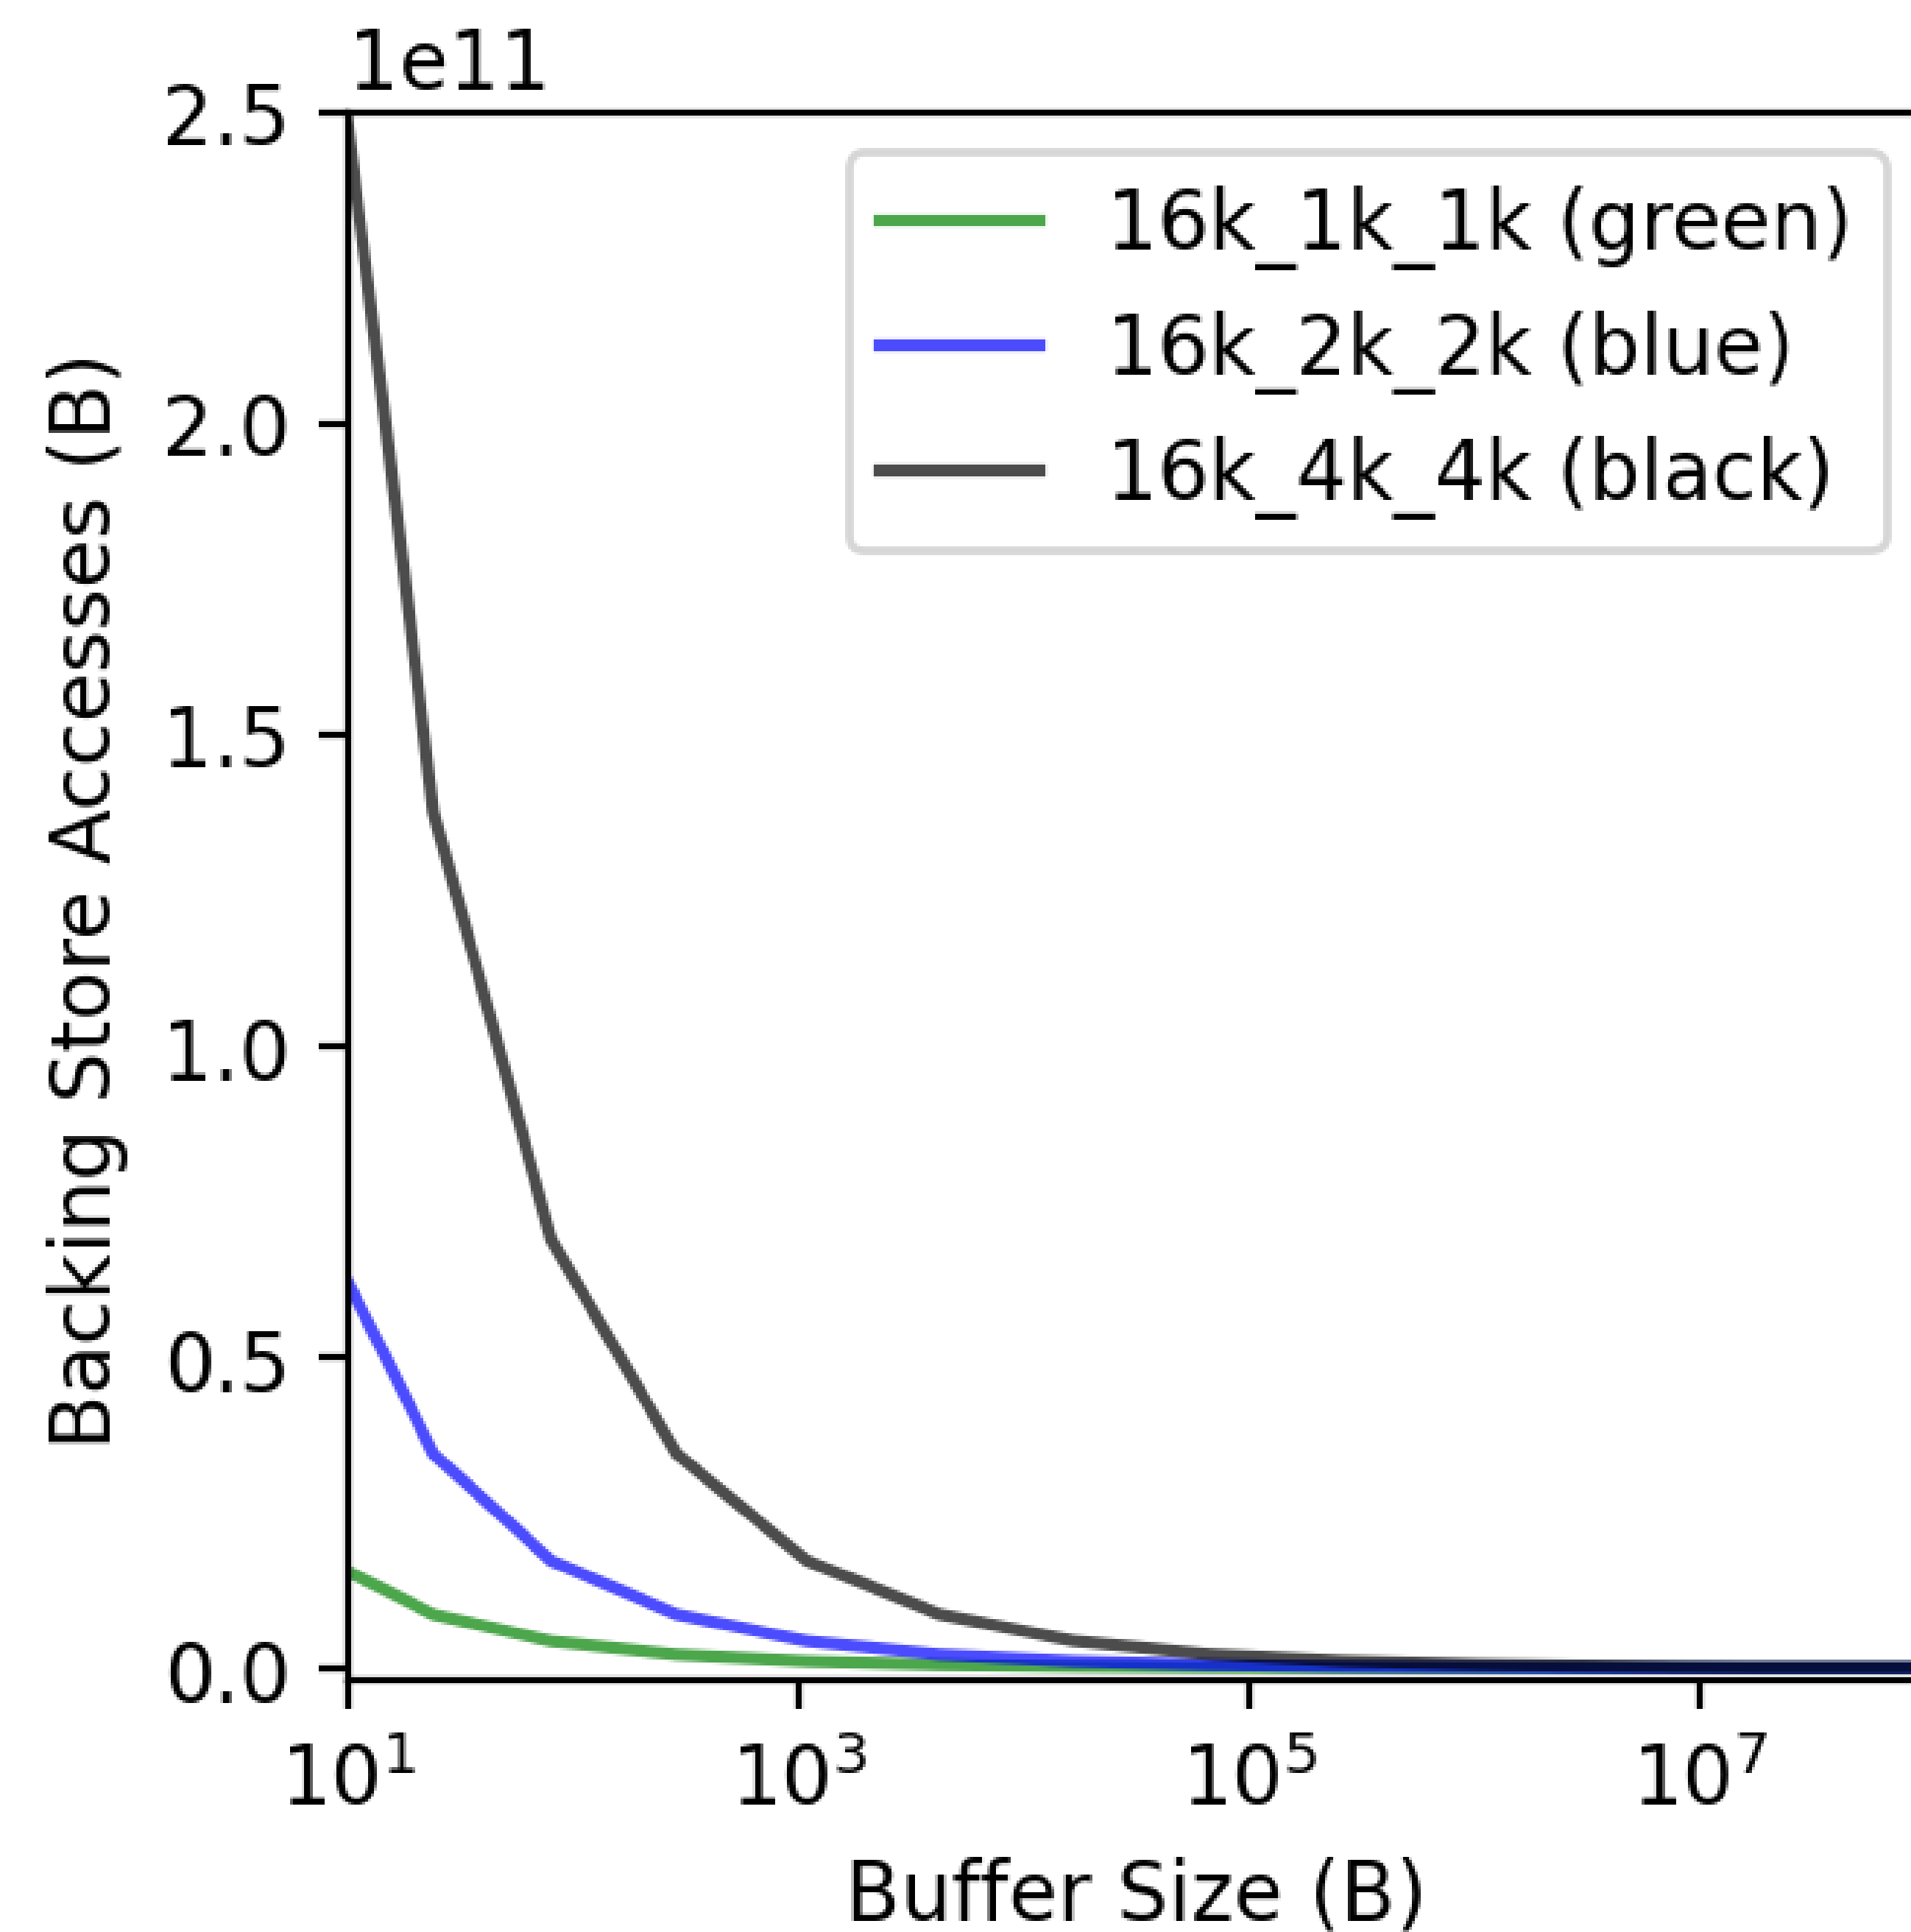
N – output column dim

Input Examples

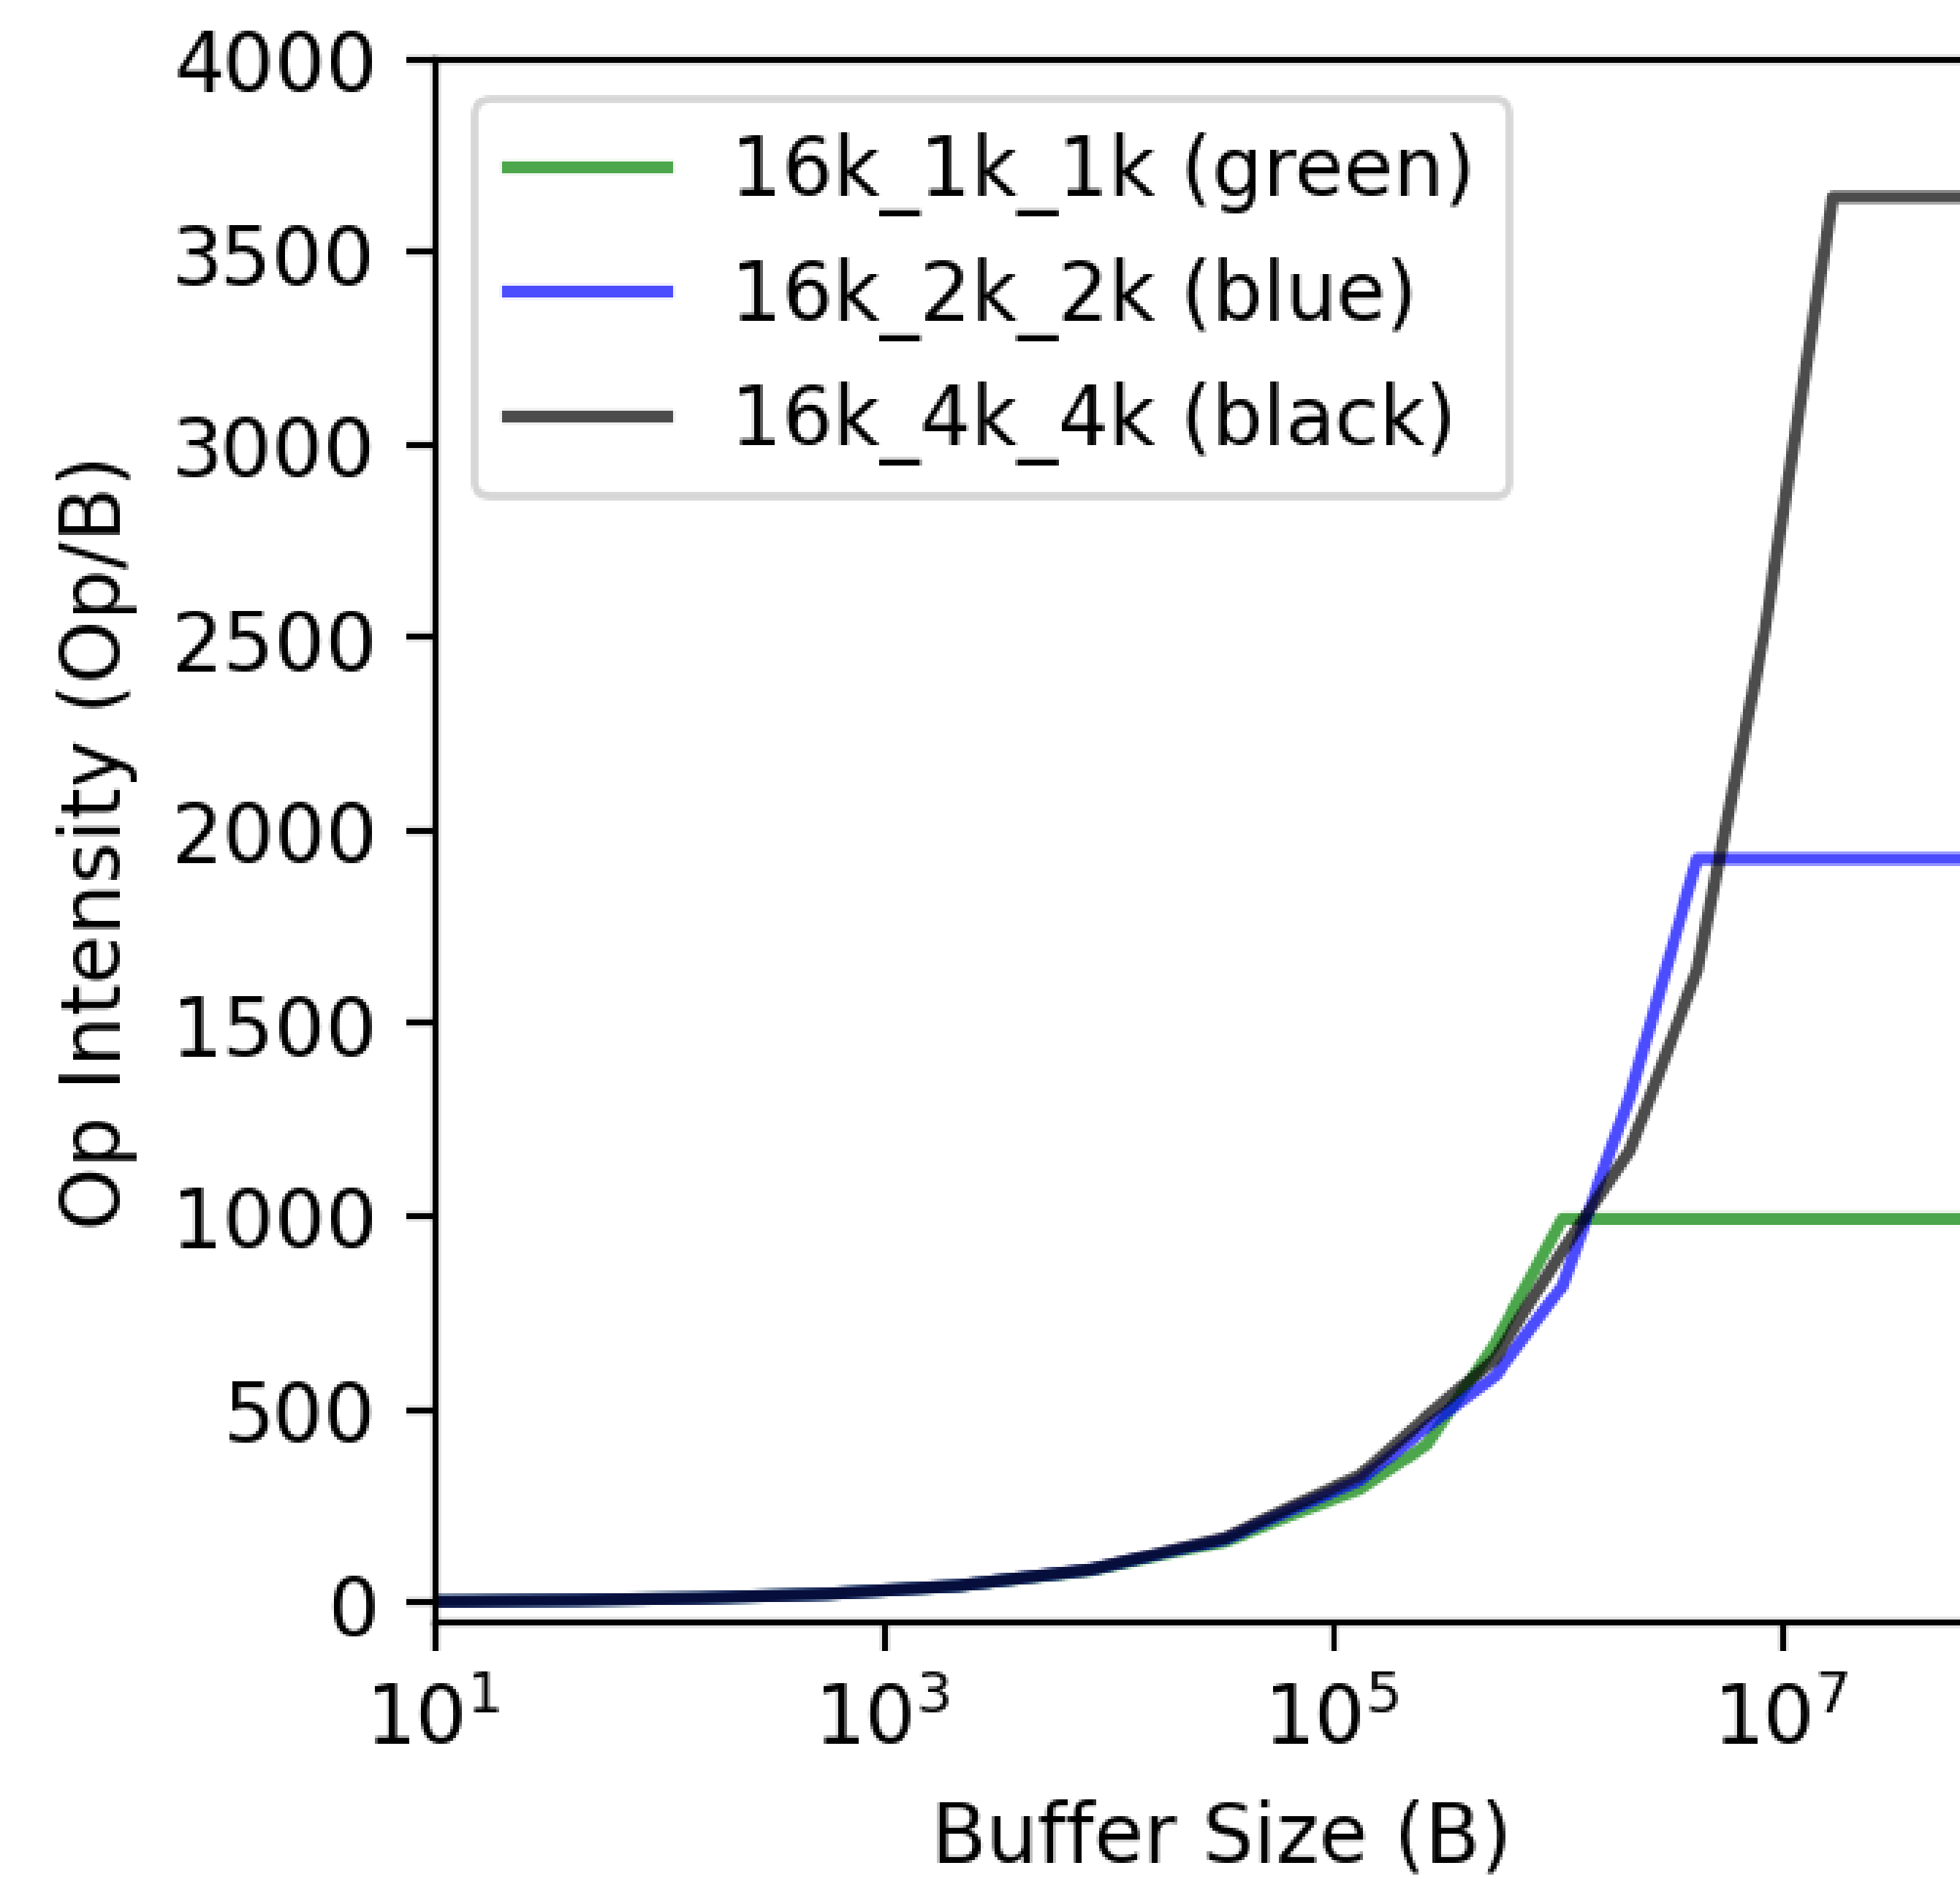


Example *Orojenes* Analysis

Data Movement Bound




OI Bound




Example *Orojenes* Analysis

Data Movement Bound




OI Bound




Example *Orojenes* Analysis

Data Movement Bound



OI Bound 1 4 16 MB



The maximal effectual buffer size of a GEMM is approximately its **smallest operand size**

#1: Orojenesi produces bounds that reveal powerful design insights


The *Orojenes* Fusion Flow

Inputs


Chain of Einsums



Unfused



Fused



Fusion is an effective technique to minimize data movement for a chain of operations

The *Orojenesis* Fusion Flow


Inputs

Chain of Einsums




Mapspace

Producer



Consumer



Unfused

The *Orojenesis* Fusion Flow

Inputs

Chain of Einsums





Mapspace

Producer

Incompatible

Consumer




w/ Fusion

The *Orojenes* Fusion Flow

Inputs

Chain of Einsums




Mapspace

Producer

Incompatible

Consumer




w/ Fusion

The *Orojenesis* Fusion Flow


Inputs

Chain of Einsums




Mapspace

Producer



Consumer




w/ Fusion


The *Orogenesis* Fusion Flow

Inputs

Chain of Einsums




Fusion Constraints



Original Flow


Snowcat Arch



Fusion imposes extra intra-layer mapping constraints


Fusion Analysis

A chain of 6 operations in GPT-6.7b block



Fusion Analysis

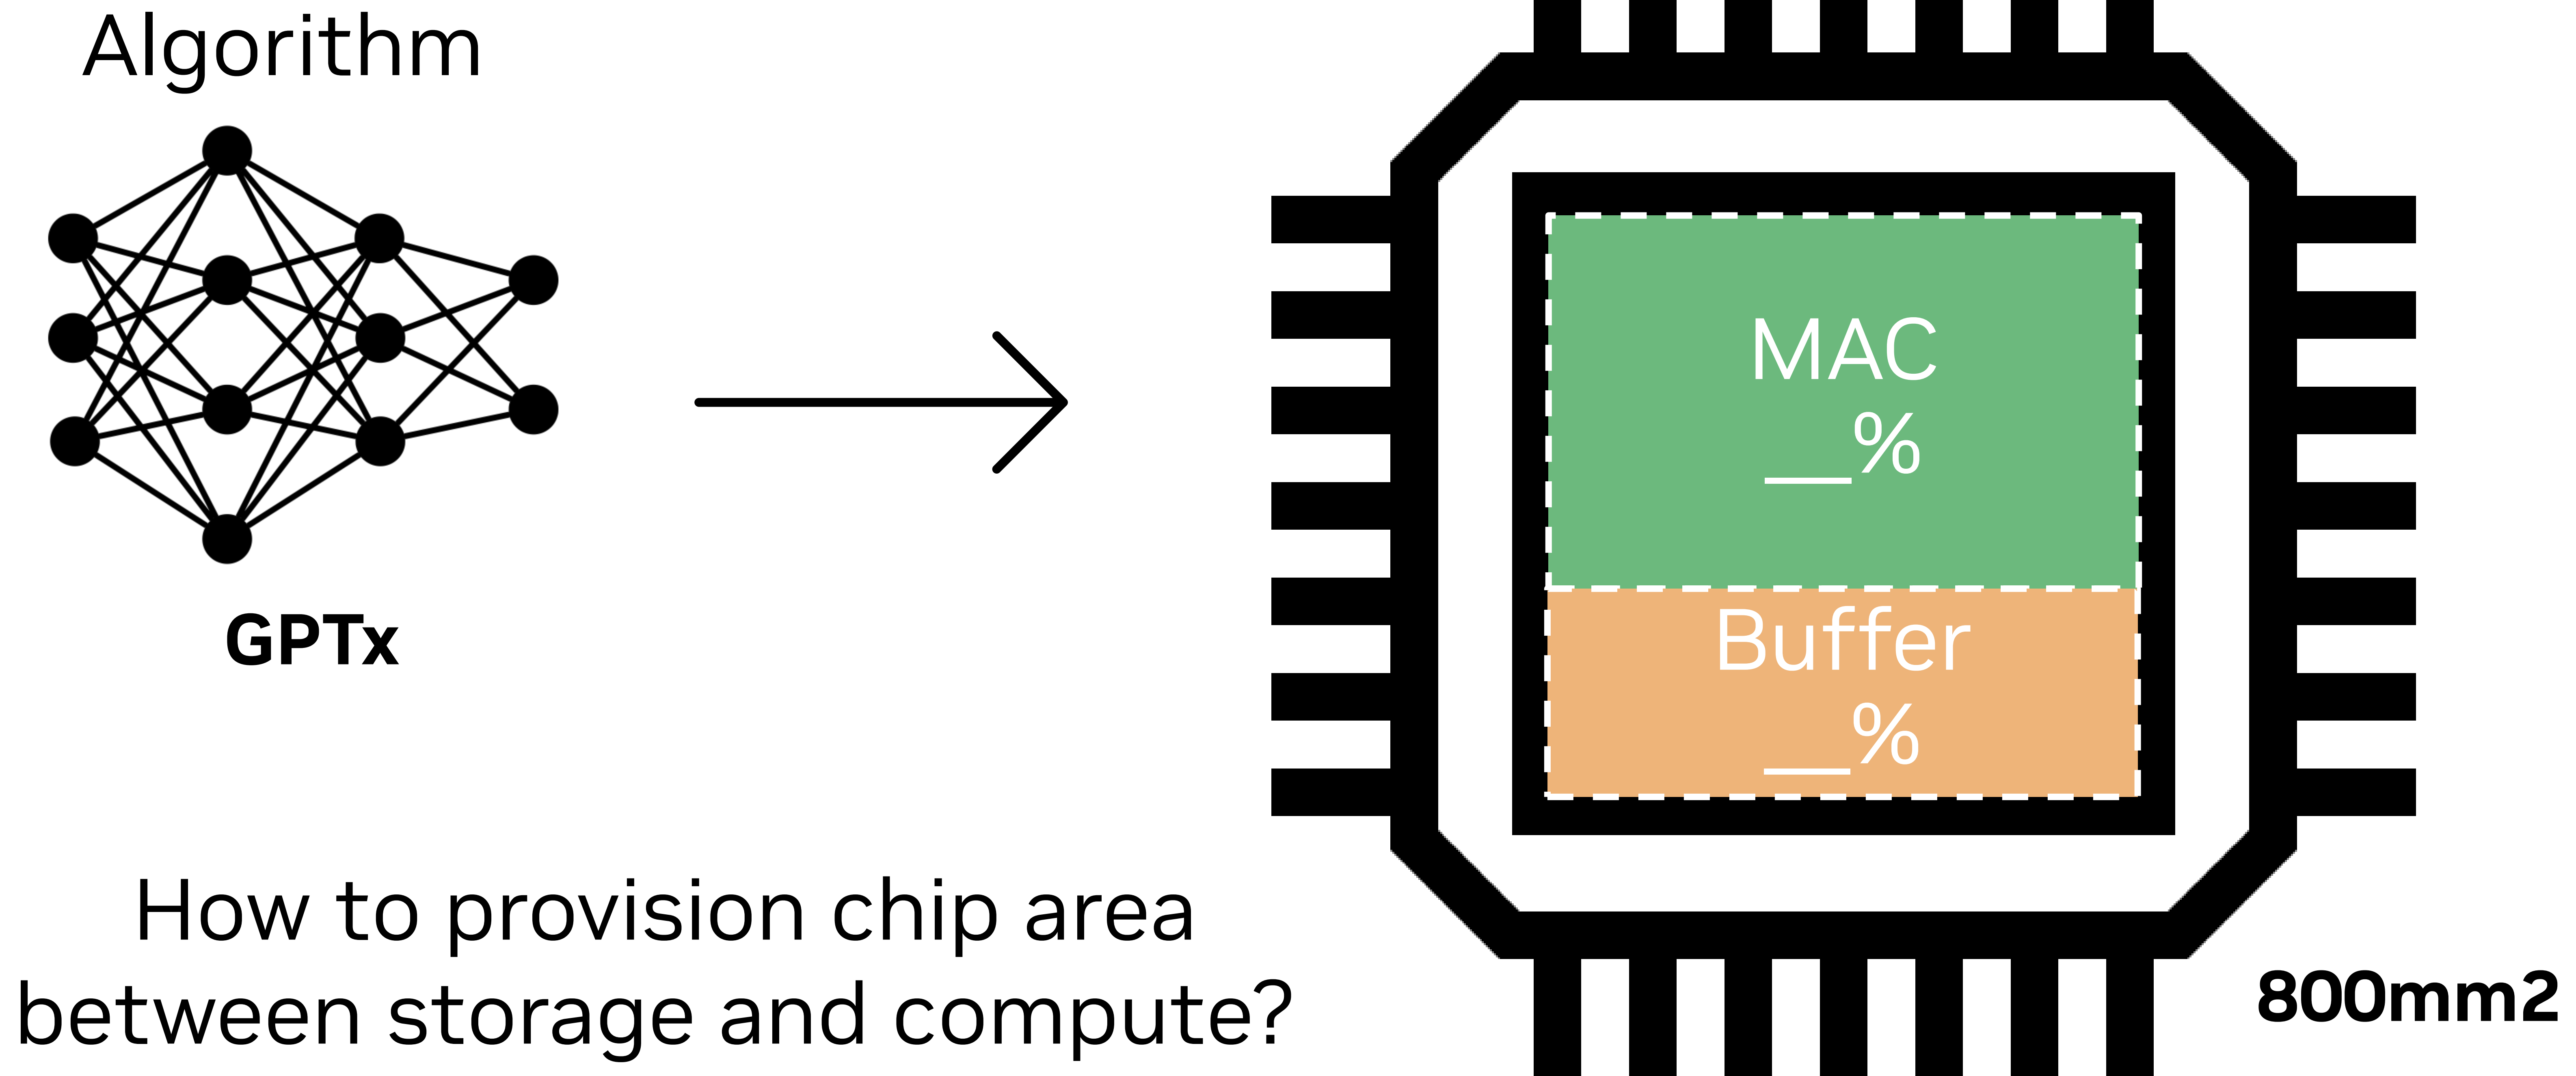
A chain of 6 operations in GPT-6.7b block



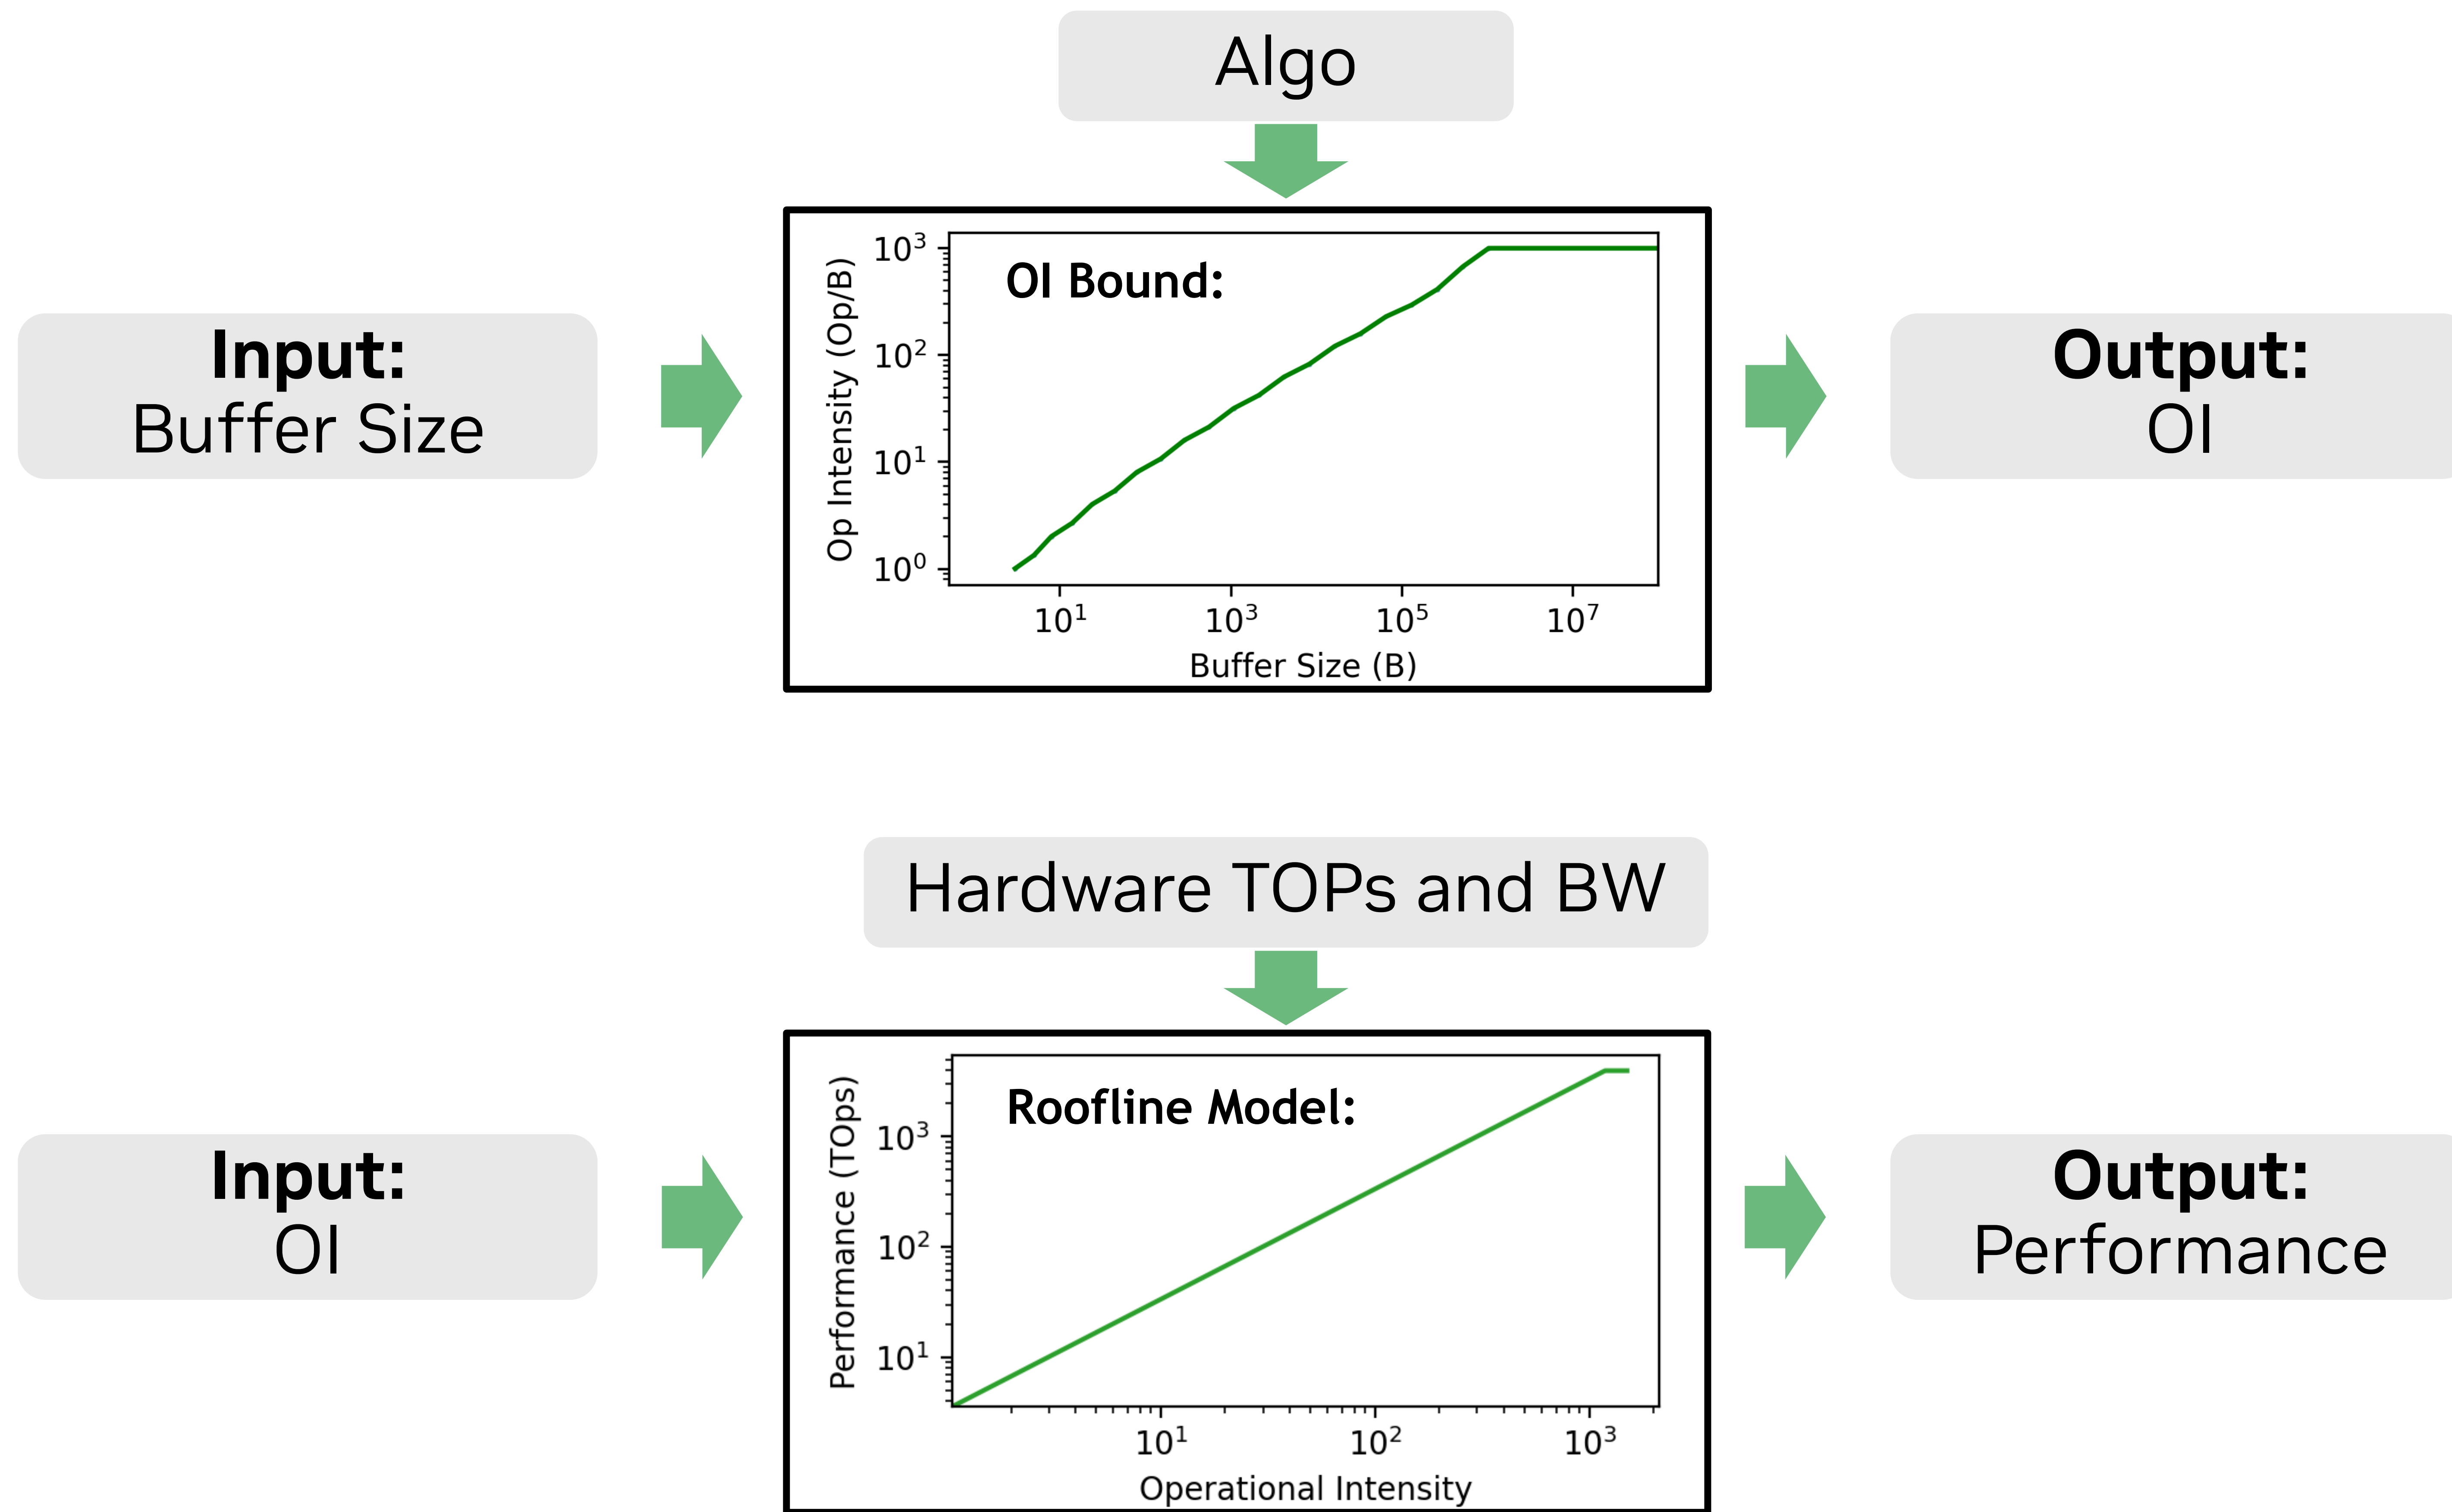
Fusion is effective when buffer size is large

#2: Orojenesis comprehends complex workload optimizations (e.g. fusion)


Motivation: A design challenge



Orojenes Performance Model





Orojenes Performance Model



Orojenesis for DSE

GPT3-6.7b





HW Specs:

Total chip area
Area per Byte
Area per MAC
Backing-store BW
Frequency

Orojenesis for DSE

GPT3-6.7b





HW Specs:

Total chip area
Area per Byte
Area per MAC
Backing-store BW
Frequency

Orojenesis for DSE



GPT3-6.7b



HW Specs:

Total chip area
Area per Byte
Area per MAC
Backing-store BW
Frequency

Motivation: A design challenge



How to provision chip area
between storage and compute?

#3: Orojenesis complements the roofline model to provide buffer size suggestions

Orojenesis

- A radically new approach for early-stage architectural DSE
- Offers **visualization** and **insights** for design tradeoffs
- Can be **a powerful addon** to the roofline performance model



Webpage: <https://timeloop.csail.mit.edu/orojenesis>